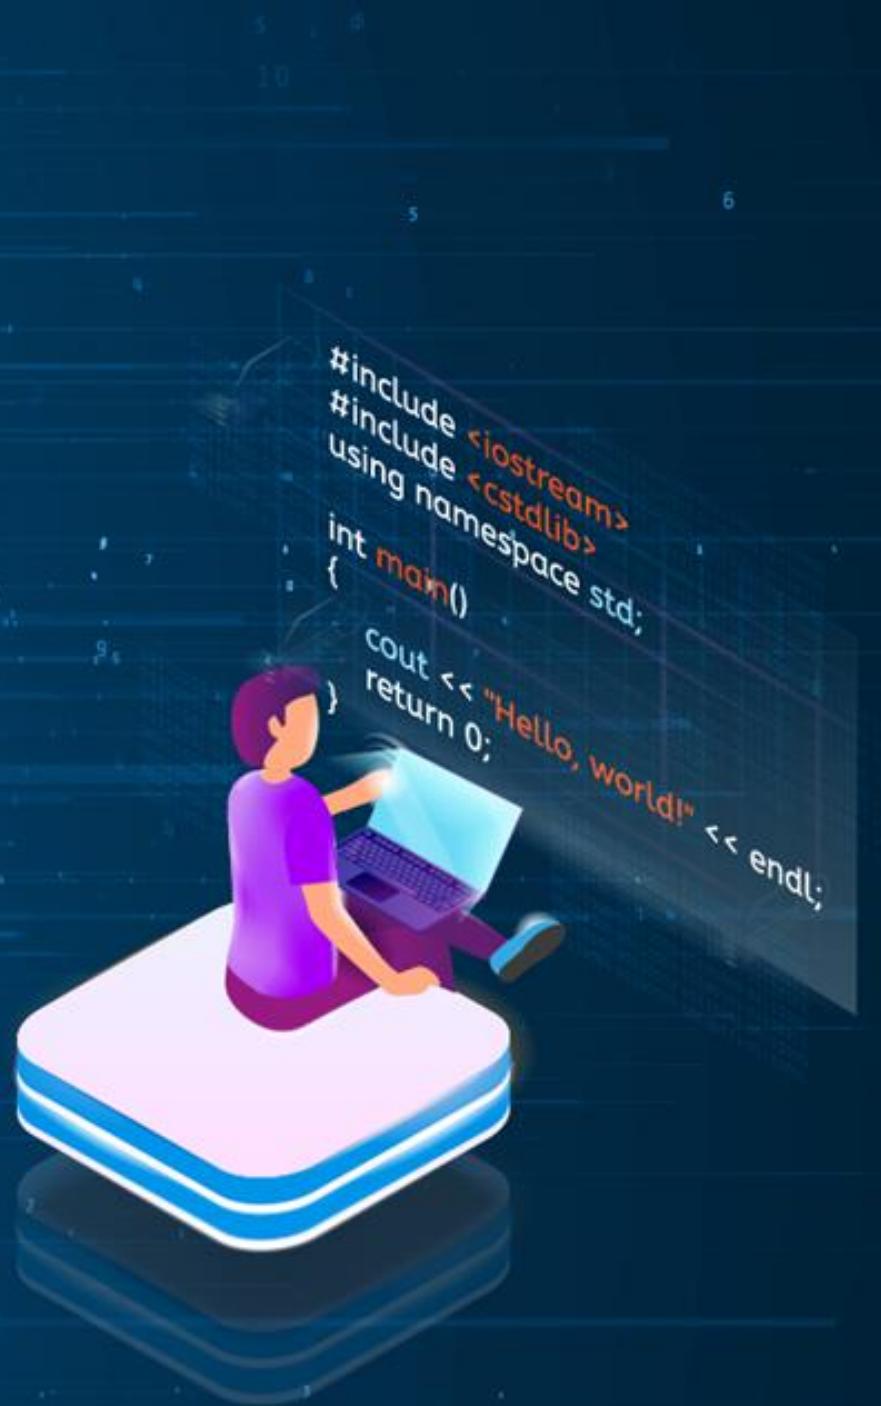


# TECHNOLOGY



## Certified Kubernetes Administrator

## Services, Load Balancing, and Networking



# Learning Objectives

By the end of this lesson, you will be able to:

- State the concerns that Kubernetes networking addresses
- Describe topology and DNS
- Present an overview of EndpointSlices, Ingress, and ingress controllers
- Define network policies and IPv4\_IPv6 dual-stack
- List the ways in which cluster networking is implemented

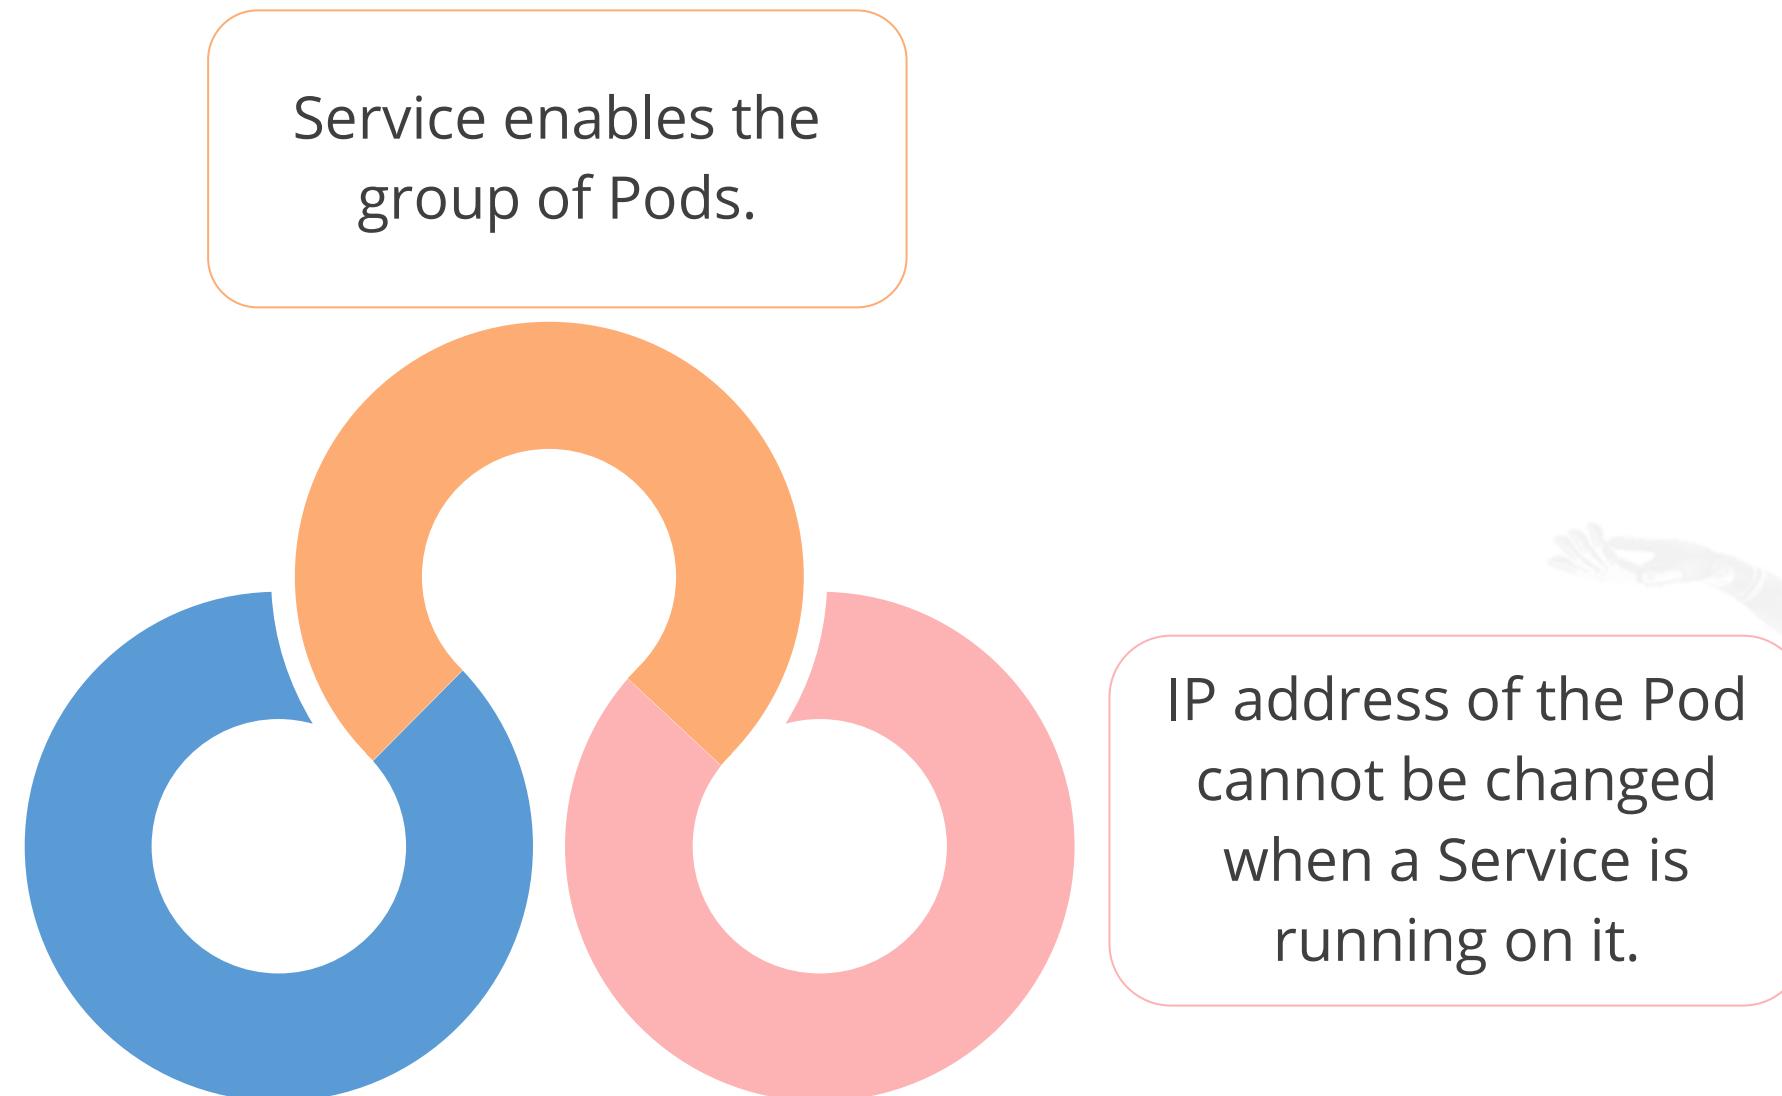


# TECHNOLOGY

## Overview

# Services

Kubernetes Service is a logical abstraction for a deployed group of Pods in a cluster.



# Networking

Kubernetes networking addresses four concerns.



Loopback is used by the Containers within a Pod for communication.



Communication between Pods happens through cluster networking.



An application running in a Pod can be made reachable from outside the cluster using the Service resource.

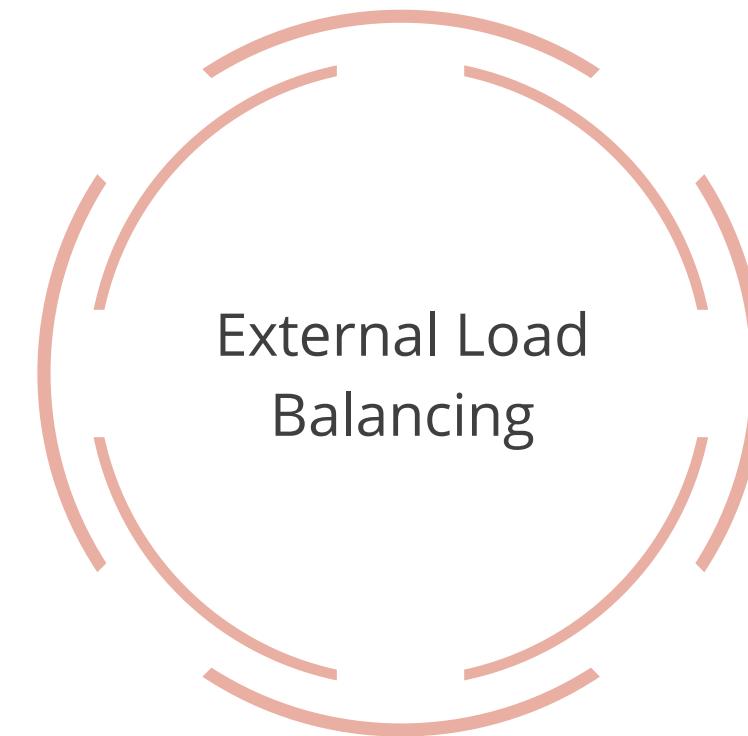
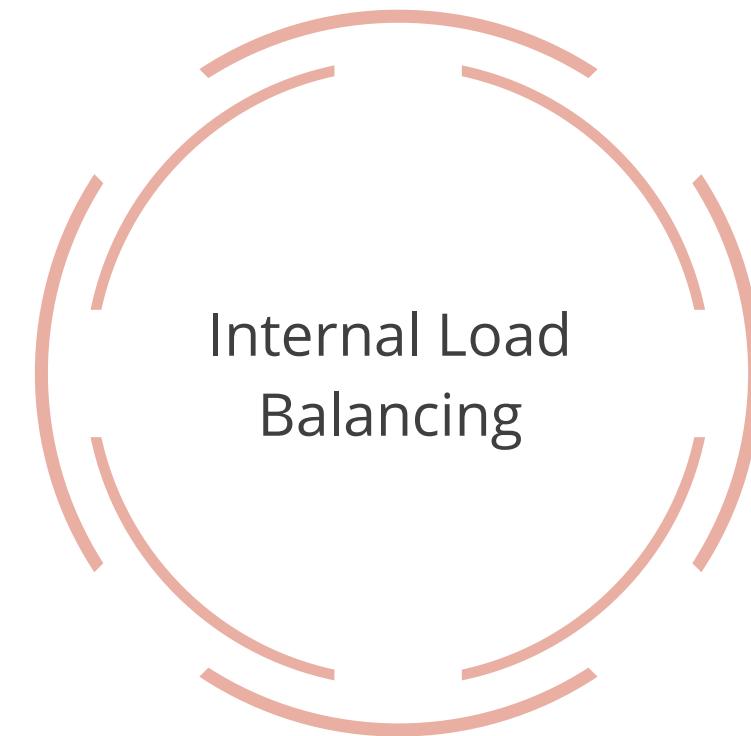


Services can be published only for consumption within the cluster.

# Load Balancing

Load Balancing helps in distributing network traffic or client's request to multiple servers.

Load balancing can be:



# TECHNOLOGY

## Services

# What Is a Service?

A Service in Kubernetes is an abstraction that exposes a set of Pods to be accessed. It also provides a policy for accessing them. The policy is referred to as a micro-service.

Kubernetes makes the Pods discoverable by providing them their own IP addresses and a single DNS name for a set of Pods.

Kubernetes also balances load between Pods.

## Define a Service

A Service is a REST object. To create a new instance of the Service, the Service definition can be POSTed to the API server.

The YAML file shown below exposes a set of Pods that listen on TCP port 9376 and contain a label app, **MyApp**. This specification creates a new Service object, named **my-service**, which targets TCP port 9376 on any Pod with the app, **MyApp** label.

### Demo

```
apiVersion: v1
Kind:     service
Metadata:
    name:   my-service
spec:
    selector:
        app:    MyApp
    ports:
        - protocol: TCP
          port:    80
          targetPort: 9376
```

# Service

1

The default protocol for Service is TCP. However, UDP, SCTP, HTTP, and PROXY are also supported.

2

Kubernetes supports multiple port definitions on a Service object.

# Service Without Selectors

Services mostly abstract access to Kubernetes Pods though they can also abstract other kinds of backends

To define a Service without a Pod Selector, add the following in the YAML file:

Demo

```
apiVersion: v1
Kind:   service
Metadata:
    name:  my-service
spec:
    ports:
        - protocol: TCP
          port: 80
          targetPort: 9376
```

# Service without Selectors

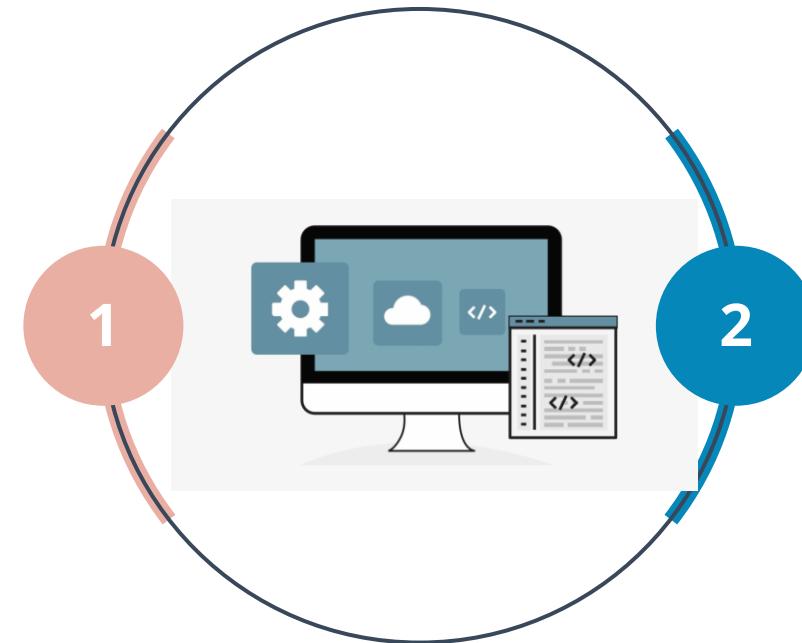
A Service can be mapped manually to the network address and the port where it is running by adding an Endpoints object.

## Demo

```
apiVersion: v1
Kind:   Endpoints
Metadata:
  name:  my-service
subsets:
  - addresses:
      - ip: 192.0.2.42
    ports:
      - port: 9376
```

# Application Protocol

The appProtocol field provides a way to specify an Application Protocol for each Service port.



It follows standard Kubernetes Label Syntax.

Values should either be IANA standard service names or domain prefixed names.

# Virtual IPs and Service Proxies

1

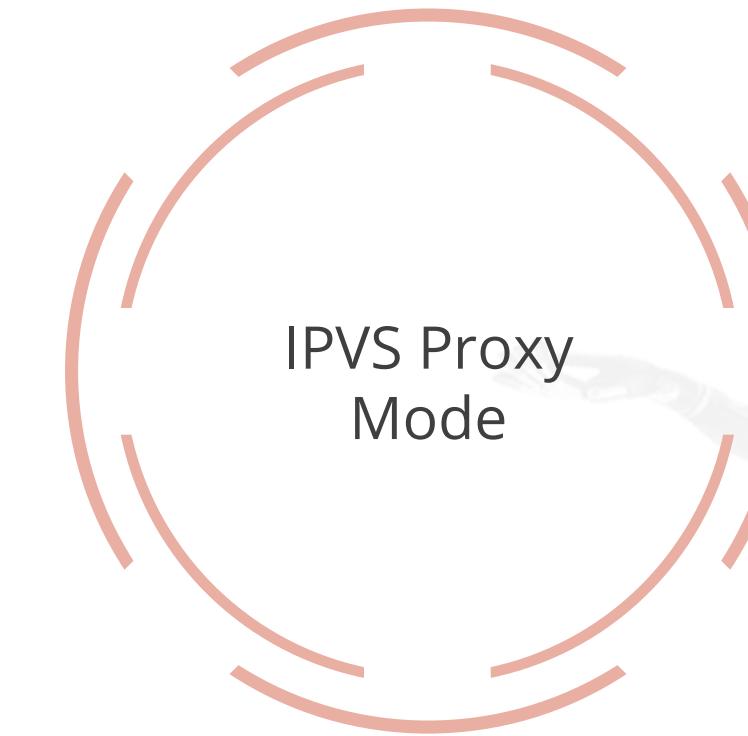
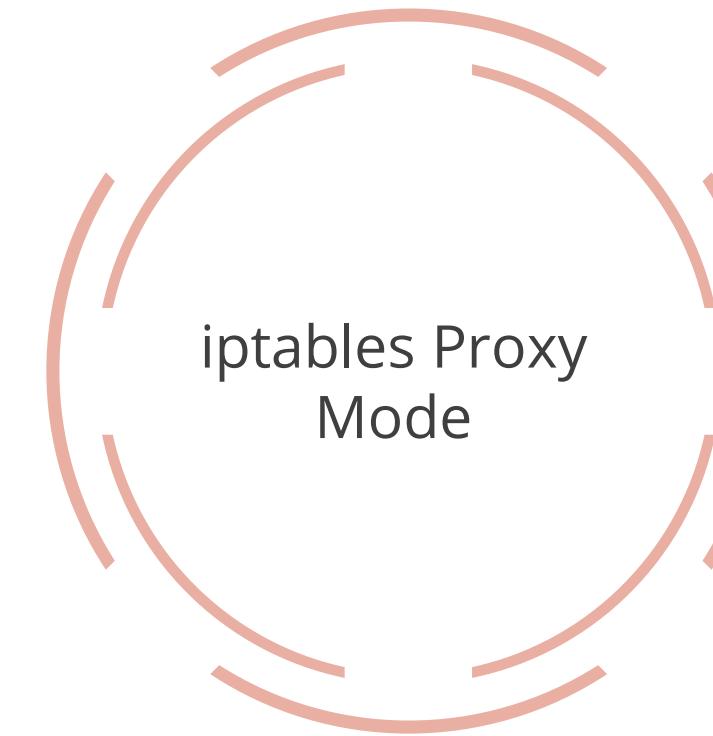
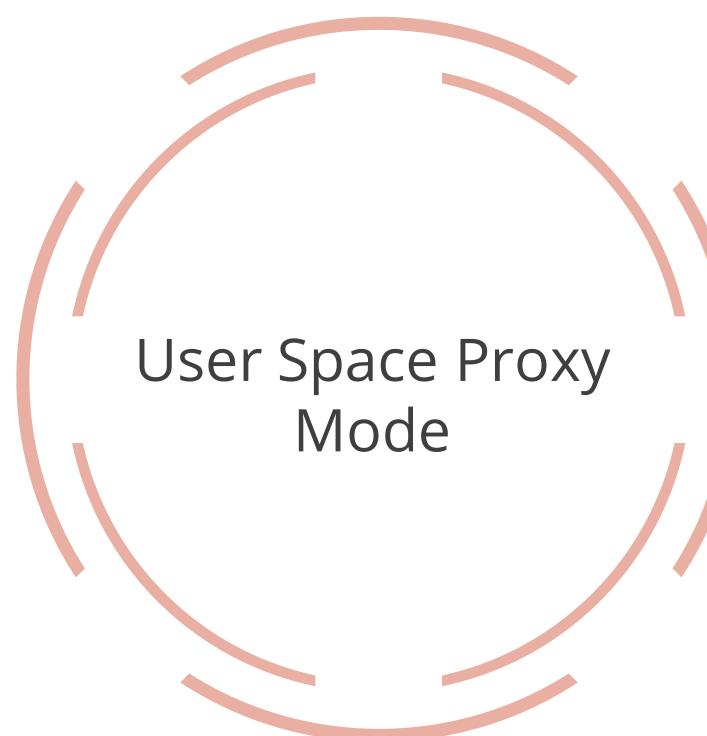
Every Node in a Kubernetes cluster runs a kube-proxy.

2

kube-proxy is responsible for implementing a form of virtual IP for Services of type other than External Name.

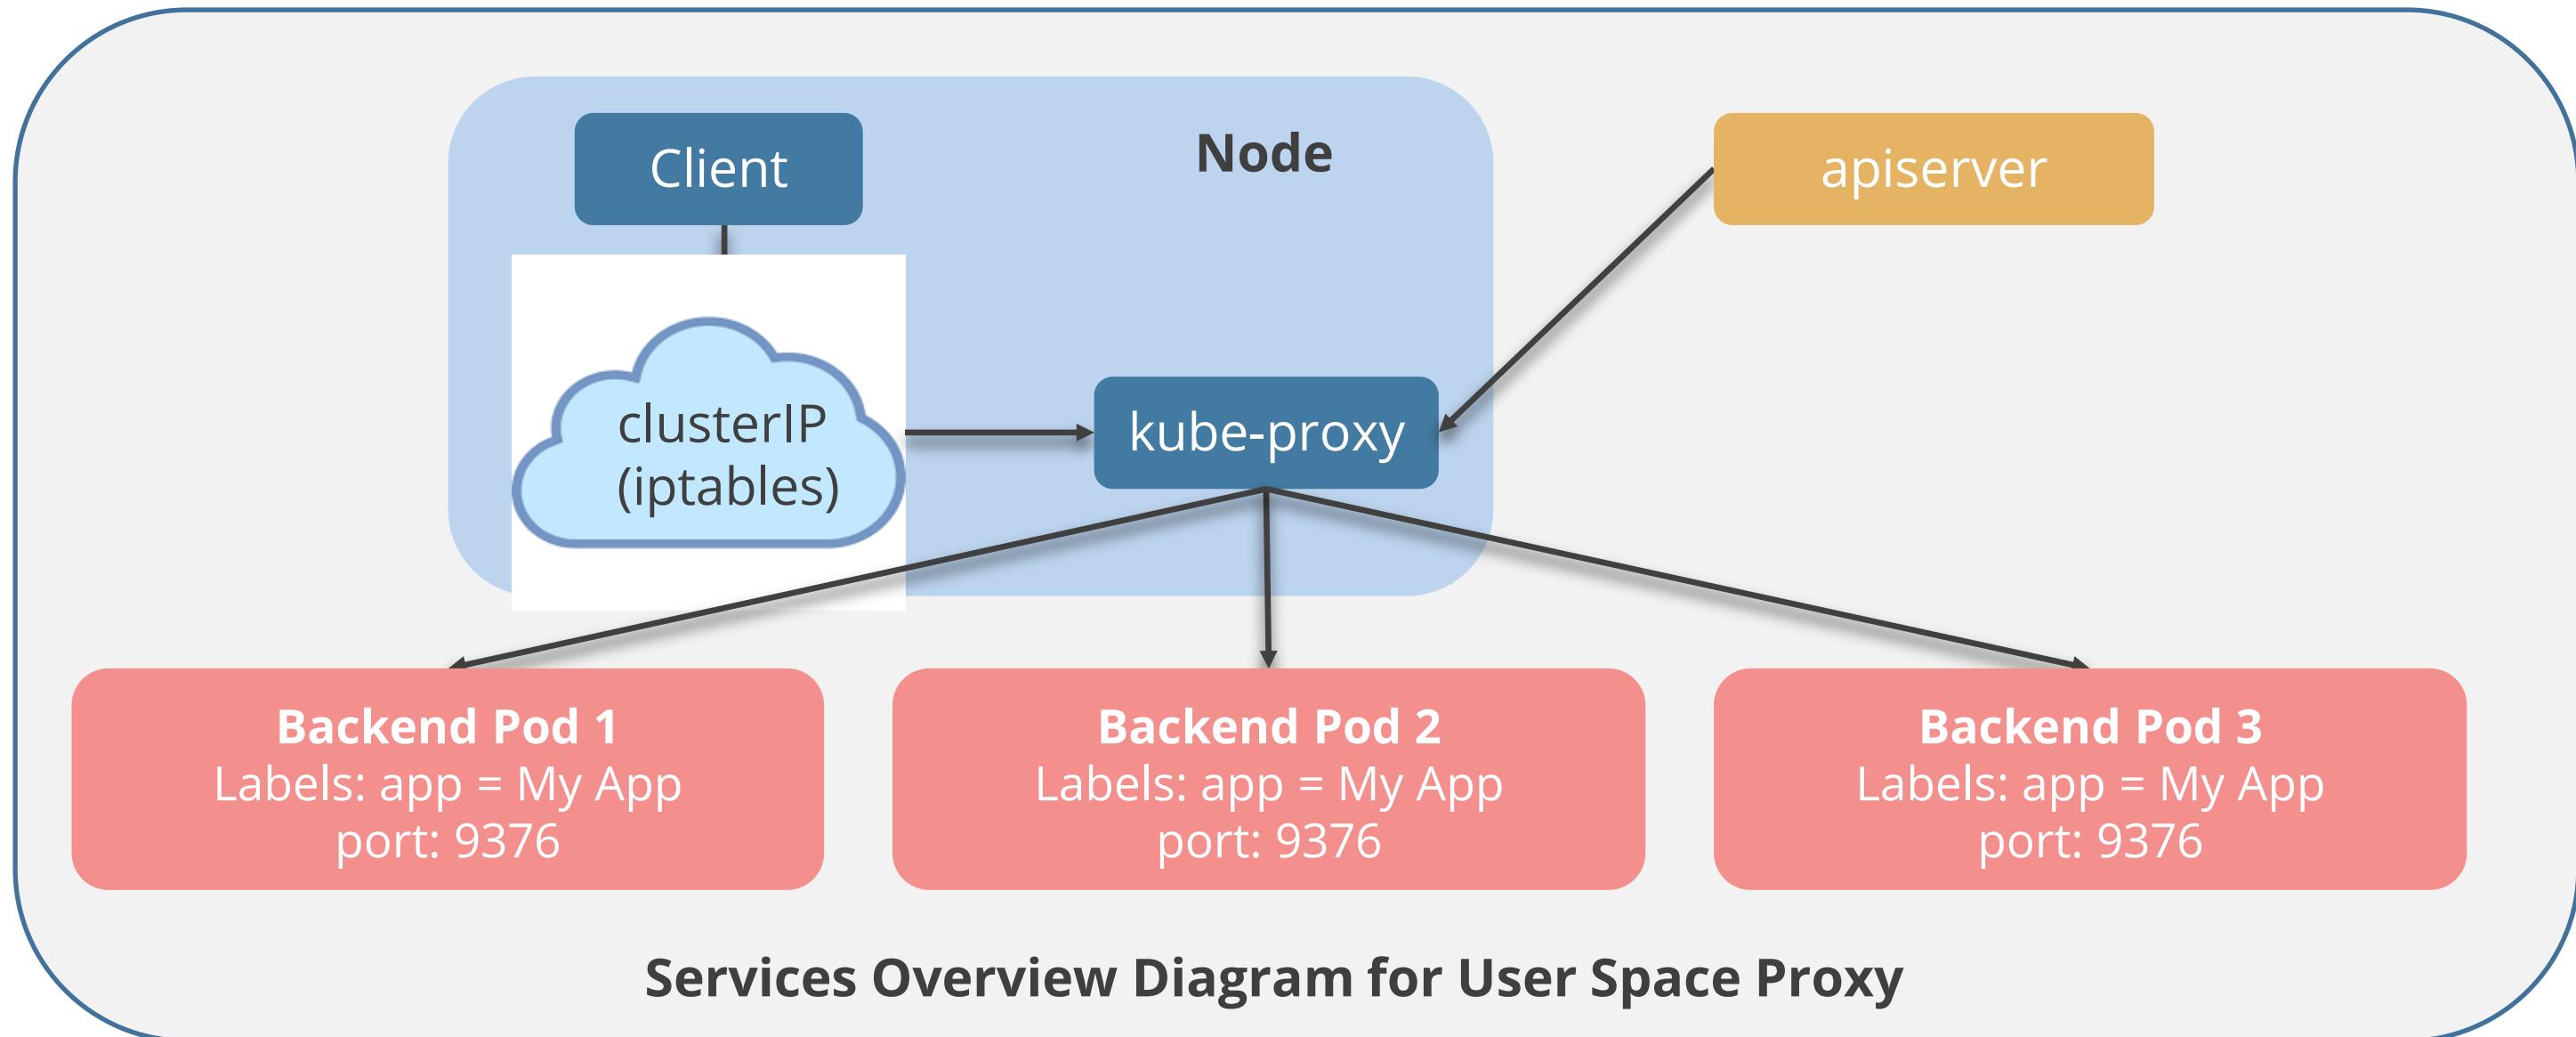
# Proxy Modes

Three types of proxy modes are supported in Kubernetes:



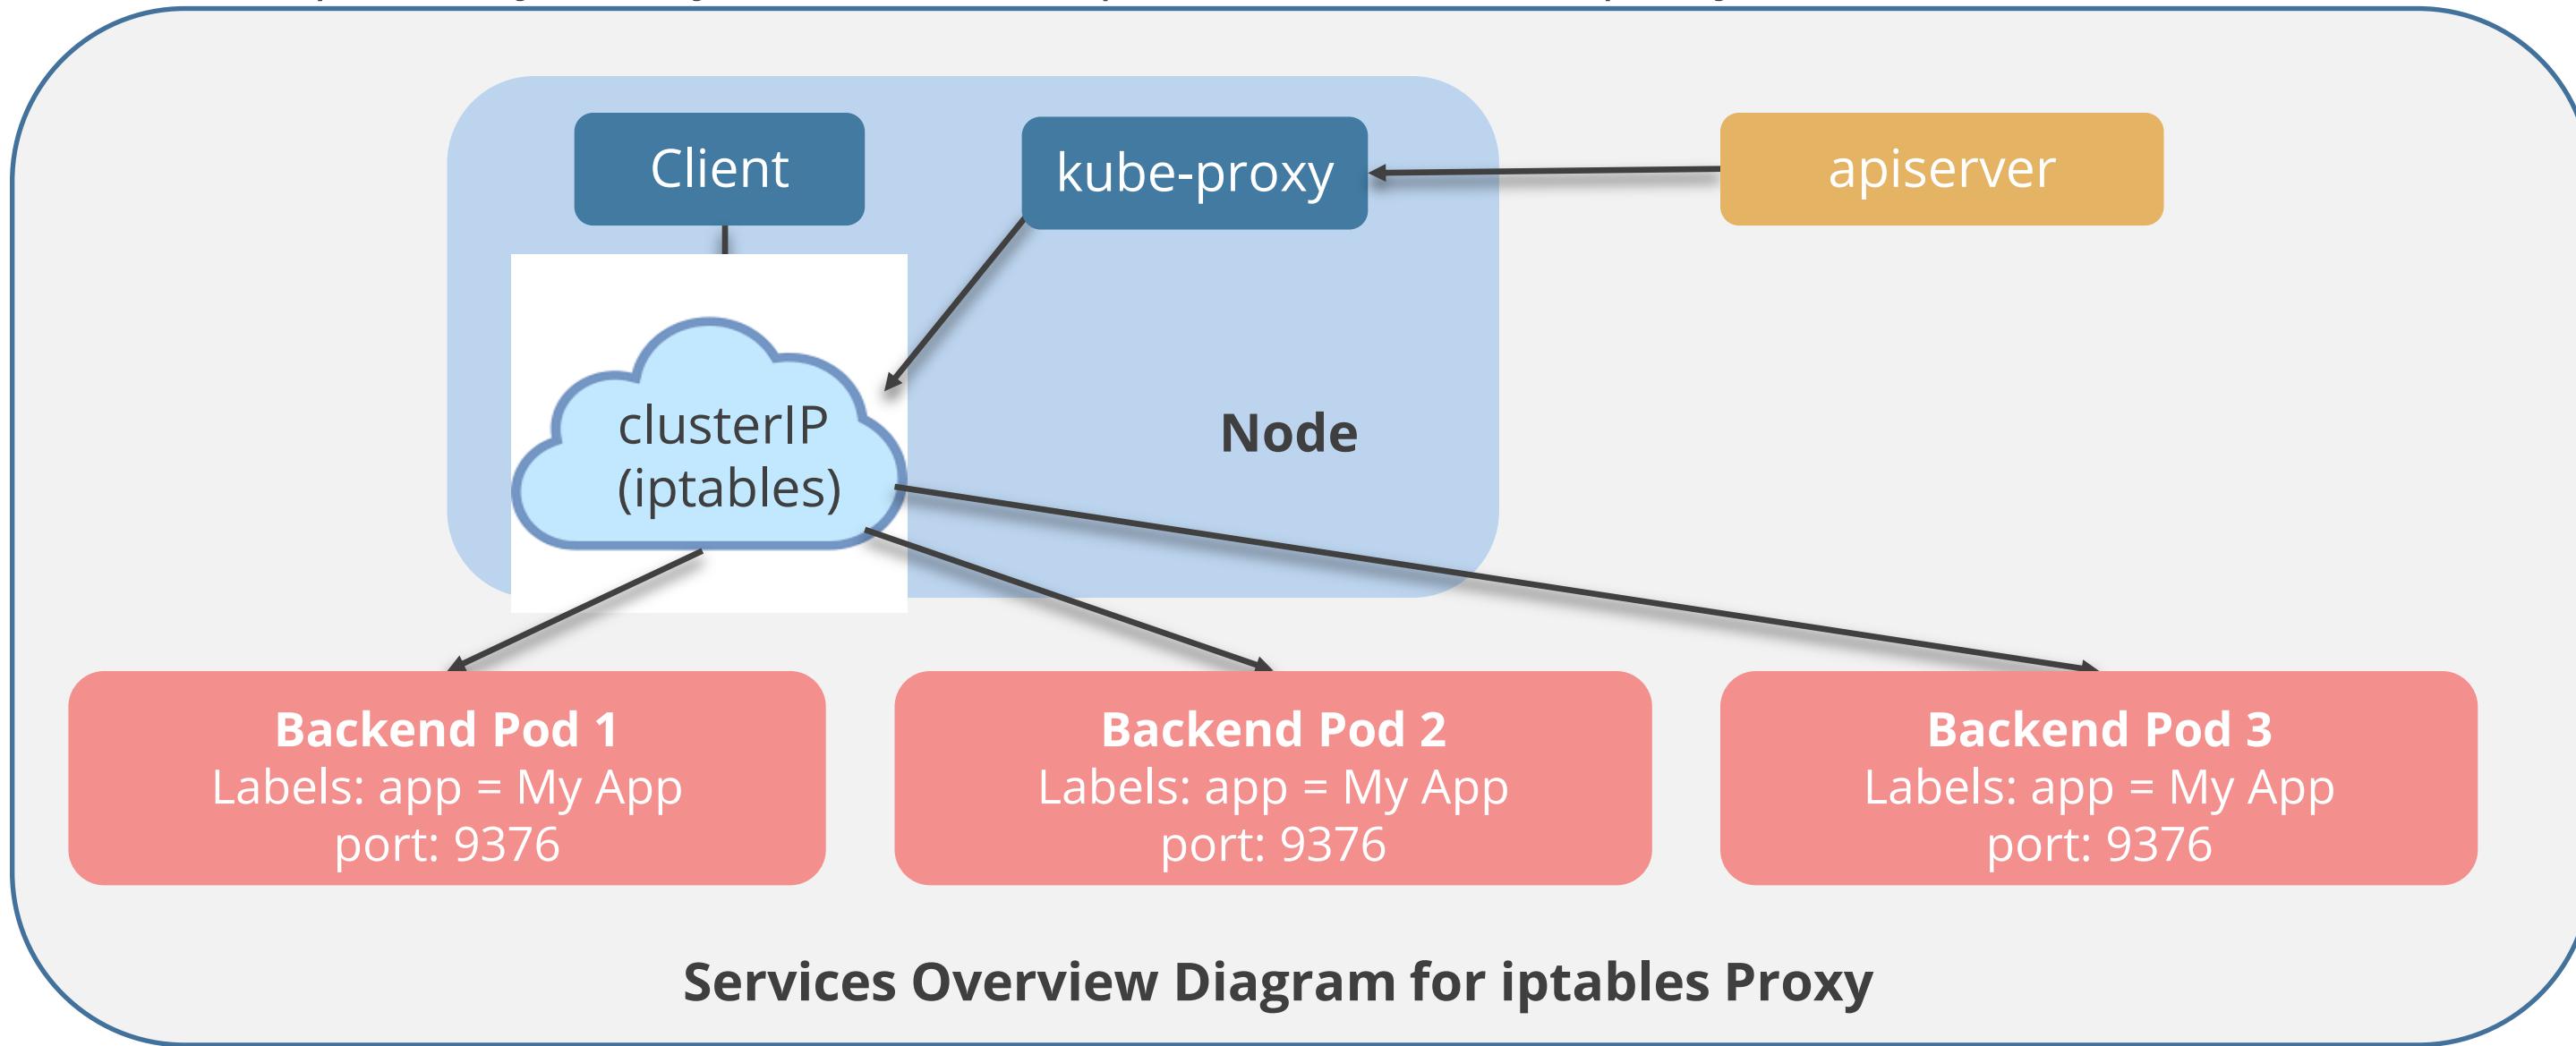
# User Space Proxy Mode

In this mode, kube-proxy watches the Kubernetes Control Plane for the addition and removal of Service and Endpoint objects. It chooses a backend via a round-robin algorithm.



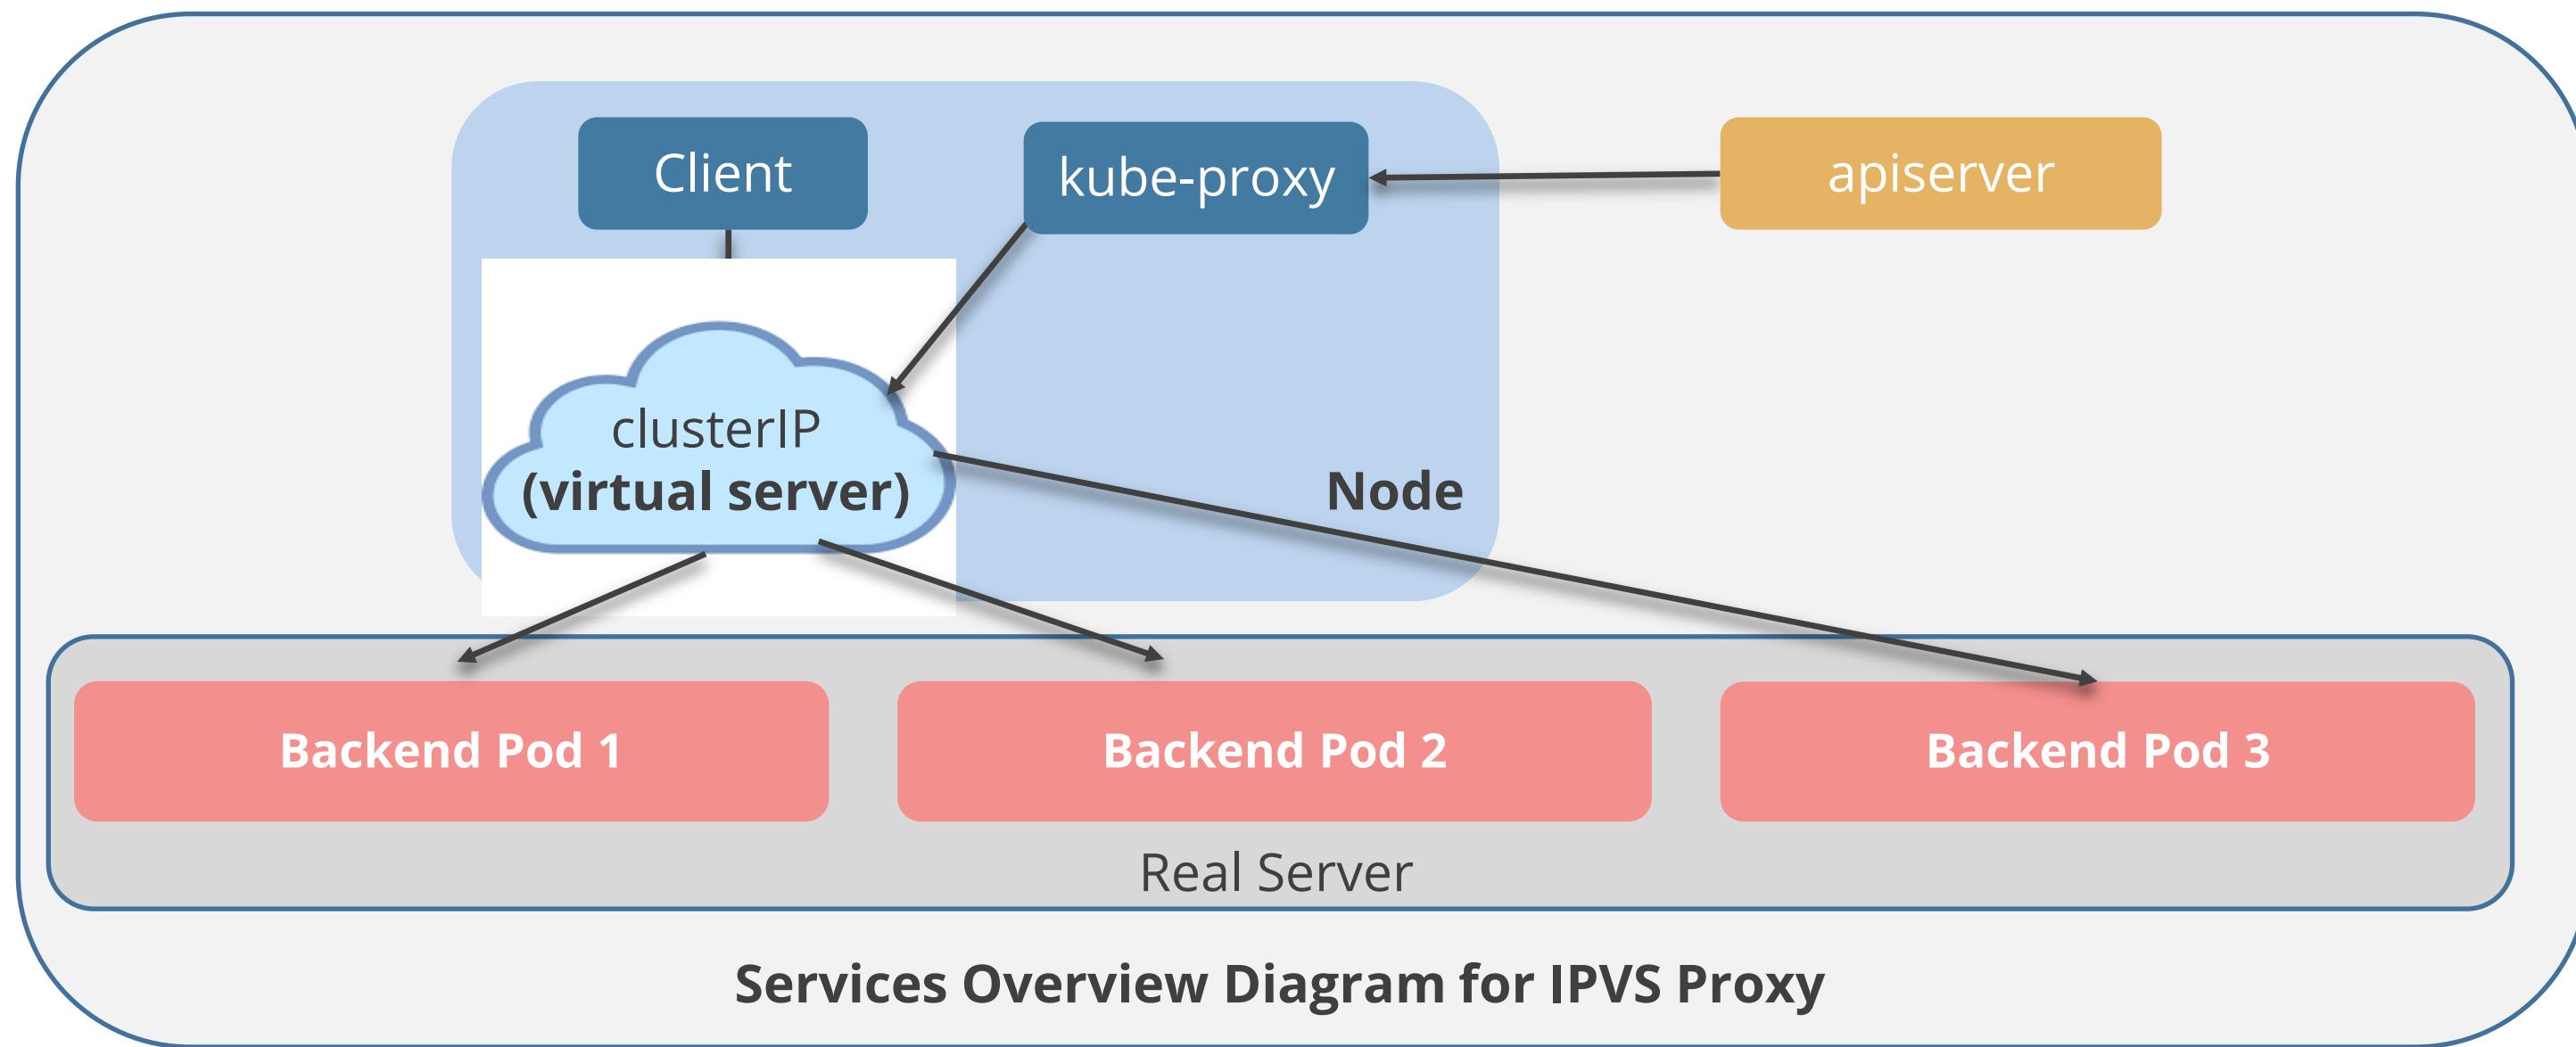
# iptables Proxy Mode

In this mode, kube-proxy watches the Kubernetes Control Plane for the addition and removal of Service and Endpoint objects. By default, in the iptables mode, kube-proxy chooses a backend at random.



## IPVS Proxy Mode

In IPVS proxy mode, kube-proxy watches Kubernetes Services and Endpoints and calls net link interface to create IPVS rules accordingly. It synchronizes IPVS rules with Kubernetes Services and Endpoints periodically.



# Multi-Port Services

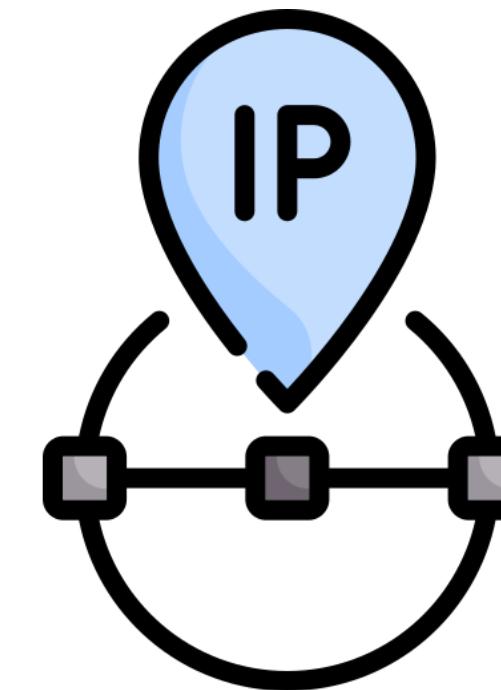
Kubernetes helps to configure multiple port definitions on a Service object. The configuration below shows the use of the **port** attribute:

Demo

```
apiVersion: v1
Kind:   service
Metadata:
  name:  my-service
spec:
  selector:
    app: MyApp
  ports:
    - name: http
      protocol: TCP
      port: 80
    - name: http
      protocol: TCP
      port: 443
      targetPort: 9377
```

## Choose IP Addresses

Cluster IP address can be specified as part of a Service creation request.

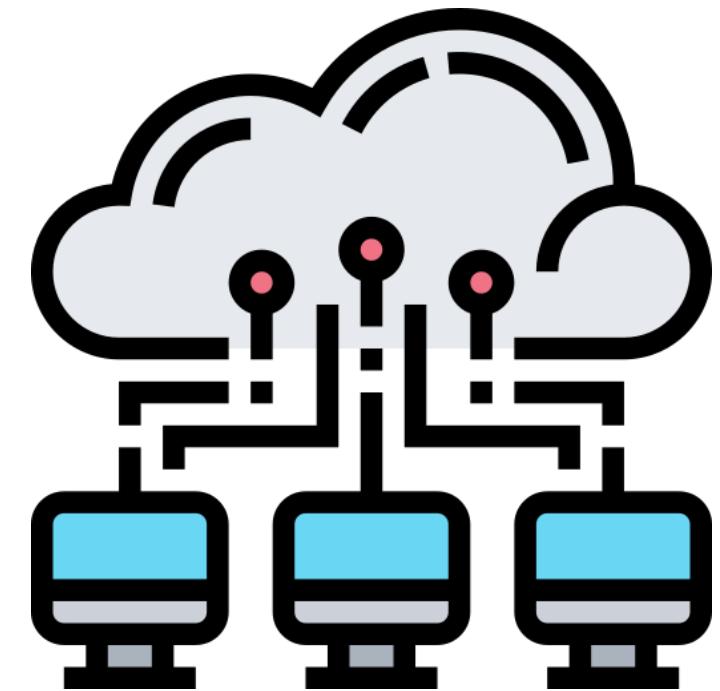


**Note**

The chosen IP address must be a valid IPv4 or IPv6 address from within the service-cluster-ip-range CIDR range that is configured for the API server.

# Discover Services

Services can be discovered using environment variables and DNS.



# Environment Variables

The Kubelet adds a set of environment variables for each active Service when a Pod is run on a Node. The example below shows the environment variables for the Service redis-master, which exposes TCP port 6379 and has been allocated cluster IP address 10.0.0.11.

```
REDIS_MASTER_SERVICE_HOST = 10.0.0.11
REDIS_MASTER_SERVICE_PORT= 6379
REDIS_MASTER_PORT= tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP = tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP_PROTO= tcp
REDIS_MASTER_PORT_6379_TCP_PORT= 6379
REDIS_MASTER_PORT_6379_TCP_ADDP= 10.0.0.11
```

## DNS

A cluster-aware DNS server, such as CoreDNS, watches the Kubernetes API for new Services and creates a set of DNS records for each. The admin must almost always set up a DNS service for the Kubernetes cluster using an add-on. The Kubernetes DNS server is the only way to access ExternalName Services.

### Example:

For a Service called **my-service** in a Kubernetes namespace **my-ns**, the Control Plane and the DNS Service together create a DNS record for **my-service.my-ns**.

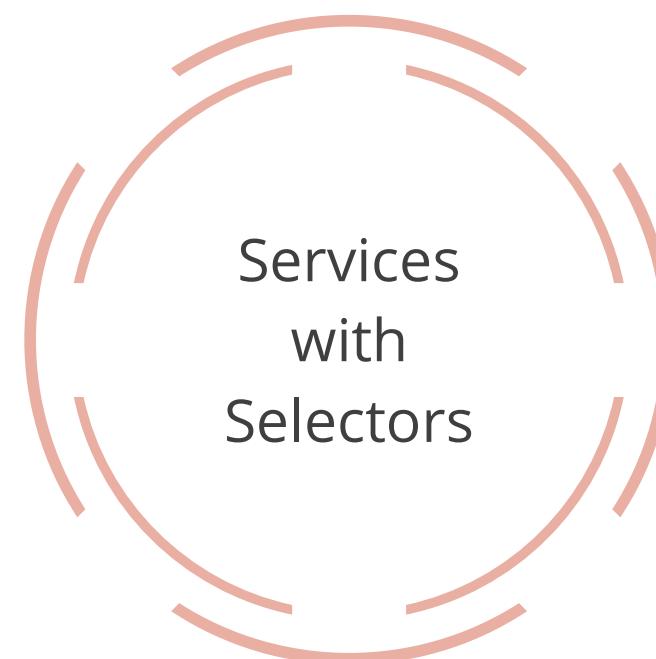
Pods in the **my-ns** namespace should be able to find the service by doing a name lookup for **my-service**.

Pods in other namespaces must qualify the name as **my-service.my-ns**.

# Headless Service

A Headless Service is a service that has no Service IP. They are used to interface with other service discovery mechanisms, without being tied to the implementation of Kubernetes.

Headless Services may be classified into two categories:



# Publishing Services (ServiceTypes)

Kubernetes ServiceTypes help to specify the required kind of service.

Type values and their behaviors include:

ClusterIP

ExternalName

NodePort

LoadBalancer

# Type NodePort

The **NodePort** Service Type allows setting up a custom Load Balancing solution for configuring environments that are not fully supported by Kubernetes. The **NodePort** type can be configured as shown below:

## Demo

```
Apiversion: v1
Kind: service
Metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app: MyApp
  ports:
    # By default, and for convenience , the 'targetPort' is
    # set to the same value as the ' port' field.
    - port: 80
      targetPort: 80
      # Optional field
      # By default, and for convenience, the kubernetes control
      plane will allocate a port from a range ( default: 30000-32767)
      NodePort: 30007
```

# Type LoadBalancer

The LoadBalancer type provisions Load Balancers for cloud providers that support external Load Balancers. Load Balancers are created asynchronously. **status.loadBalancer** field contains the provisioned balancer's information.

## Demo

```
apiVersion: v1
Kind:   service
Metadata:
  name:  my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
  clusterIP: 10.0.171.239
  type: LoadBalancer
status:
  LoadBalancer: 9377
  ingress:
    -ip: 192.0.2.127
```

To direct traffic from the external Load Balancer to the backend Pods, use the configuration shown.

## Load Balancers with Mixed Protocol Types

For **LoadBalancer** type of Services, when there is more than one port defined, all ports, by default, must have the same protocol. The protocol must be supported by the cloud provider.

If the feature gate **MixedProtocolILBService** is enabled for the kube-apiserver, it can use different protocols when there is more than one port defined.

# Disable Load Balancer NodePort Allocation

Starting in v1.20, Node port allocation or a **Service Type=LoadBalancer** can be optionally disabled. This can be done by setting the field **spec.allocateLoadBalancerNodePorts** to false. It can only be used for Load Balancer implementations that route traffic directly to Pods.

By default, **spec.allocateLoadBalancerNodePorts** is true and type LoadBalancer Services will continue to allocate Node ports.

If **spec.allocateLoadBalancerNodePorts** is set to false on an existing Service with allocated Node ports, these Node ports will not be de-allocated automatically.

**ServiceLBNodePortControl** feature gate must be enabled to use this field.

## Topology

# Topology-Aware Traffic Routing

Service Topology enables a Service to route traffic, based on the Node topology of the cluster.

By default, traffic sent to a **ClusterIP** or **NodePort** Service can be routed to any backend address for the Service.

The label matching between the source and destination lets the cluster operator designate sets of Nodes that are closer and farther from one another.

# Usage of Service Topology

Service traffic routing can be controlled if the ServiceTopology feature of the cluster is enabled. This can be done by specifying the **topologyKeys** field on the service spec, a preference-order list of Node labels.

Traffic will be directed to the Node whose first label value matches the originating Node's value.

Topology constraints will not be applied if topologyKeys is empty or unspecified.



The second label will be considered if there is no backend for the Service on a matching Node. It will move until no labels remain.

# Constraints

1

Service topology is not compatible with **externalTrafficPolicy=Local**, and therefore, a Service cannot use both these features.

2

Valid Topology keys are currently limited to kubernetes.io/hostname, topology.kubernetes.io/zone, and topology.kubernetes.io/region.

3

Topology keys must be valid label keys; at most, 16 keys may be specified.

4

If the **catch-all value \*** is used, it must be the last value in the Topology keys.

# Understanding Service Topology



**Problem Statement:** Understand the working of service topology in Kubernetes.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

---

## Steps to demonstrate topology in Kubernetes:

1. Route through local Endpoints only
2. Route through local Endpoints as a preference

# TECHNOLOGY

## Service Catalog

## Overview

Service Catalog is an extension API that enables applications running in Kubernetes clusters to easily use externally managed software offerings.



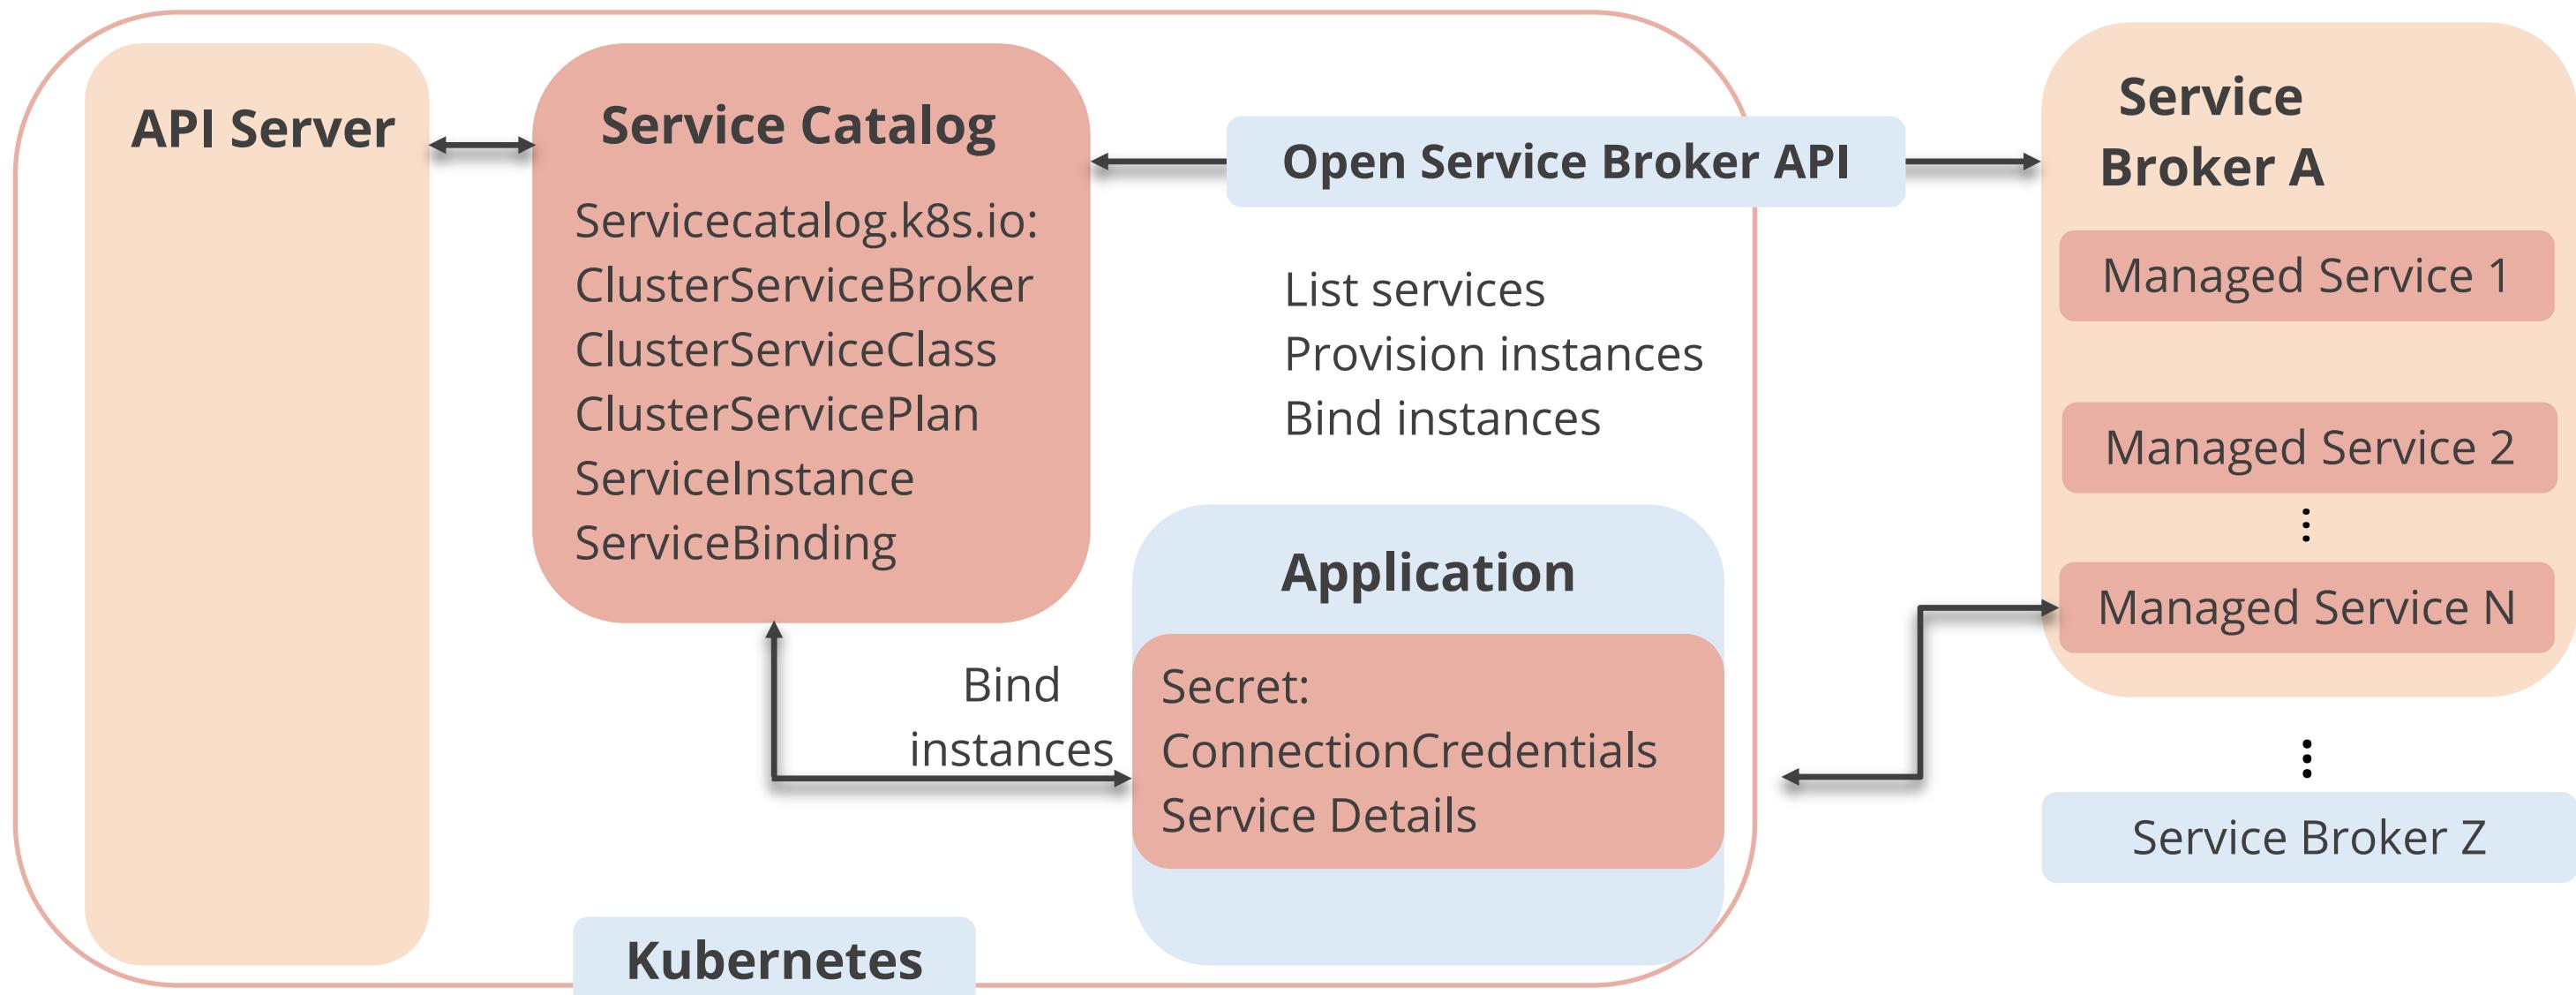
Service Catalog helps to list, provision, and bind with external Managed Services from Service Brokers by abstracting service creation and management.



It acts as an intermediary for the Kubernetes API Server to negotiate the initial provisioning and retrieve the credentials necessary for the application to use a Managed Service.

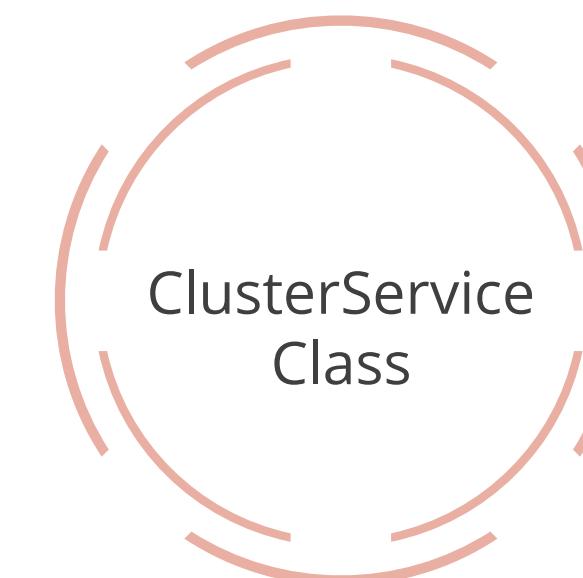
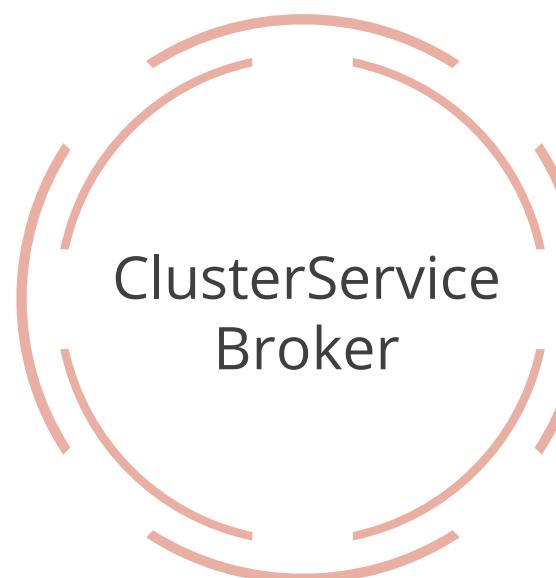
# Architecture

Service Catalog uses the Open Service Broker API to communicate with Service Brokers.



## API Resources

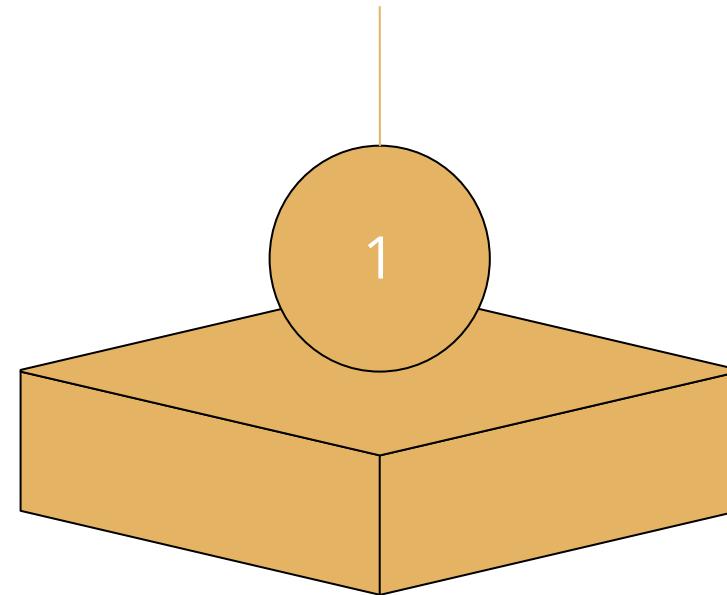
Service Catalog installs the `servicecatalog.k8s.io` API and provides the following Kubernetes resources:



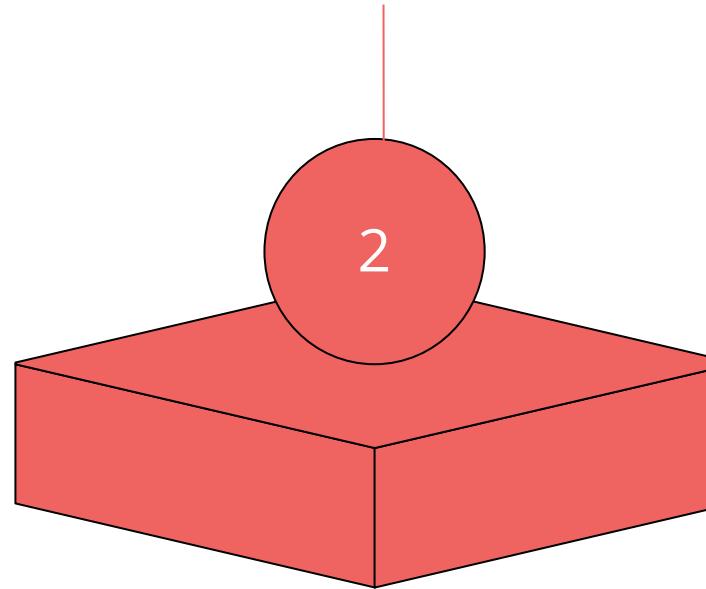
# Authentication

Service Catalog supports two methods of authentication:

Basic (username/password)

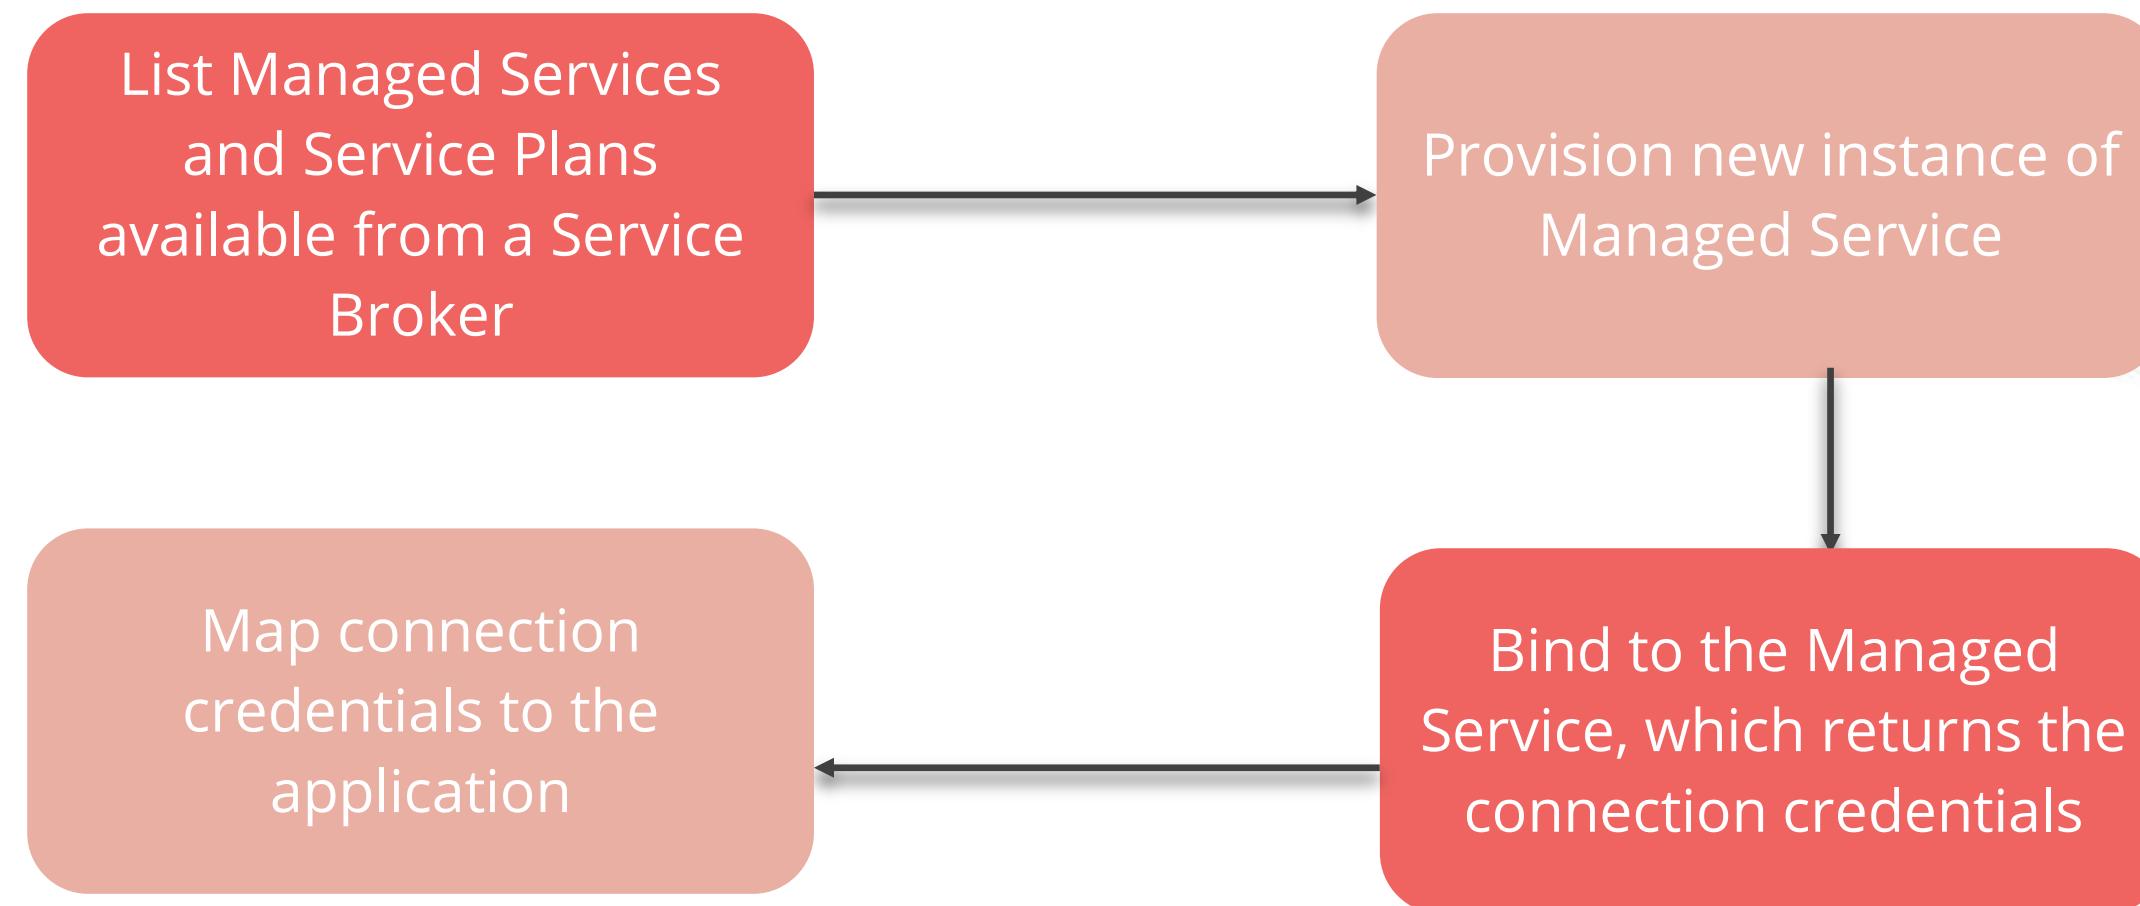


OAuth 2.0 Bearer Token



## Usage

Service Catalog API Resources are used to provision Managed Services and make them available within a Kubernetes cluster, employing a four-step process:



# List Managed Services and Service Plans

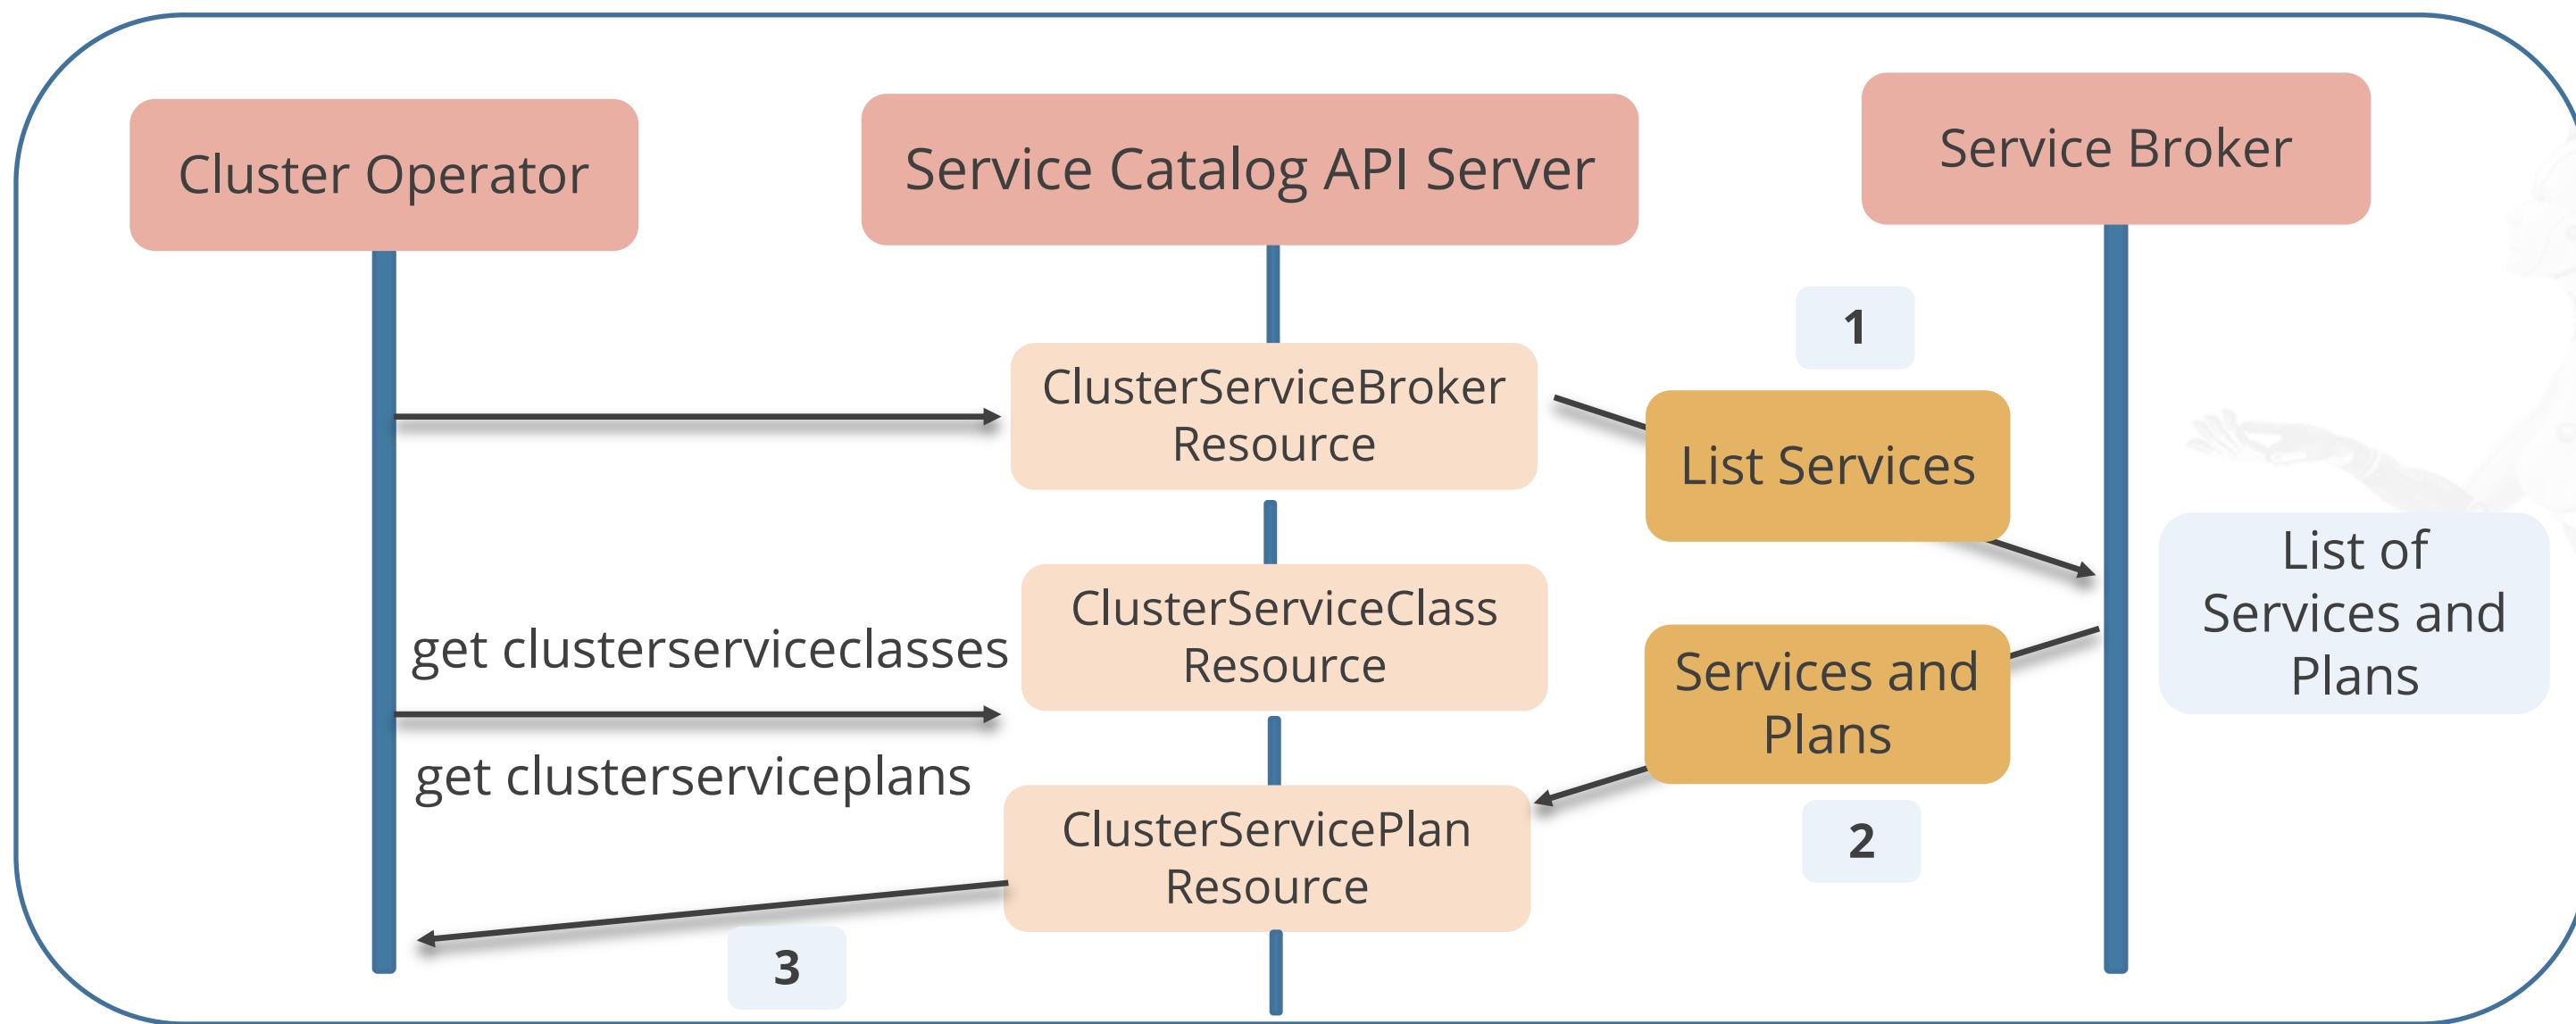
A ClusterServiceBroker resource is created within the servicecatalog.k8s.io group. This resource contains the URL and connection details necessary to access a Service Broker Endpoint. Shown below is an example of a ClusterServiceBroker resource:

## Demo

```
apiVersion: servicecatalog.k8s.io/v1beta1
Kind: ClusterServiceBroker
Metadata:
  name: cloud-broker
spec:
#Points to the endpoint of a service broker. ( This example is not a
working URL.)
  URL:
https://servicebroker.somecloudprovider.com/v1alpha1/projects/service-
catalog/brokers/default
  #####
  #Additional values can be added here, which may be used to
communicate
  #with the service broker, such as bearer token info or a caBundle for
TLS
  #####
```

# List Managed Services and Service Plans

Here is a sequence diagram illustrating the steps involved in listing Managed Services and Plans available from a Service Broker:



# Provision a New Instance

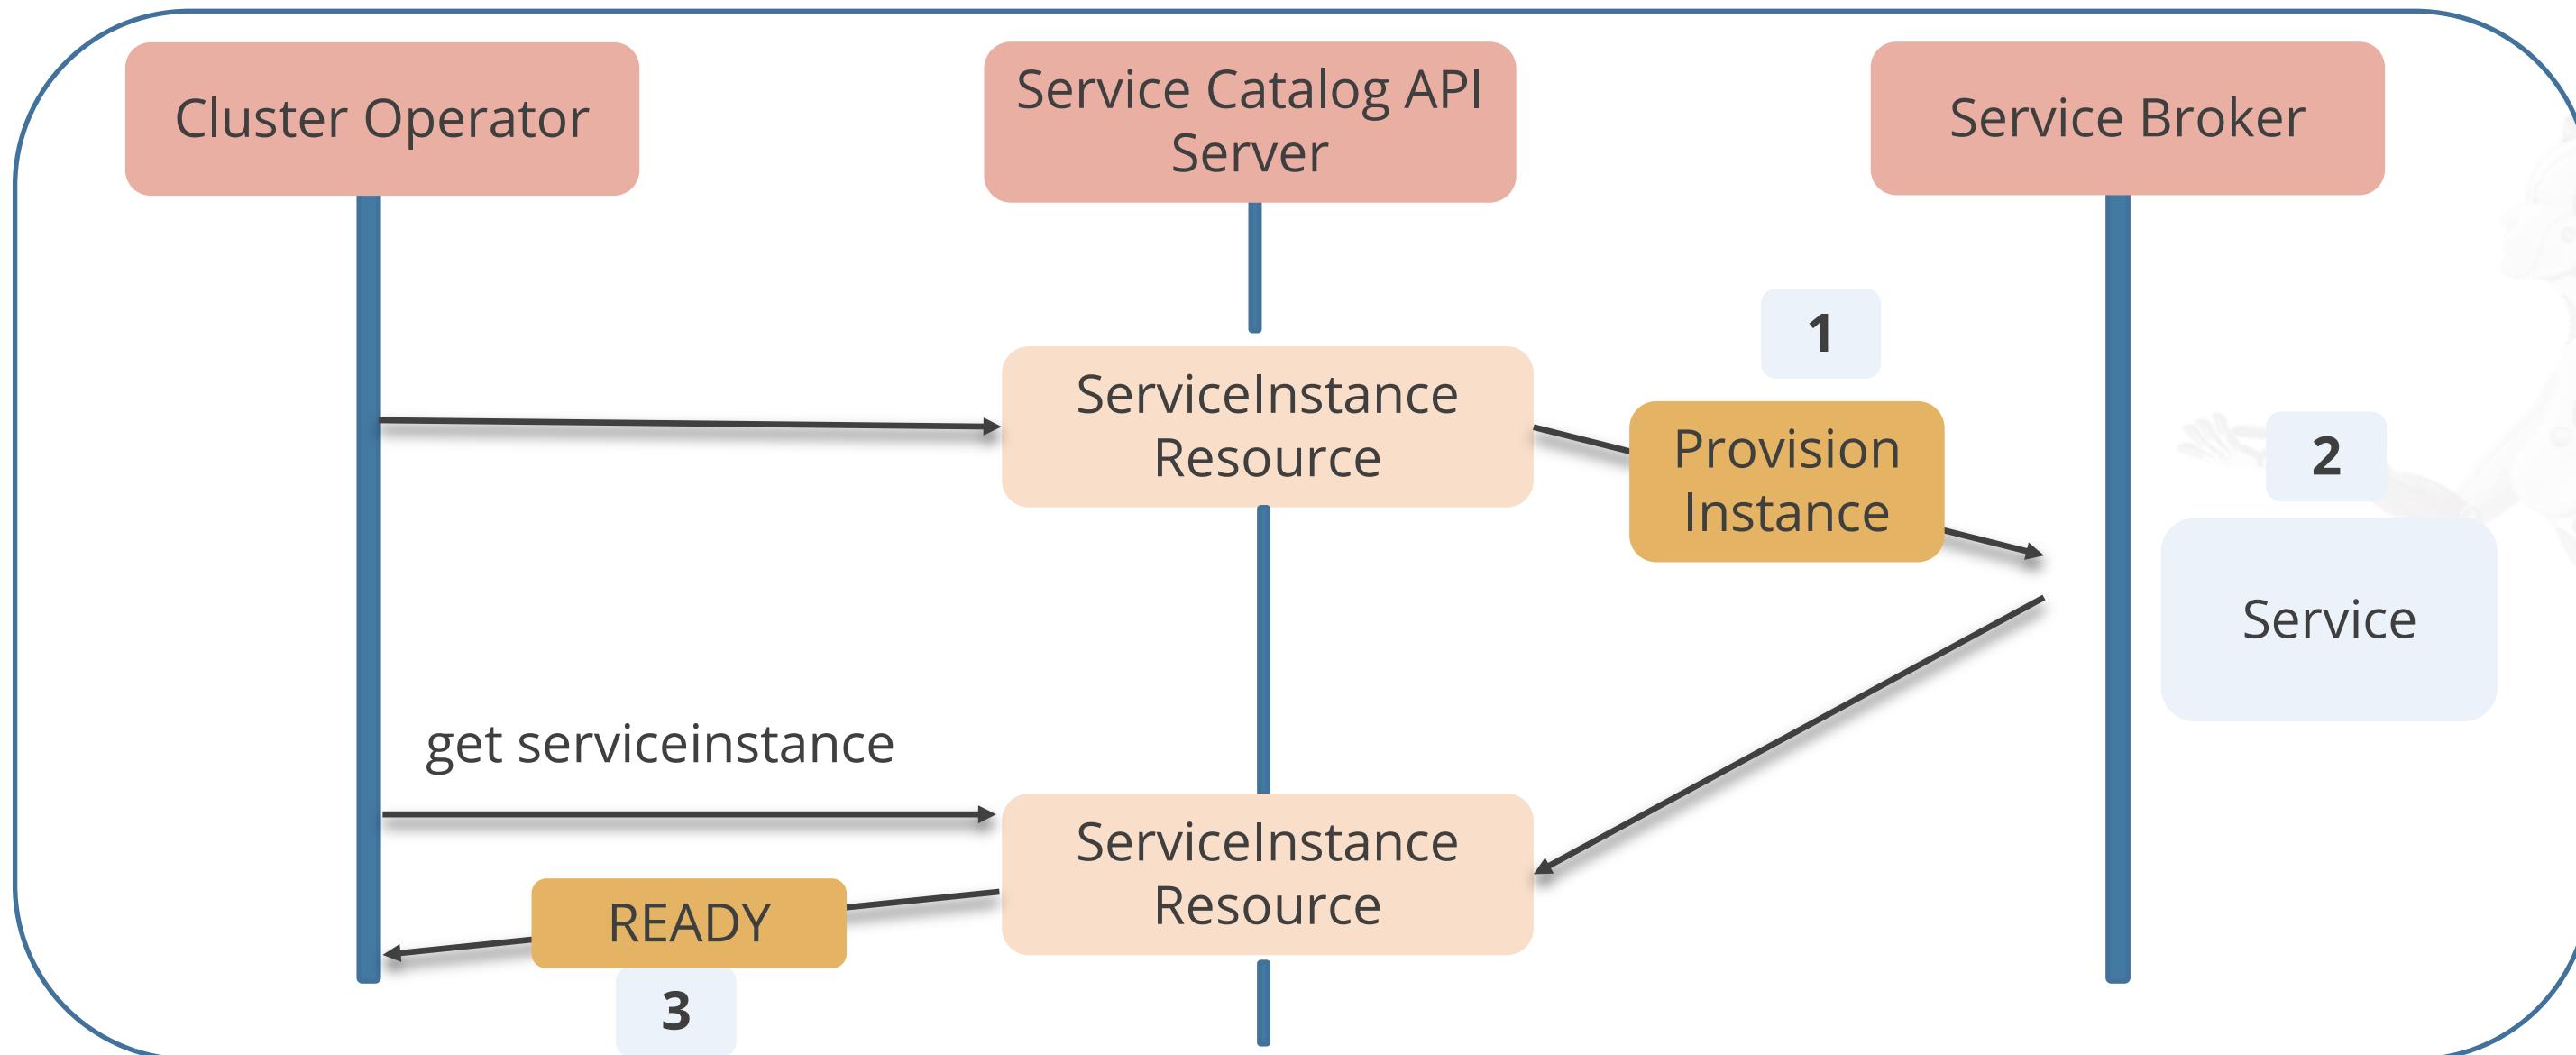
Cluster operator can initiate the provisioning of a new instance by creating a ServiceInstance resource. Given below is an example:

Demo

```
apiVersion: servicecatalog.k8s.io/v1beta1
Kind:   ServiceInstance
Metadata:
  name:  cloud-queue-instance
  namespace: cloud-apps
spec:
  # Reference one of the previously returned services
  clusterServiceClassExternalName: cloud-provider-name
  clusterServicePlanExternalName: service-plan-name
  #####
  # Additional parameters can be added here
  # which may be used by the service broker
  #####
```

# Provision a New Instance

The sequence diagram below illustrates the steps involved in provisioning a new instance of Managed Service.



# Bind to a Managed Service

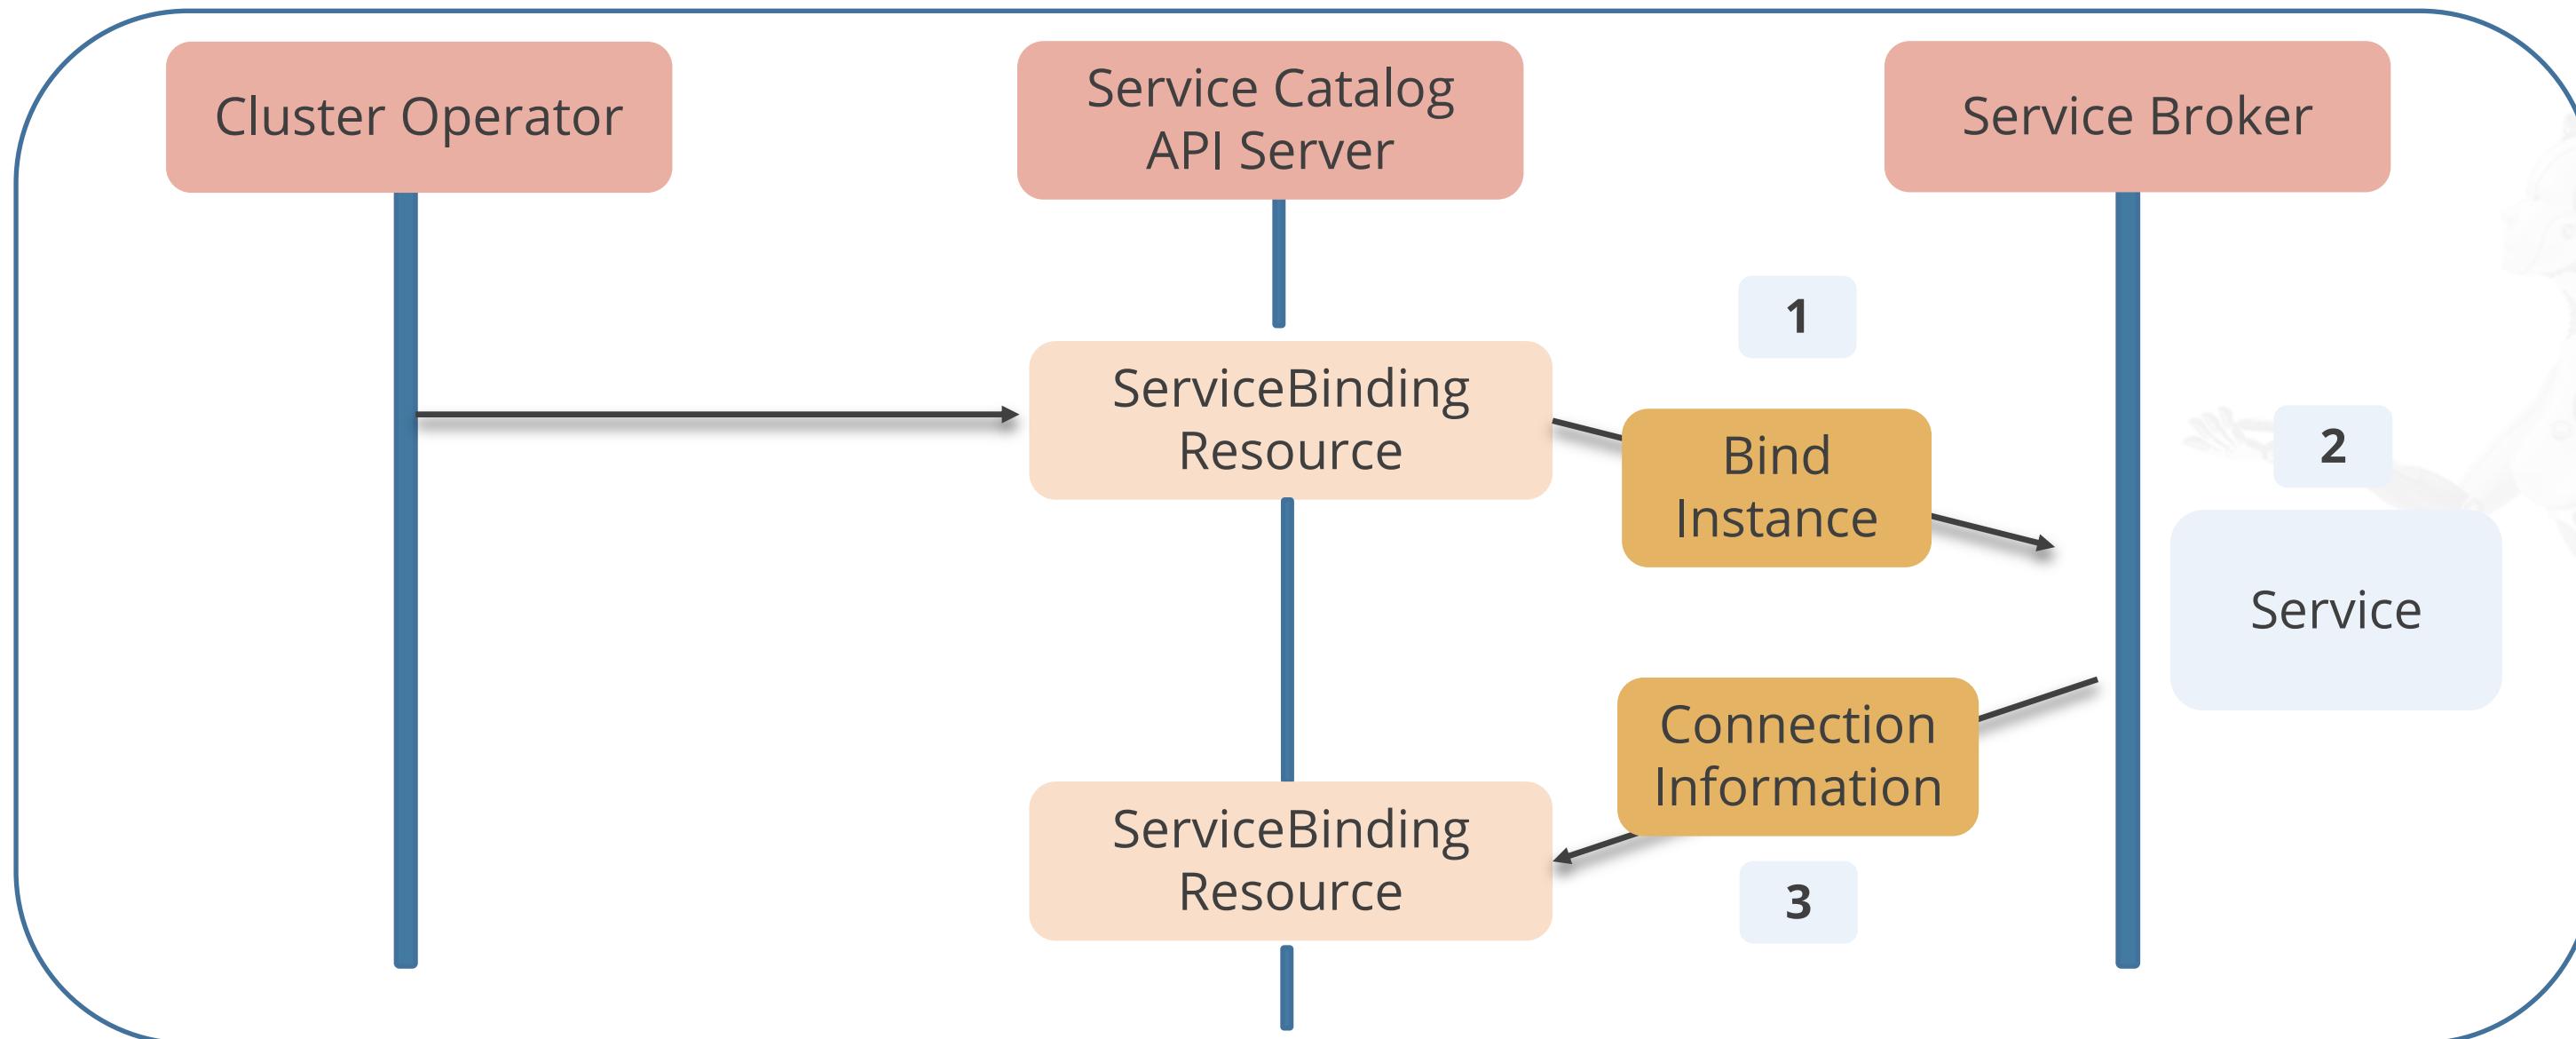
A cluster operator must bind to the Managed Service to get the connection credentials and service account details necessary for the application to use the service. An example of a **ServiceBinding** resource is shown below:

Demo

```
apiVersion: servicecatalog.k8s.io/v1beta1
Kind:   ServiceBinding
Metadata:
  name:  cloud-queue-binding
  namespace: cloud-apps
spec:
  instanceRef:
    name: cloud-provider-name
  #####
  # Additional information can be added here, such as a secretName or
  # service account parameters, which may be used by the service
  broker.
  #####
```

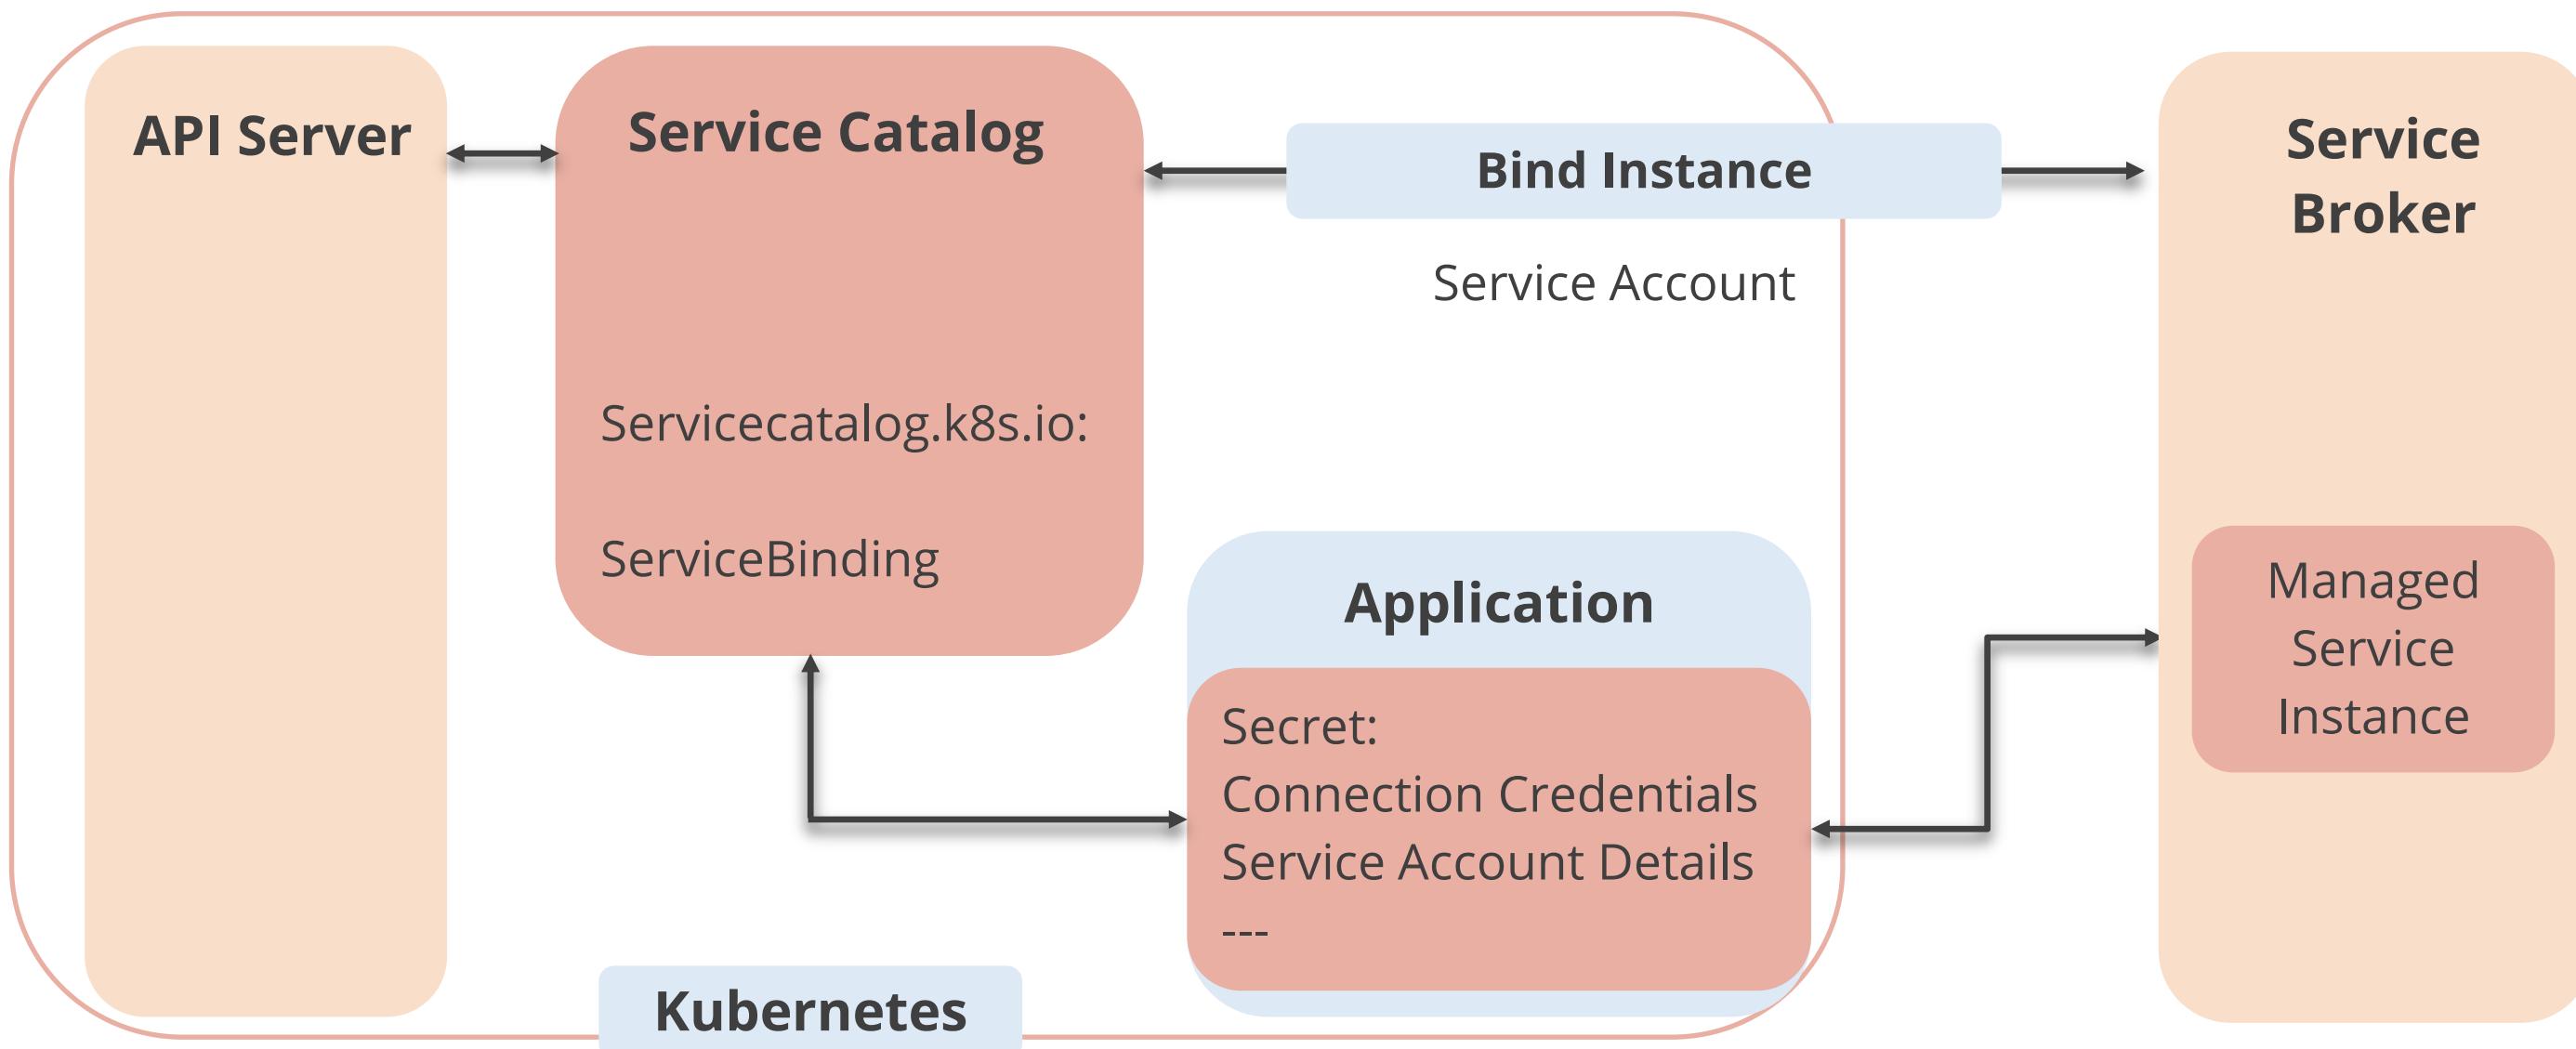
# Bind to a Managed Service

The sequence diagram given below illustrates the steps involved in binding to a Managed Service instance.



# Map the Connection Credentials

The final step involves mapping the connection credentials and service-specific information into the application.



# Pod Configuration File

Declarative Pod configuration can be used to perform mapping. The example here shows how to map service account credentials into the application.

Demo

```
...
spec:
  volumes:
    - name: provider-cloud-key
      secret:
        secretName: sa-key
  containers:
    ...
      volumeMounts:
        - name: provider-cloud-key
          mountpath: /var/secrets/provider
      env:
        - name: PROVIDER_APPLICATION_CREDENTIALS
          value: "/var/secrets/provider/key.json"
```

# Pod Configuration File

The example below describes how to map secret values into application environment variables:

```
Demo
...
env:
- name: "TOPIC"
  valueFrom:
    secretKeyRef:
      name: provider-queue-credentials
      key: topic
"
```

# Understanding Service Catalog



**Problem Statement:** Understand the working of Service Catalog in Kubernetes.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

---

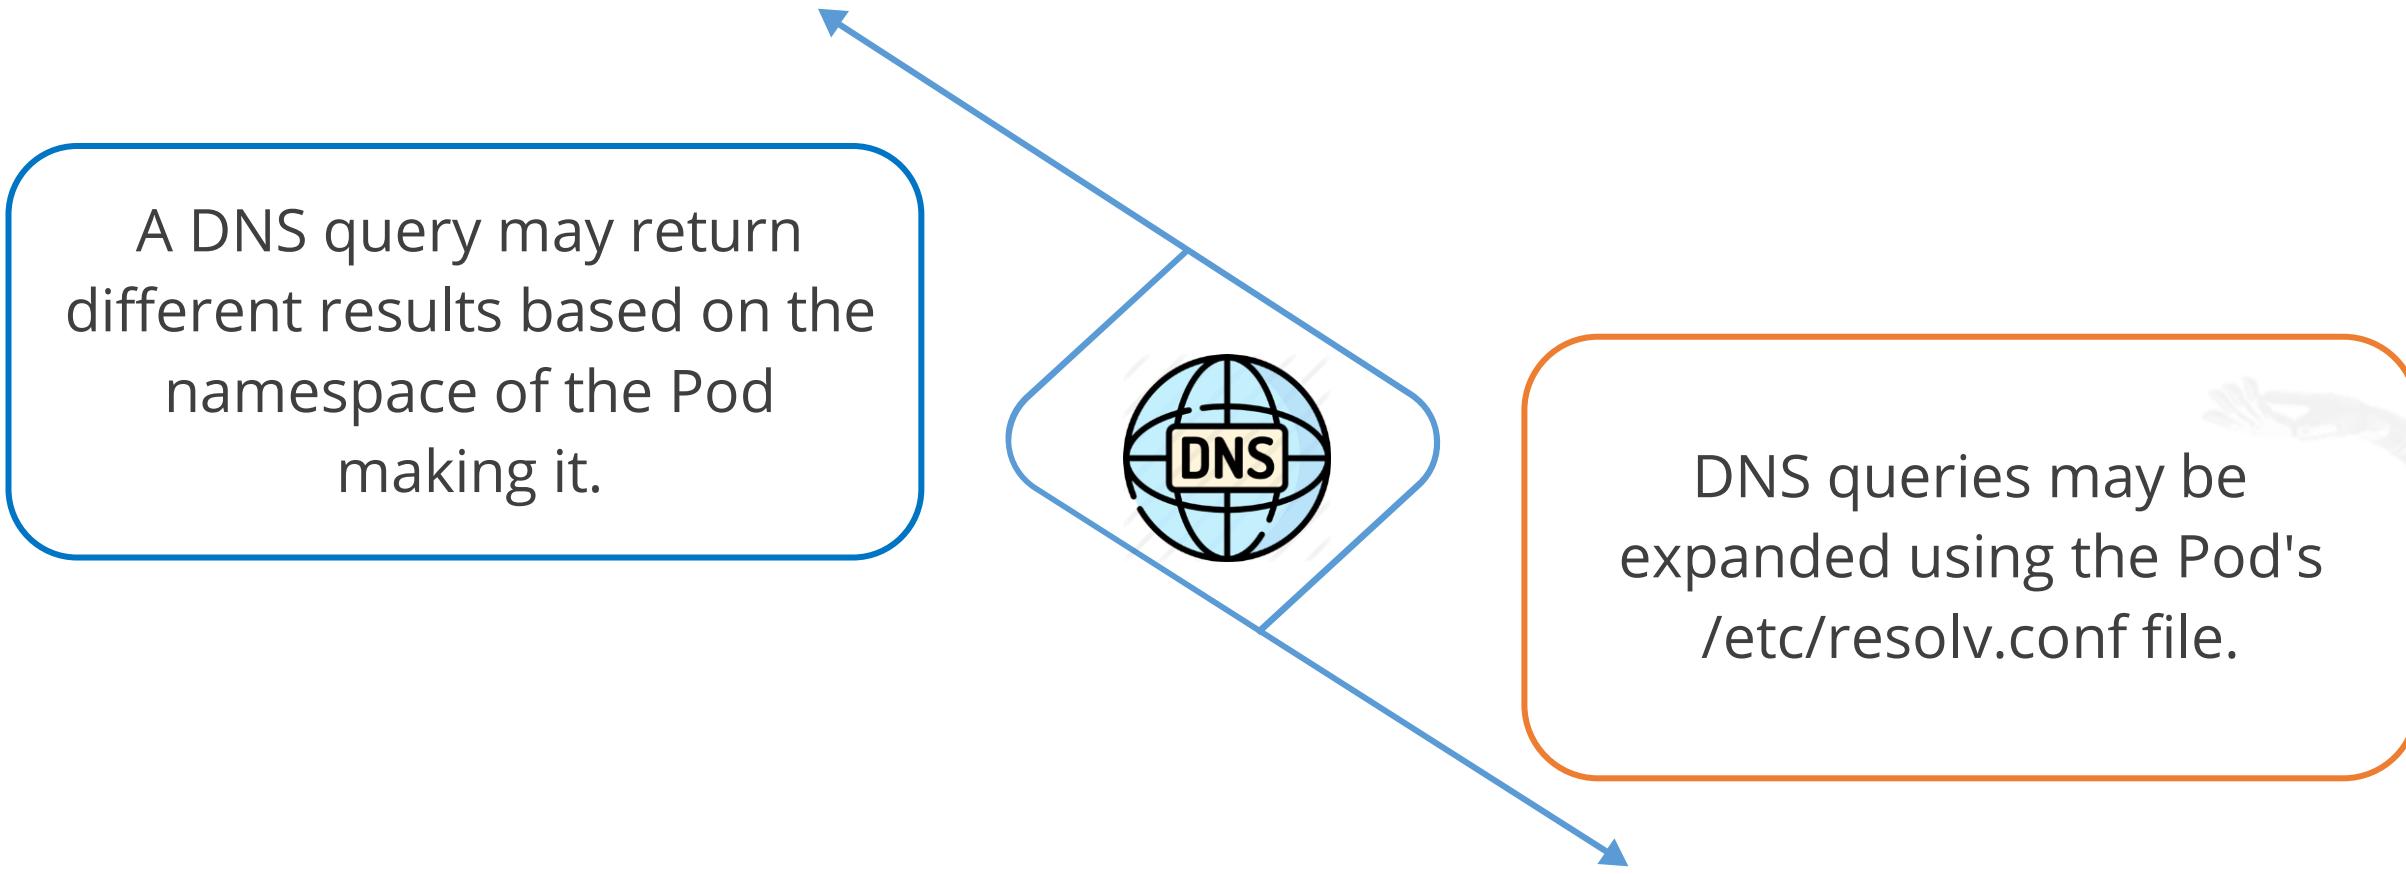
## Steps to demonstrate Service Catalog in Kubernetes:

1. Install service catalog through helm
2. Add service catalog helm repository
3. Check whether it is installed
4. Check whether RBAC is enabled
5. Install service catalog into the cluster

## DNS for Services and Pods

# Introduction

Kubernetes creates DNS records for Services and Pods. Kubernetes DNS schedules a DNS Pod and Service on the cluster and configures the Kubelets to tell individual Containers to use the DNS Service's IP to resolve DNS names.

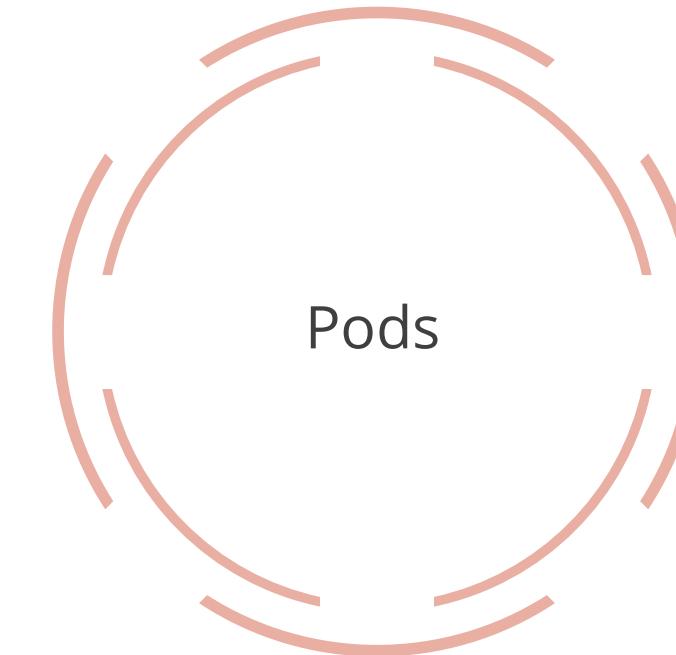
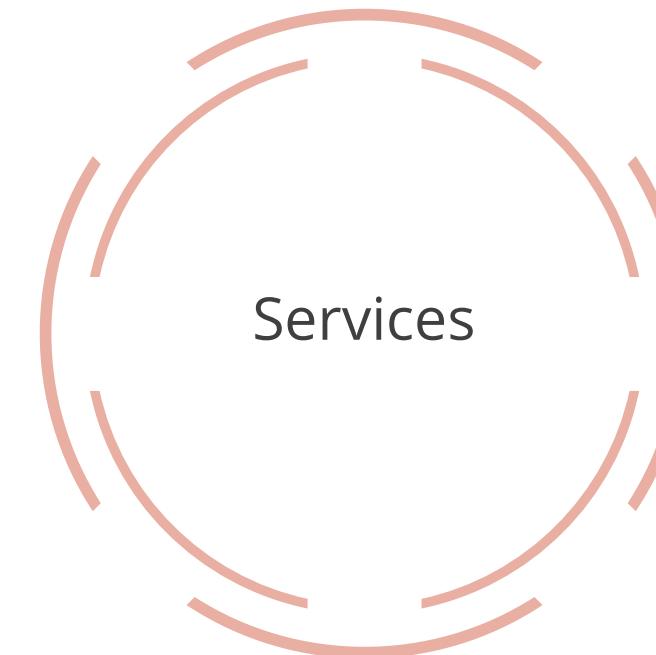


A DNS query may return different results based on the namespace of the Pod making it.

DNS queries may be expanded using the Pod's /etc/resolv.conf file.

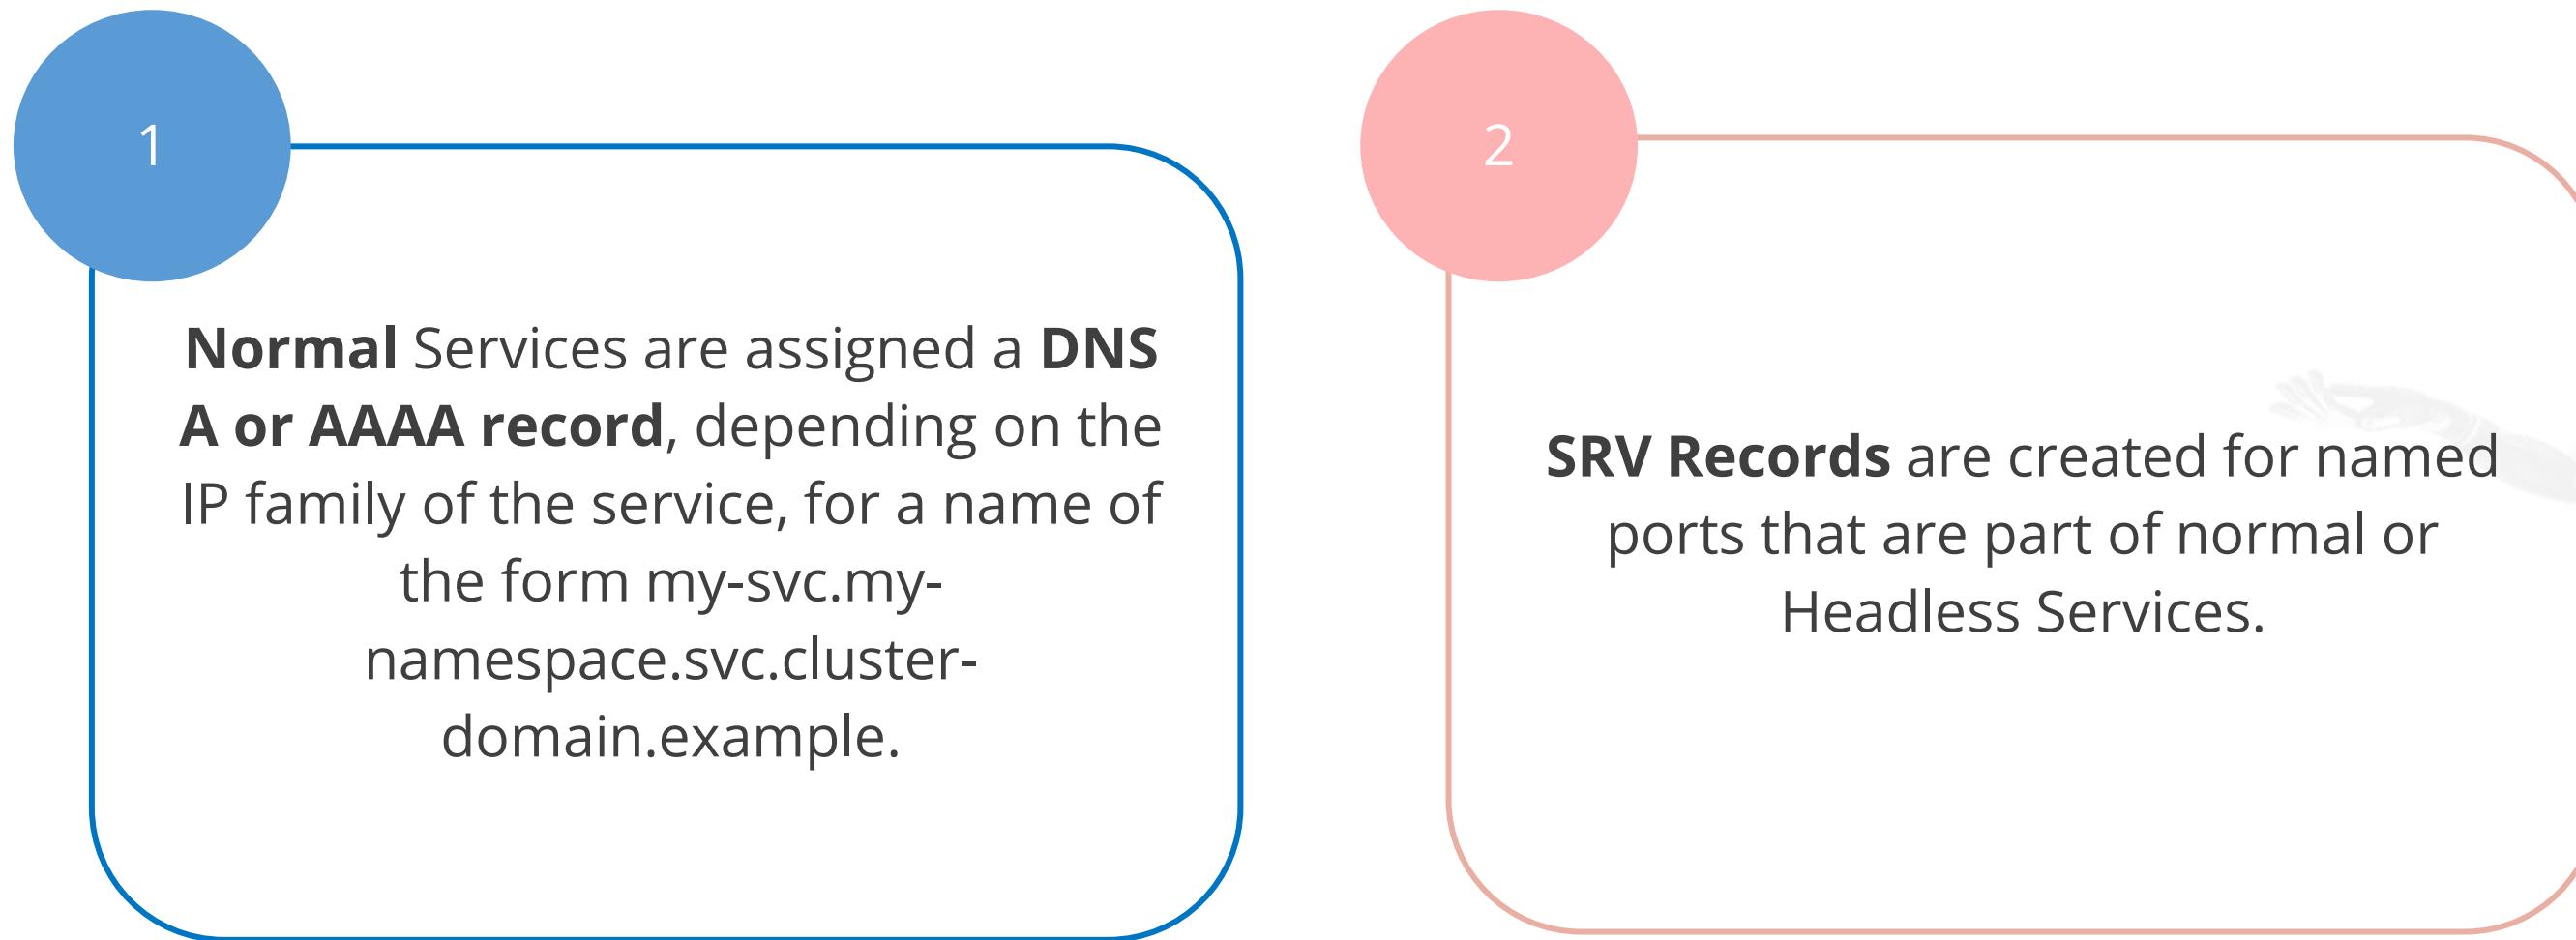
# DNS Records

The following objects get DNS records:



# Services

Snapshots can be provisioned in two ways:



## Pods

A Pod has the following DNS resolution:

Pod-ip-address.my-namespace.Pod.cluster-domain.example

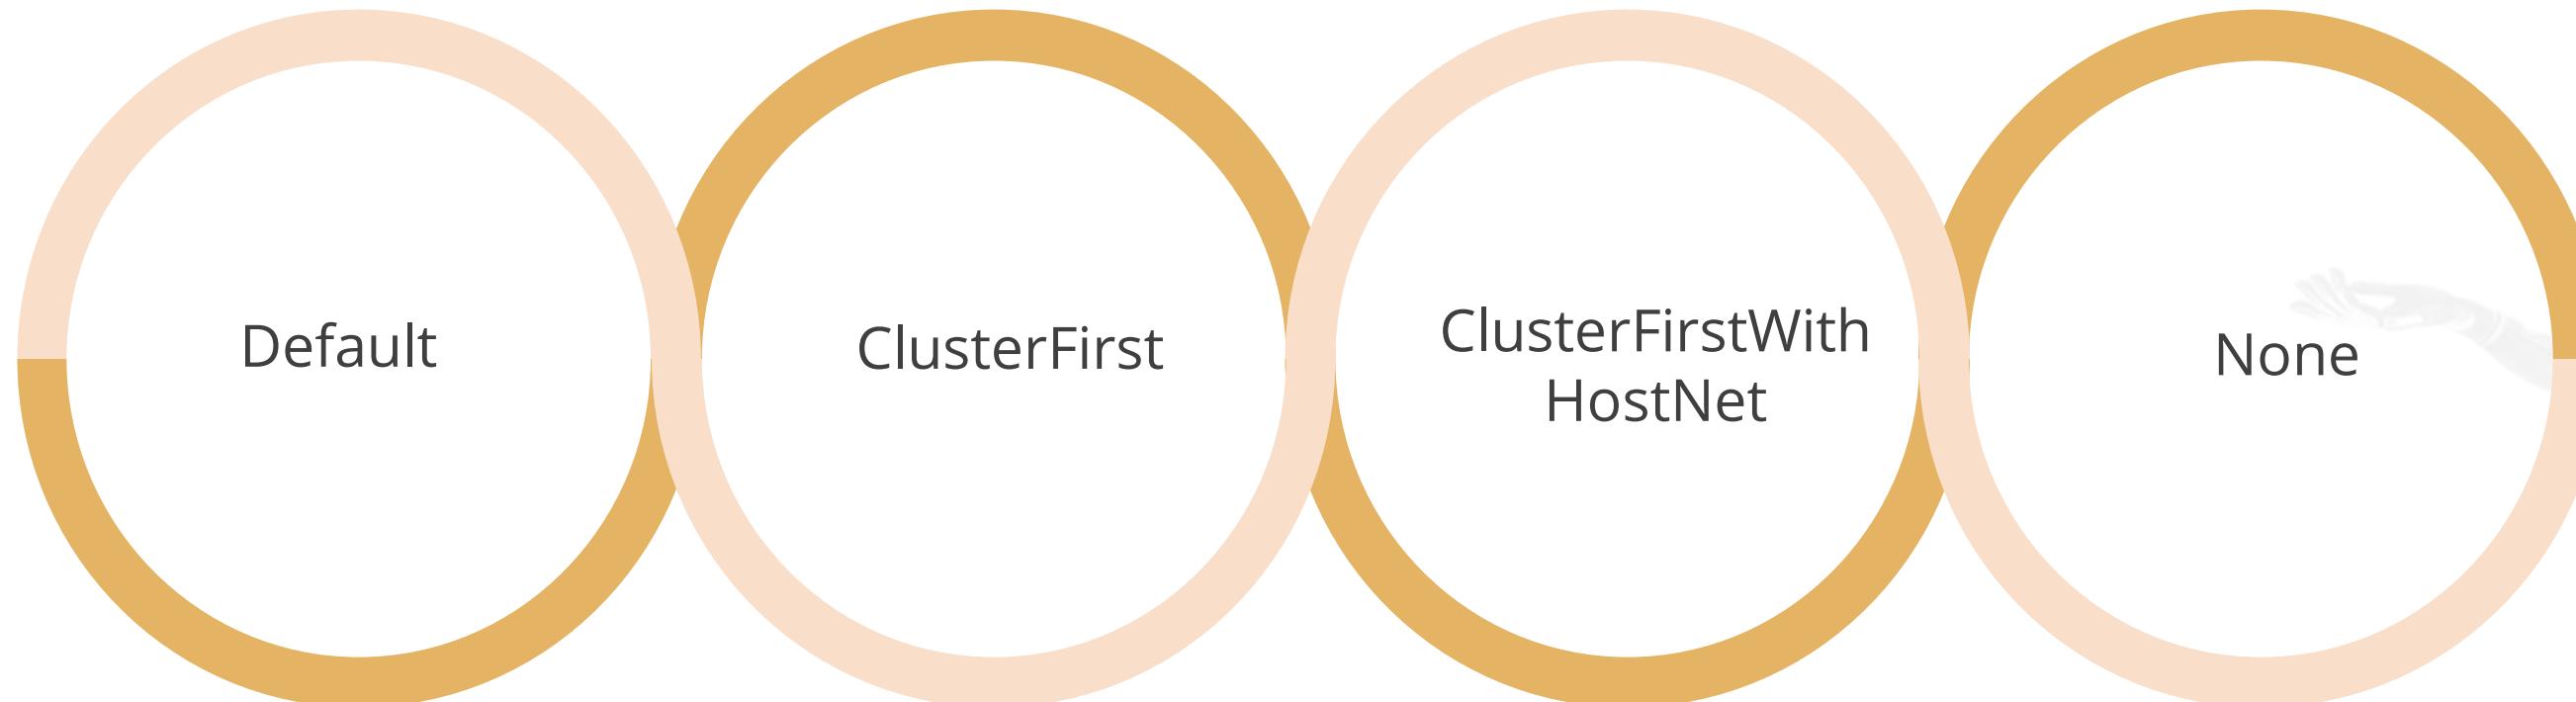
The Pod spec has an optional hostname field, which can be used to specify the Pod's hostname.

It also has an optional subdomain field that can be used to specify its subdomain.

## Pod's DNS Policy

DNS policies can be set on a per-Pod basis.

Kubernetes supports four Pod-Specific DNS Policies:



## Pod's DNS Policy

The example below shows a Pod with its DNS policy set to **ClusterFirstWithHostNet** because it has **hostNetwork** set to **true**.

Demo

```
apiVersion: v1
Kind: Pod
Metadata:
  name: busybox
  namespace: default
spec:
  containers:
  - image: busybox:1.28
    command:
      - sleep:
        - "3600"
    imagePullPolicy: IfNotPresent
    name: busybox
  restartPolicy: Always
  hostNetwork: true
  dnsPolicy: ClusterFirstWithHostNet
```

## Pod's DNS Config

Pod's DNS Config allows users more control on the DNS settings for a Pod.

A user can specify three properties in the dnsConfig field:

**nameservers**

**searches**

**options**

# Pod's DNS Config

Here is an example Pod with custom DNS settings:

Demo

```
apiVersion: v1
Kind: Pod
Metadata:
  namespace: default
  name: dns-example
spec:
  containers:
  - name:
    image: nginx
    dnsPolicy: "None"
    dnsConfig:
      nameservers:
      - 1.2.3.4
      searches:
      - ns1.svc.cluater-domain.example
      - my.dns.search.suffix
      options: busybox
      - name: ndots
        valuek: "2"
      - name: edns0
```

# Understanding the Configuration of DNS for Services and Pods



**Problem Statement:** Learn how to configure DNS for Services and Pods.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

---

## Steps to demonstrate DNS for Services and Pods in Kubernetes:

1. Configure DNS policy
2. Create DNS custom configuration

## Connecting Applications with Services

# Kubernetes Model for Connecting Containers

1

Kubernetes assumes that Pods can communicate with other Pods, regardless of the host on which host they land.

2

Kubernetes gives every Pod its own cluster-private IP address. There is no need to explicitly create links between Pods or map Container ports and host ports.

# Expose Pods to a Cluster

When Pods are exposed to a cluster, it makes them accessible from any Node in the cluster.  
This is done as shown below:

Demo

```
apiVersion: apps/v1
  Kind: Deployment
  Metadata:
    name: my-nginx
  spec:
    selector:
      matchLabels
      run: my-nginx
    replicas: 2
    template:
      metadata:
        labels:
          run: my-nginx
    spec:
      containers:
        - name: my-nginx
          image: nginx
        ports:
          - containerPort: 80
```

# Expose Pods to a Cluster

How to check the Nodes that the Pod is running on:

Demo

Command:

```
kubectl apply -f ./run-my-nginx.yaml  
kubectl get Pods -l run=my-nginx -o wide
```

Result:

| NAME                      | READY | STATUS  | RESTARTS | AGE | IP         | NODE                   |
|---------------------------|-------|---------|----------|-----|------------|------------------------|
| my-nginx-3800858182-jr4a2 | 1/1   | Running | 0        | 13s | 10.244.3.4 | kubernetes-minion-905m |
| my-nginx-3800858182-kna2y | 1/1   | Running | 0        | 13s | 10.244.2.5 | kubernetes-minion-ljyd |

# Create a Service

A Kubernetes Service refers to the abstraction which defines a logical set of Pods running in a cluster with the same functionality. The configuration below helps create a Service for **nginx** replica:

Demo

```
#Create a Service for your 2 nginx replicas with
#kubectl expose:

Command:
kubectl expose deployment/my-nginx

Result:
service/my-nginx exposed
```

Demo

```
#This is equivalent to kubectl apply -f the following
#YAML:
apiVersion: v1
  Kind: Service
  Metadata:
    name: my-nginx
    labels:
      run: my-nginx
  spec:
    ports:
      - port: 80
        protocol: TCP
        selector:
          run: my-nginx
```

# Create a Service

This specification will create a Service which targets TCP port 80 on any Pod with the run my-nginx label. It exposes it on an abstracted Service port.

Demo

```
Command:  
kubectl get svc my-nginx  
  
Result:  
NAME      TYPE      CLUSTER-IP      EXTERNAL-  
IP        PORT(S)    AGE  
my-  
nginx   ClusterIP  10.0.162.149  <none>      80/TC  
P          21s
```

## Create a Service

Service is backed by a group of Pods. These Pods are exposed through **Endpoints**.

The Service's Selector is evaluated continuously and the results are POSTed to an Endpoints object, also named **my-nginx**.

When a Pod dies, it is automatically removed from the Endpoints and new Pods matching the Service's Selector are automatically added to the Endpoints.

# Create a Service

Check the Endpoints. Note that the IPs are the same as the Pods created in the first step.

## Demo

Command:

```
kubectl describe svc my-nginx
```

Result:

```
Name:           my-nginx
Namespace:      default
Labels:         run=my-nginx
Annotations:    <none>
Selector:       run=my-nginx
Type:          ClusterIP
IP:            10.0.162.149
Port:          <unset> 80/TCP
Endpoints:     10.244.2.5:80,10.244.3.4:80
Session Affinity: None
Events:        <none>
```

## Demo

Command:

```
kubectl get ep my-nginx
```

Result:

| NAME     | ENDPOINTS                   | AGE |
|----------|-----------------------------|-----|
| my-nginx | 10.244.2.5:80,10.244.3.4:80 | 1m  |

# Access a Service Using Environment Variables

When a Pod runs on a Node, the kubelet adds a set of environment variables for each active Service.

Demo

#Command:

```
kubectl exec my-nginx-3800858182-jr4a2 -- printenv | grep SERVICE
```

#Result:

```
KUBER    ES_SERVICE_HOST= 10.0.0.1
KUBERNETES_SERVICE_PORT= 443
KUBERNETES_SERVICE_PORT_HTTPS= 443
```

# Access a Service Using DNS

Kubernetes offers a DNS cluster addon Service that automatically assigns DNS names to other Services.

Demo

```
#Command:  
kubectl get services kube-dns --namespace=kube-system  
  
#Result:  
NAME      TYPE        CLUSTER-IP    EXTERNAL-IP   PORT(S)        AGE  
kube-dns   ClusterIP   10.0.0.10   <none>       53/UDP,53/TCP   8m
```

# Secure a Service

These are the requirements of a secure communication channel:

Self-signed  
certificates for  
https

An nginx server  
configured to use  
the certificates

A secret that  
makes the  
certificates  
accessible to  
Pods

# Secure a Service

## nginx https example

Demo

Command:  
`make keys KEY=/tmp/nginx.key  
CERT=/tmp/nginx.crt  
kubectl create secret tls nginxsecret --key /tmp/nginx.key --cert /tmp/nginx.crt`

Result:  
`secret/nginxsecret created`

## configmap example

Demo

Command:  
`kubectl get secrets`

Result:  

| NAME                | TYPE                                | DATA | AGE |
|---------------------|-------------------------------------|------|-----|
| default-token-il9rc | kubernetes.io/service-account-token | 1    | 1d  |
| nginxsecret         | kubernetes.io/tls                   | 2    | 1m  |

Command:  
`kubectl create configmap nginxconfigmap --from-file=default.conf`

Result:  
`configmap/nginxconfigmap created`

Command:  
`kubectl get configmaps`

Result:  

| NAME           | DATA | AGE  |
|----------------|------|------|
| nginxconfigmap | 1    | 114s |

# Expose a Service

Kubernetes supports two ways of exposing a service: **NodePorts** and **LoadBalancers**. The YAML configuration given below creates an **nginx HTTPS** replica to serve the traffic on the internet if the Node has a **public IP**:

## Demo

```
Kubectl get svc my-nginx -o yaml | grep NodePort. -c 5
  uid: 07191fb3-f61a-11e5-42010f00002
spec:
  clusterIP: 10.0.162.149
  ports:
    - name : http
      Nodeport: 31704
      port: 8080
      protocol: TCP
      targetPort: 80
    - name: https
      Nodeport: 32453
      port: 443
      protocol: TCP
      targetPort: 443
  selector:
    run: my-nginx
```

## Demo

```
Kubectl get Nodes -o yaml | grep ExternalIP -c 1
  - address: 104.197.41.11
    type: ExternalIP
    allocable:
  -
  - address : 23.251.152.56
    type: ExternalIP
    allocatable:
  .
  .
$ curl https://EXTENAL-IP:<NODE-PORT> -l
  .
  .
<h1> Welcome to nginx!</h1>
```

# Expose a Service

The IP address in the EXTERNAL-IP column is the one that is available on the public internet.  
The CLUSTER-IP is only available inside the cluster/private cloud network.

```
#Command:  
kubectl edit svc my-nginx  
kubectl get svc my-nginx  
  
#Result:  


| NAME     | TYPE         | CLUSTER-IP   | EXTERNAL-IP    | PORT(S)        | AGE |
|----------|--------------|--------------|----------------|----------------|-----|
| my-nginx | LoadBalancer | 10.0.162.149 | xx.xxx.xxx.xxx | 8080:30163/TCP | 21  |

  
curl https://<EXTERNAL-IP> -k  
...  
<title>Welcome to nginx!</title>
```

# Connecting Applications with Services



**Problem Statement:** Learn how to connect and configure applications with Services in Kubernetes.

## Assisted Practice: Guidelines

---

### **Steps to demonstrate connecting applications with Services in Kubernetes:**

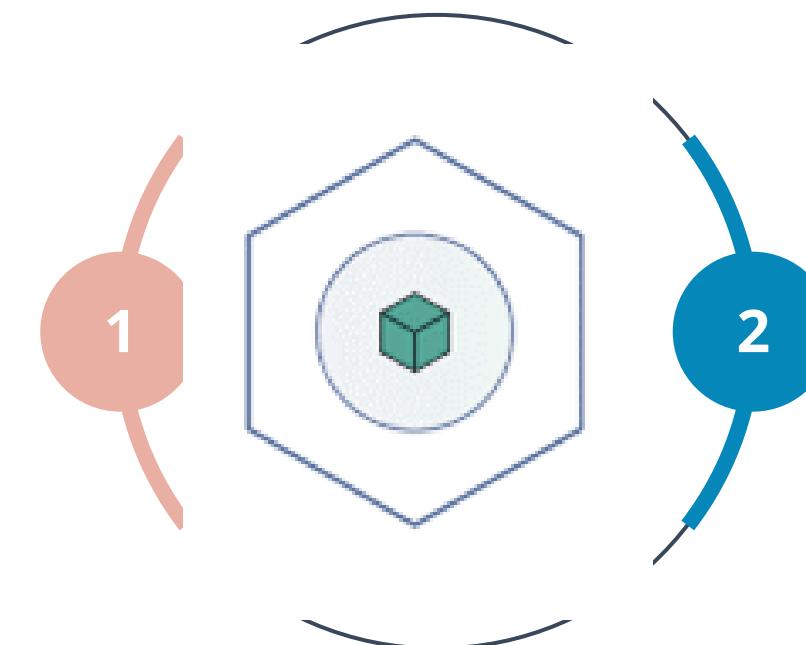
1. Expose Pod to cluster
2. Create Pod
3. Check the Nodes on which the Pod is running
4. Create a Service
5. Check the Service
6. Fetch environment variables

## EndpointSlices

# Introduction

EndpointSlices provide a simple way to track network Endpoints within a Kubernetes cluster.

EndpointSlices help to mitigate issues and provide an extensible platform for additional features such as Topological Routing.



They offer a more scalable and extensible alternative to Endpoints.

# EndpointSlice Resources

An EndpointSlice contains references to a set of network Endpoints. Here is a sample EndpointSlice resource:

Demo

```
apiVersion: discovery.k8s.io/v1
Kind:   EndpointSlice
Metadata:
  name:  example-abc
  labels:
    kubernetes>io/servicename: example
  addressType: IPv4
  ports:
    - name: http
      protocol: TCP
      port: 80
  endpoints :
    - addressess:
        -"10.1.2.3"
      conditions:
        ready: true
      host name : Pod-1
      NodeName: Node-1
      zone: us-west2-a
```

# Address Types and Conditions

EndpointSlices support three address types, namely, IPv4, IPv6, and Fully Qualified Domain Name (FQDN).

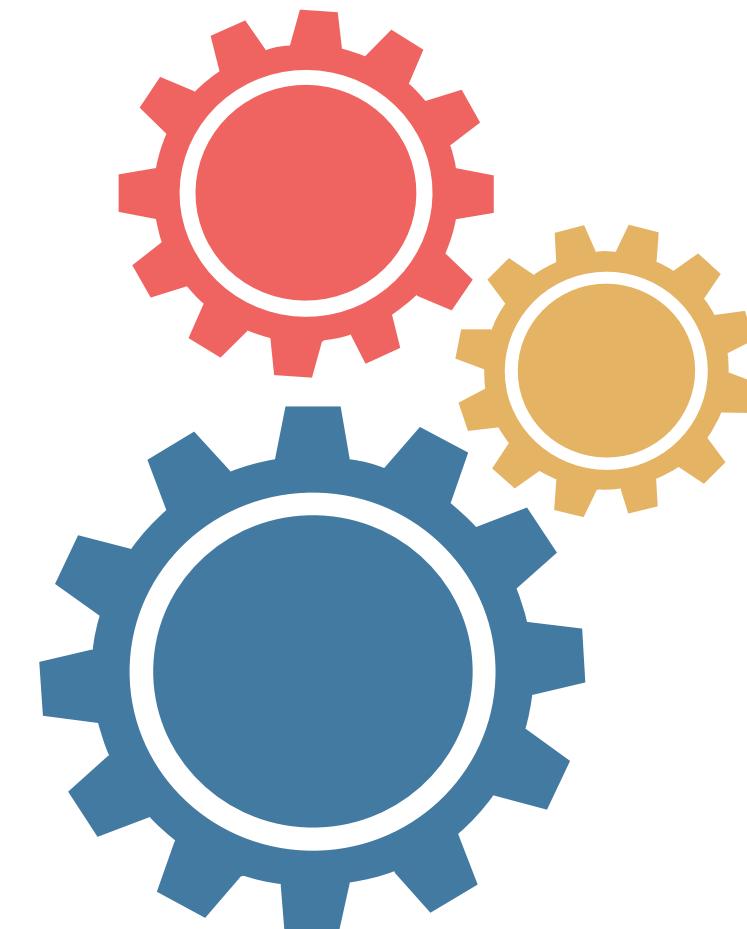
The EndpointSlice API stores three Endpoint conditions:

## Ready

A condition that maps to a Pod's Ready condition

## Terminating

A condition that indicates whether an Endpoint is terminating



## Serving

A condition that does not account for terminating states

# Topology Information, Management, and Ownership

Topology information includes the location of the Endpoint and information about the corresponding Node and zone.

EndpointSlices are owned by the Service that the EndpointSlice object tracks Endpoints for.

Most often, the Control Plane (specifically the EndpointSlice Controller) creates and manages EndpointSlice objects.



# EndpointSlice Mirroring

When applications create custom Endpoint resources, the cluster's Control Plane mirrors most Endpoint resources to corresponding EndpointSlices. This helps to avoid concurrently writing to both Endpoints and EndpointSlice resources.

## The Control Plane mirrors Endpoints resources unless:



The Endpoints resource has an `Endpointslice.kubernetes.io/skip-mirror` label set to true.



The Endpoints resource has a `control-plane.alpha.kubernetes.io/leader` annotation.



The corresponding Service resource does not exist.



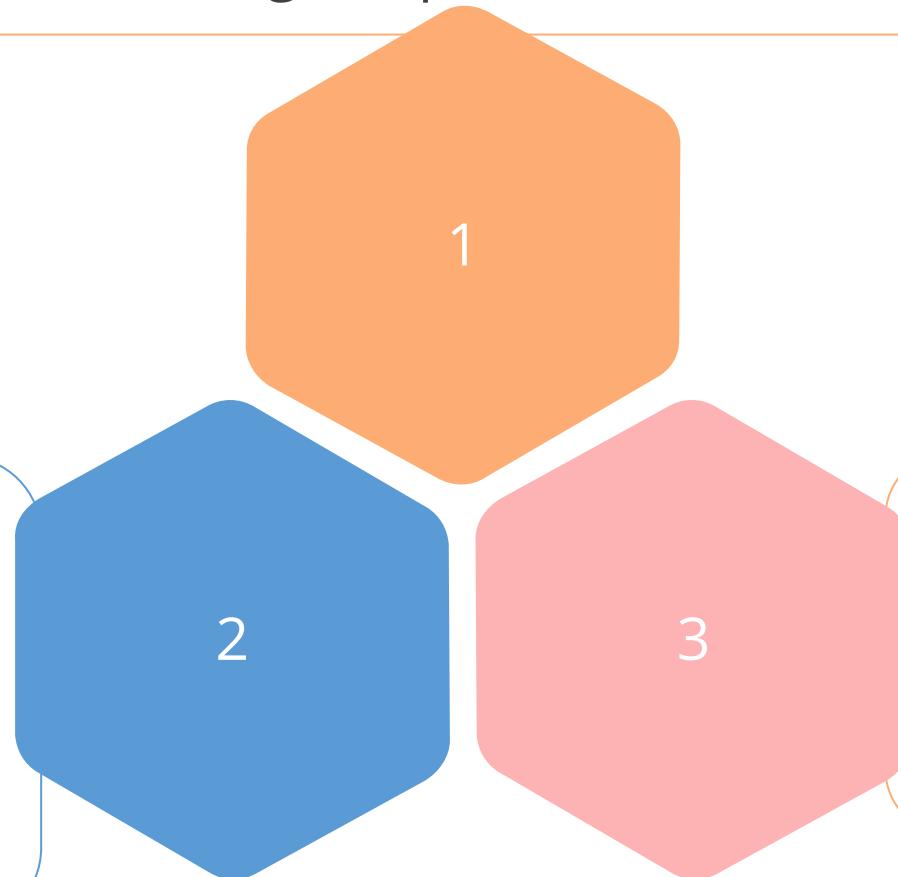
The corresponding Service resource has a non-nil Selector.

# Distribution of EndpointSlices

To fill the EndpointSlices, the Control Plane does the following:

Iterate through existing EndpointSlices, remove Endpoints that are no longer desired, and update matching Endpoints that have changed

If there are still new Endpoints left to be added, try to fit them into a previously unchanged slice and/or create new ones



Iterate through EndpointSlices that have been modified in the first step and fill them up with any new Endpoints needed

# Duplicate Endpoints

Due to the nature of EndpointSlice changes, Endpoints may be represented in more than one EndpointSlice at the same time.

This naturally occurs as changes to different EndpointSlice objects can arrive at the Kubernetes client watch or cache at different times.

Implementations using EndpointSlice must be able to have the Endpoint appear in more than one slice.

# Understanding the Configuration of EndpointSlices



**Problem Statement:** Learn how to configure Endpointslices in Kubernetes.

# Assisted Practice: Guidelines

---

## Steps to demonstrate EndpointSlices in Kubernetes:

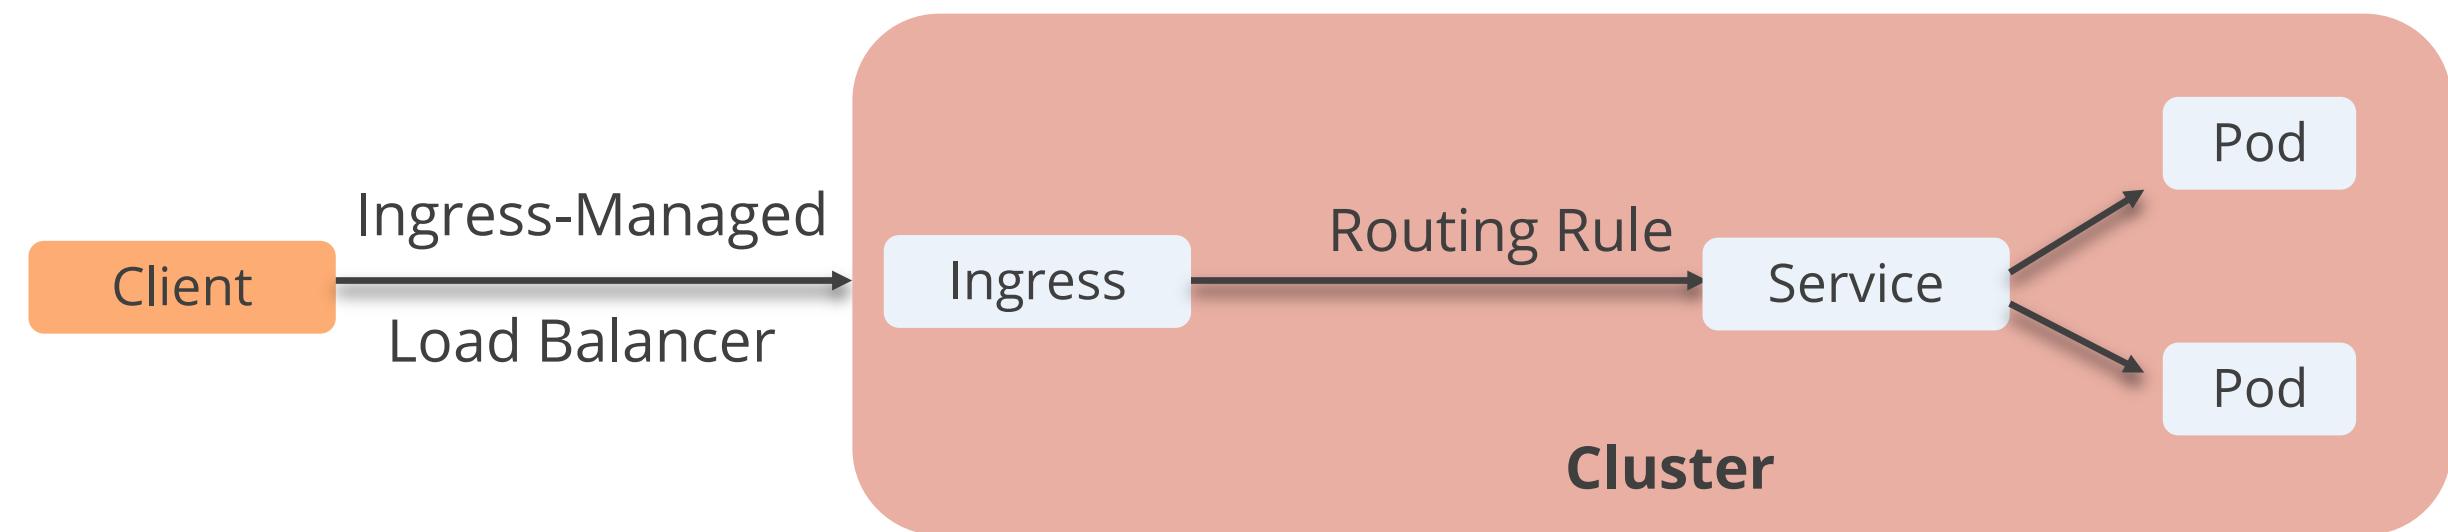
1. Create and configure a simple Endpointslice
2. Check the resources
3. Describe EndpointSlices

## Ingress

# What Is Ingress?

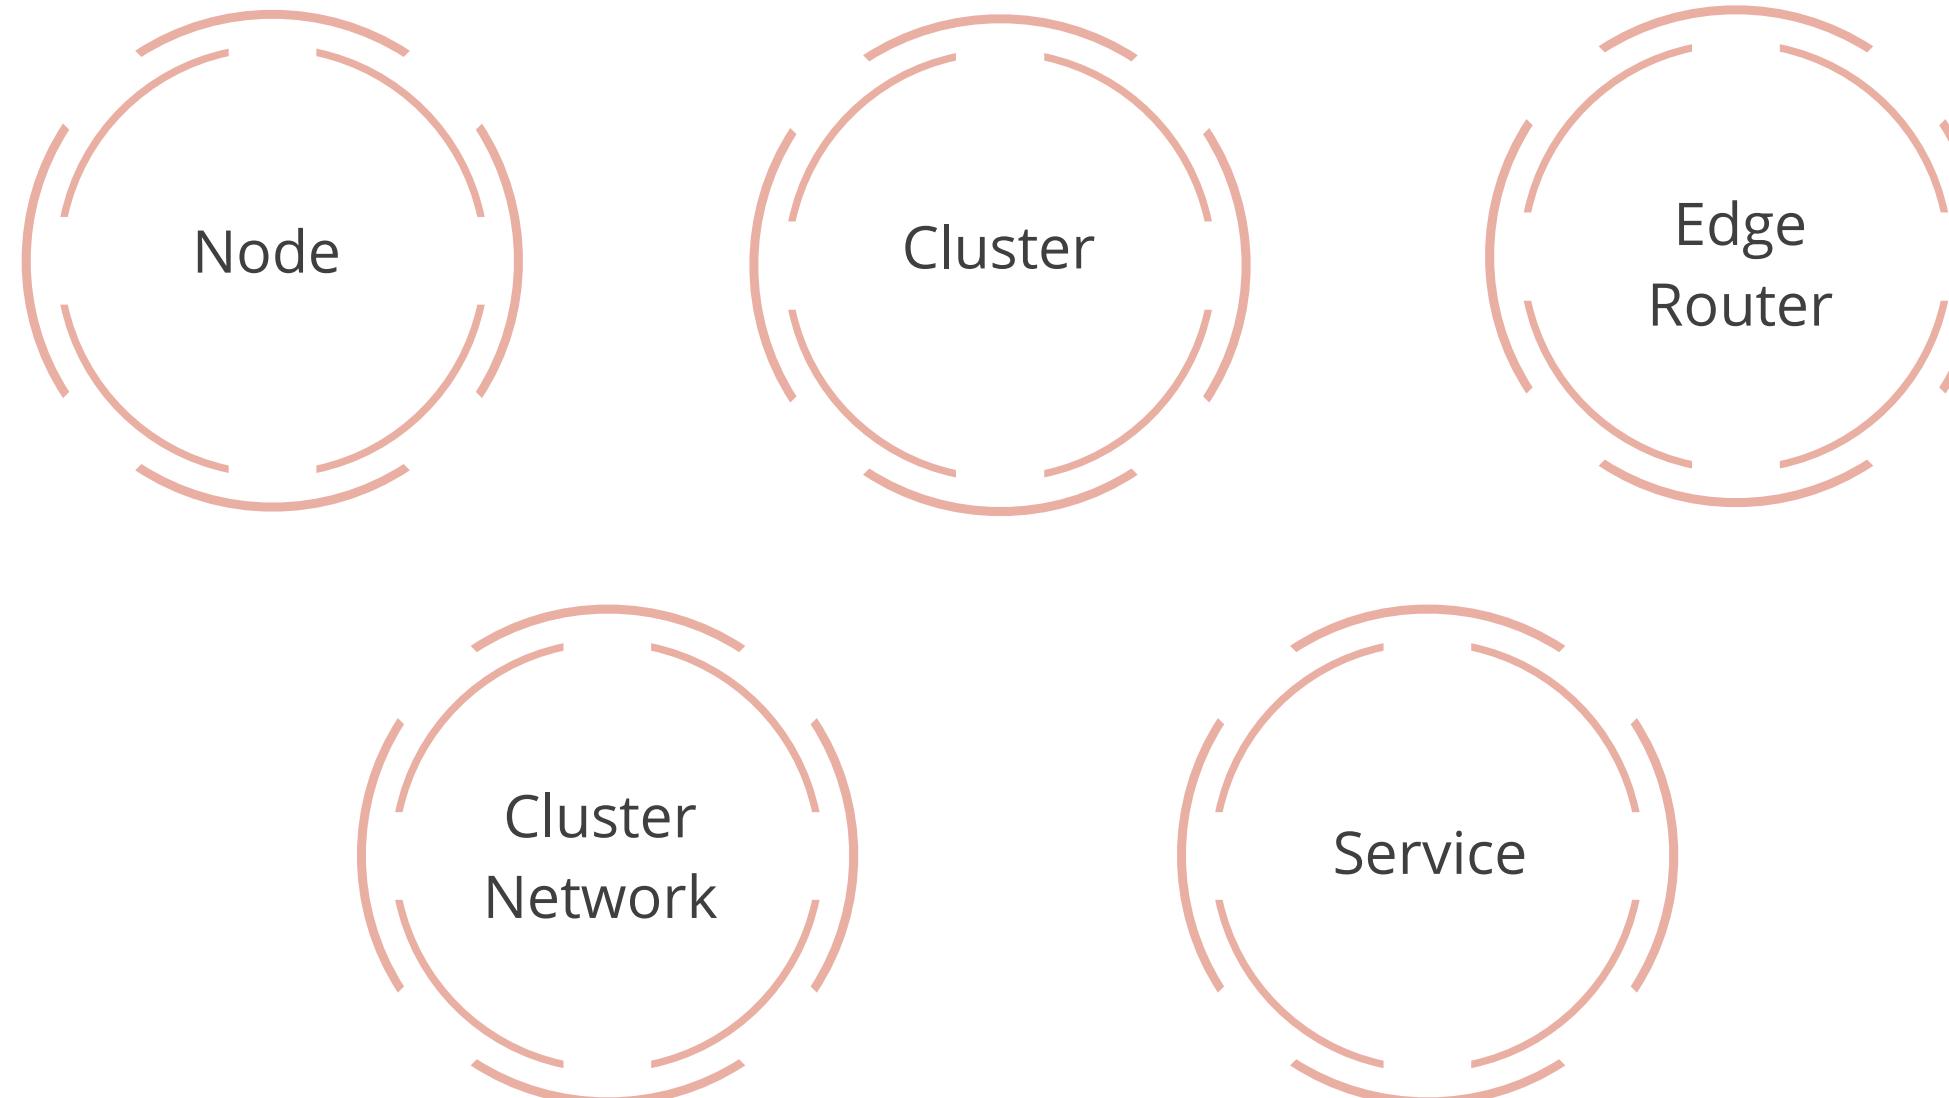
Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.

Here is a simple example where an Ingress sends all its traffic to one Service:



An Ingress may be configured to give Services externally reachable URLs and Load Balance traffic, terminate SSL/TLS, and offer name-based virtual hosting.

# Terminologies Used in Ingress



# Ingress Resource

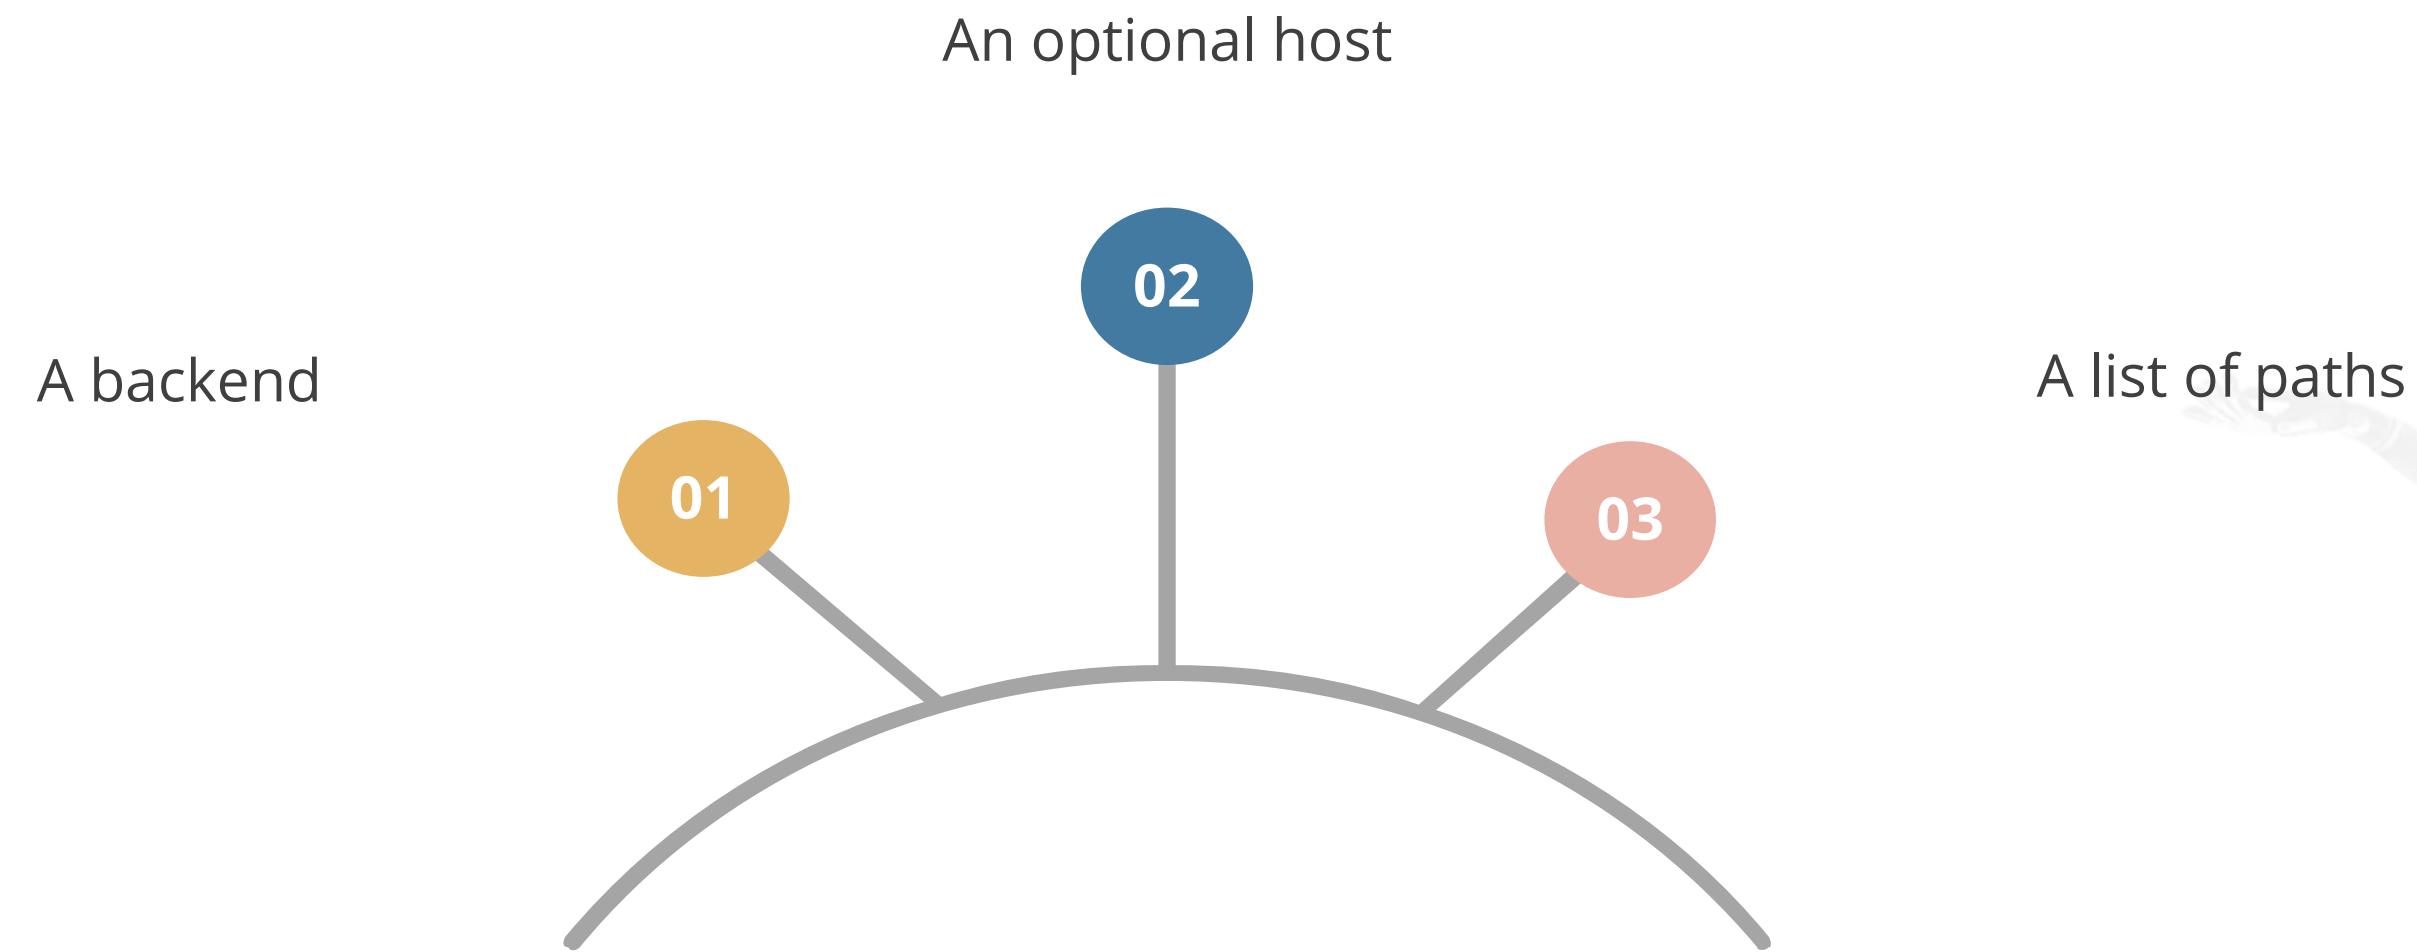
An Ingress needs apiVersion, kind, and metadata fields.  
Shown below is an example of a minimal Ingress resource:

## Demo

```
apiVersion: networking.k8s.io/v1
Kind:   Ingress
Metadata:
  name:  minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          service:
            name: test
            port : Pod-1
            number: 80
```

# Ingress Rules

Each HTTP rule contains the following information:



# Default and Resource Backend

The default backend is conventionally a configuration option of the Ingress Controller and is not specified in the Ingress resources.

Default backend

A Resource backend is an ObjectRef to another Kubernetes resource within the same namespace as the Ingress object.

Resource backend

# Resource Backend

A common usage for a Resource backend is to ingress data to an object storage backend with static assets.

## Demo

```
apiVersion: networking.k8s.io/v1
Kind: Ingress
Metadata:
  name: ingress-resource-backend
  defaultBackend:
    resource: IPv4
    apiGroup:
    kind:
    name: static-assets
  - http
    paths:
    - path: /icons
      pathType: ImplementationSpecific
      backend:
        resource:
          apigroup: k8s.example.com
          kind : StorageBucket
          Name: icon-assets
```

## Command to View Ingress

To view the created Ingress, use the command shown below:

Demo

```
kubectl describe ingress ingress-resource-backend
```

# Path Types

Each path in an Ingress must have a corresponding path type. There are three supported path types:

## ImplementationSpecific

Is dependent on the **IngressClass** for matching

## Exact

Matches the **URL path exactly**, with case sensitivity

## Prefix

Matches based on a **URL path prefix** split by the / separator

# Hostname Wildcards

Precise matches require that the HTTP host header matches the host field. Wildcard matches require the HTTP host header to be equal to the suffix of the wildcard rule.

Hosts can be precise matches (for example **foo.bar.com**) or a wildcard (for example **\*.foo.com**).

| Host      | Host header     | Match   |
|-----------|-----------------|---|
| *.foo.com | bar.foo.com     | Matches based on shared suffix                    |
| *.foo.com | baz.bar.foo.com | No match, wildcard only covers a single DNS label |
| *.foo.com | foo.com         | No match, wildcard only covers a single DNS label |

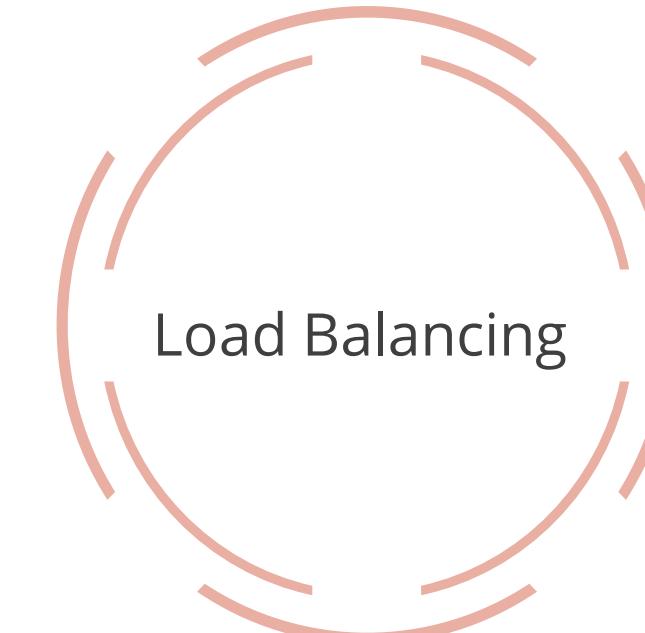
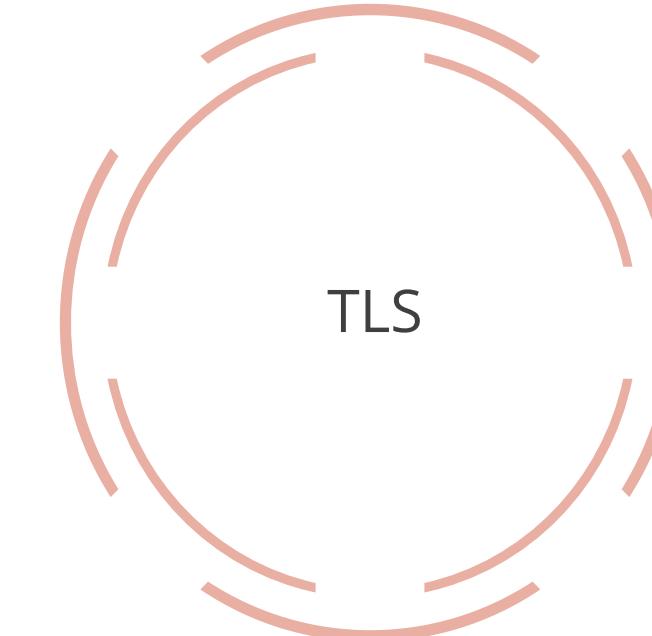
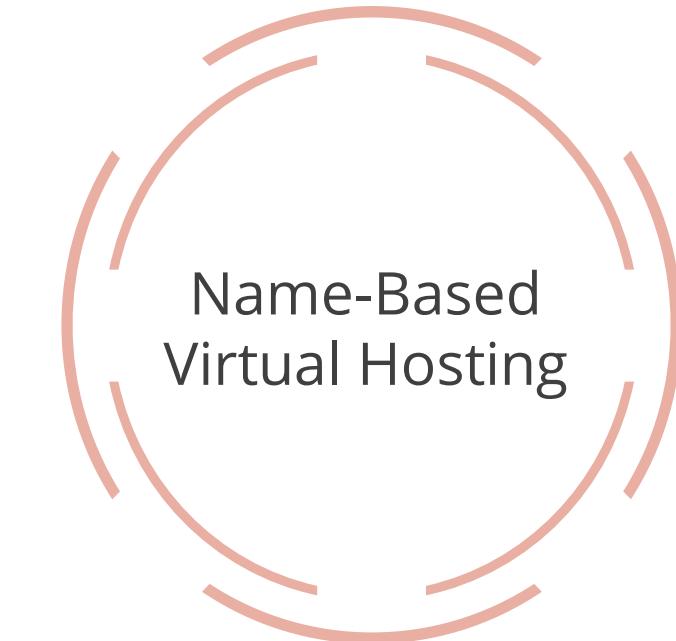
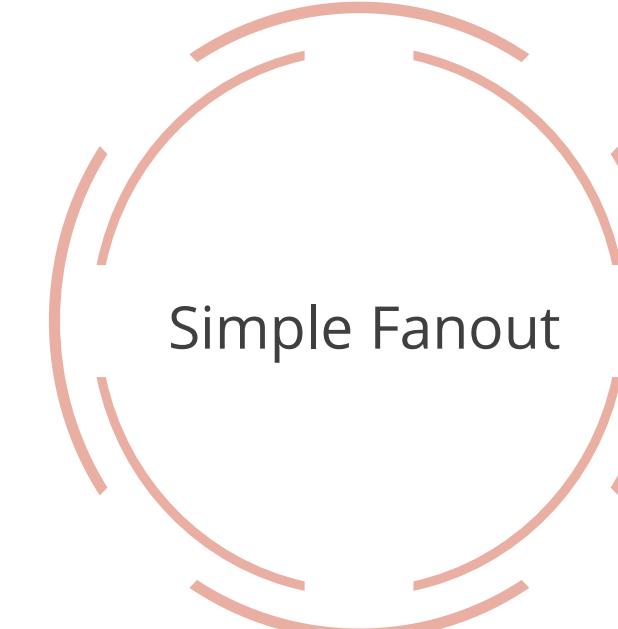
# Hostname Wildcards

## Demo

```
apiVersion: networking.k8s.io/v1
Kind: Ingress
Metadata:
  name: ingress-wildcard-host
spec:
  rules:
    - host: "foo.bar.com"
      http:
        paths:
          - pathType: Prefix
            path: "/bar"
            backend:
              service:
                Name: service1
                port:
                  number: 80
    - host: "*foo.com"
      http:
        paths:
          - pathType: Prefix
            path: "/foo"
            backend:
              service:
                Name: service2
                port:
                  number: 80
```

An example of the configuration file showing the use of wildcards

# Types of Ingress



# Ingress Backed by a Single Service

A single service can be exposed using an Ingress by specifying a default backend with no rules.

Demo

```
apiVersion: networking.k8s.io/v1
Kind:   Ingress
Metadata:
  name:  test-ingress
spec:
  defaultBackend
  service:
    name: test
  port:
    number: 80
```

Demo

```
#Command view the state of the created Ingress:
kubectl get ingress test-ingress
```

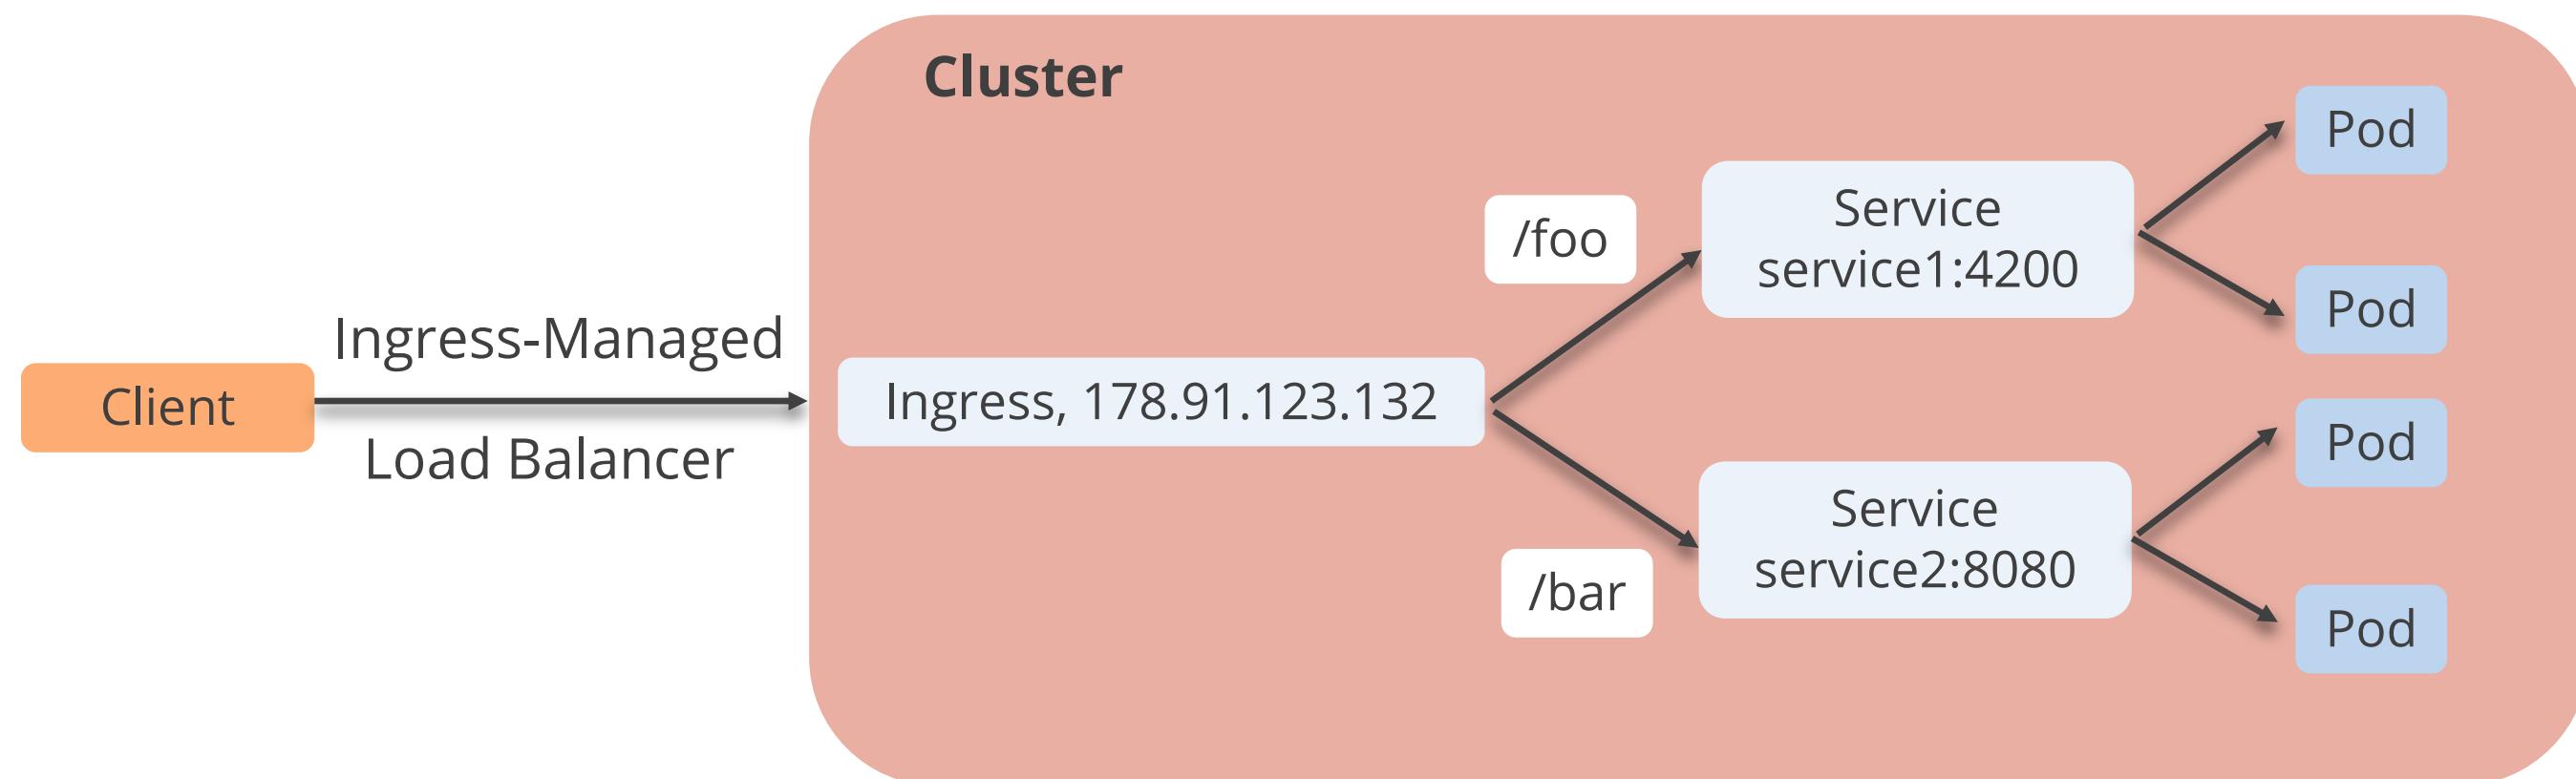
Result:

| NAME         | CLASS        | HOSTS | ADDRESS       |
|--------------|--------------|-------|---------------|
|              | PORTS        | AGE   |               |
| test-ingress | external- lb | *     | 203.0.113.123 |
|              |              |       | 80            |
|              |              |       | 59s           |

```
Where 203.0.113.123 is the IP allocated by the
Ingress controller to satisfy this Ingress.
```

# Simple Fanout

A Fanout configuration routes traffic from a single IP address to more than one Service, based on the HTTP URI being requested.



# Simple Fanout

It would require an Ingress such as:

## Demo

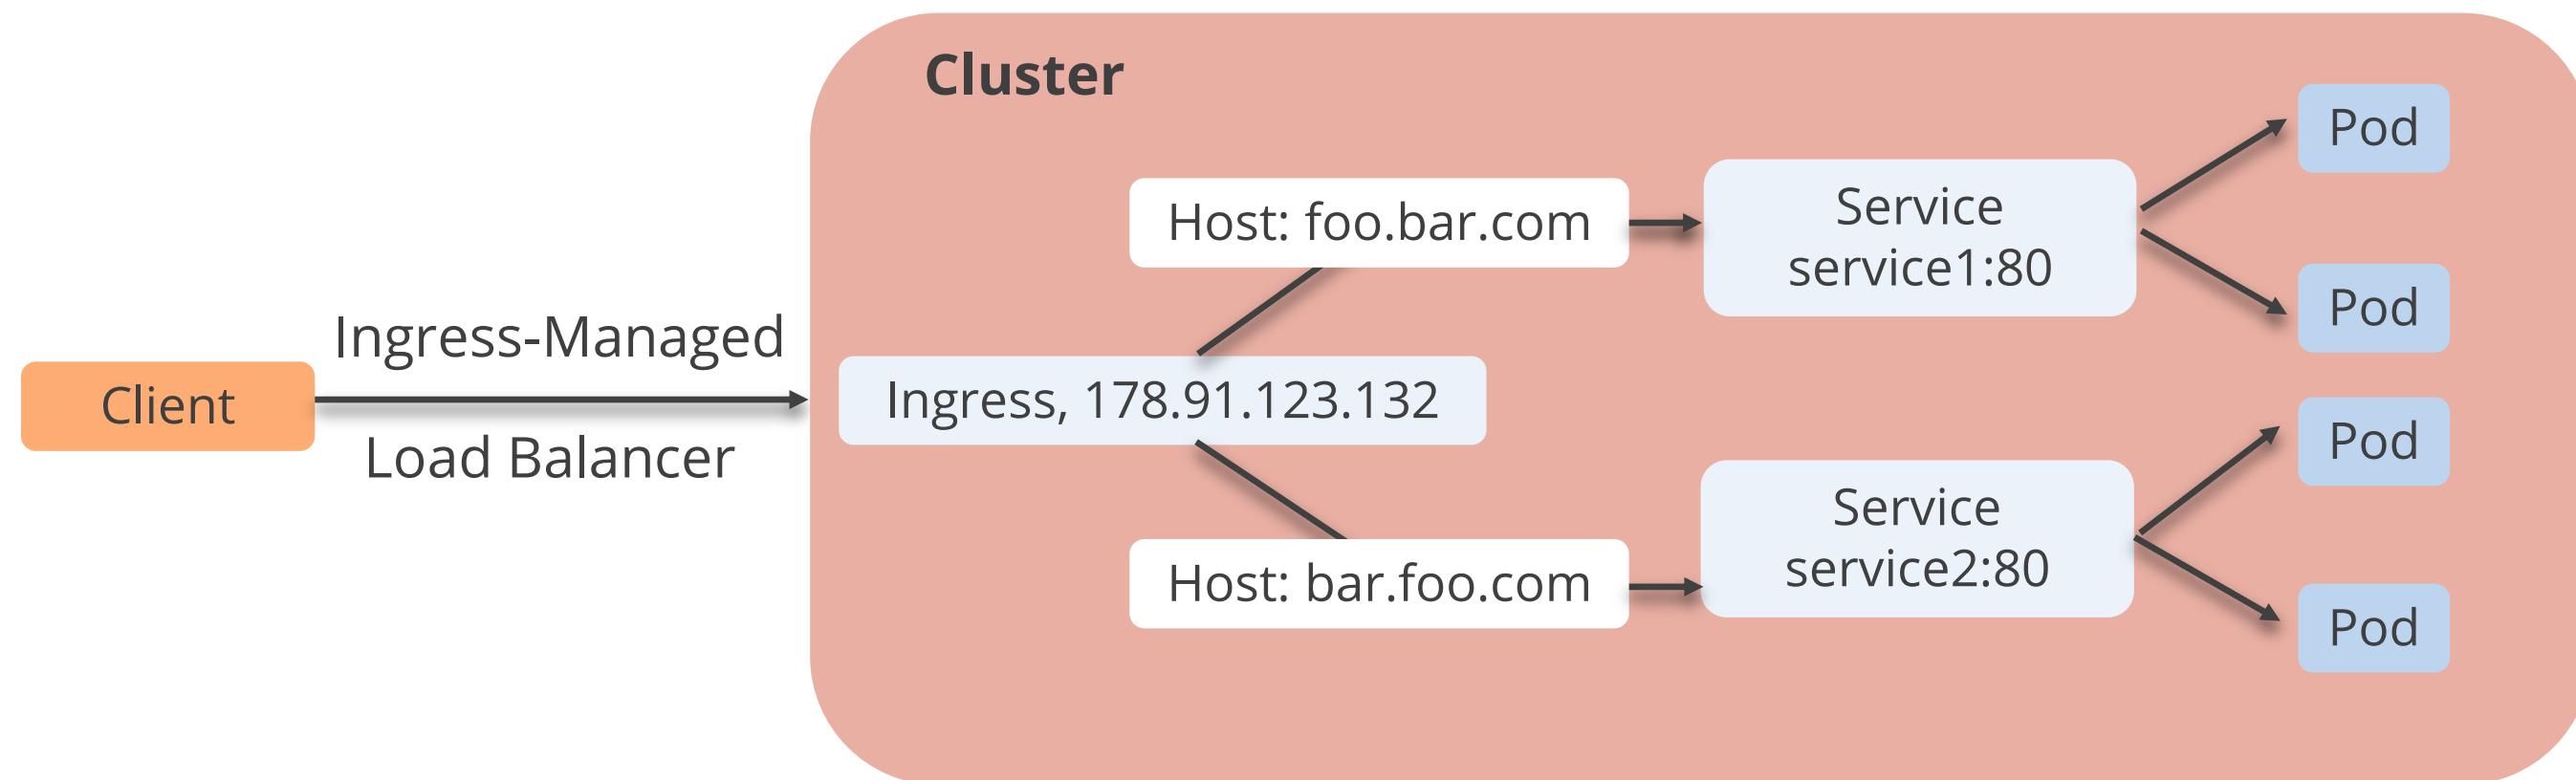
```
apiVersion: networking.k8s.io/v1
Kind: Ingress
Metadata:
  name: name-virtual-host-ingress
spec:
  rules:
    - host: foo.bar.com
      http:
        paths:
          - pathType: Prefix
            path: "/"
            backend:
              service:
                Name: service1
                port:
                  number: 80
    - host: bar.foo.com
      http:
        paths:
          - pathType: Prefix
            path: "/"
            backend:
              service:
                Name: service2
                port:
                  number: 80
```

## Demo

```
#Command to describe Ingress:
kubectl describe ingress simple-fanout-example
Result:
Name: simple-fanout-example
Namespace: default
Address: 178.91.123.132
Default backend: default-http-backend:80 (10.8.2.3:8080)
Rules:
Host Path Backends
--- ---
foo.bar.com
  /foo  service1:4200 (10.8.0.90:4200)
  /bar  service2:8080 (10.8.0.91:8080)
Events:
Type Reason Age From
Message
---- -
Normal ADD 22s loadbalancer-controller
default/test
```

# Name-Based Virtual Hosting

Name-based virtual hosts support routing HTTP traffic to multiple host names at the same IP address.



# Name-Based Virtual Hosting

The Ingress shown below tells the backing Load Balancer to route requests based on the Host Header:

Demo

```
apiVersion: networking.k8s.io/v1
Kind:   Ingress
Metadata:
  name:  name-virtual-host-ingress
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            Name: service1
            port:
              number 80
```

Demo

```
- host: bar.foo.com
  http:
    paths:
    - pathType: Prefix
      path: "/"
      backend:
        service:
          Name: service2
          port: 80
```

## TLS

The `tls.crt` and `tls.key` contain the certificate and private key to be used for TLS. An Ingress can be secured using the configuration shown below:

Demo

```
apiVersion: v1
Kind: Secret
Metadata:
  name: testsecret-tls
  namespace: default
  data
    tls.crt: base64 encoded cert
    tls.key: base64 encoded key
    type: kubernetes.in/tls
```

## TLS

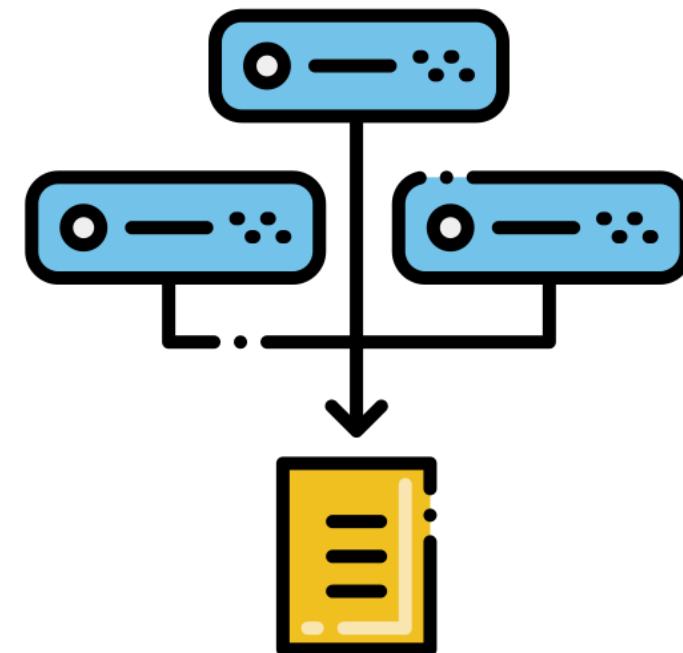
Ensure that the TLS secret created has come from a certificate that contains a Common Name (CN), also known as a Fully Qualified Domain Name (FQDN) for https-example.foo.com.

Demo

```
apiVersion: networking.k8s.io/v1
Kind:   Ingress
Metadata:
  name:  tls-example-ingress
spec:
  tls:
  - host:
      - https-example.foo.com
    rules:
    - host: bar.foo.com
      http:
        paths:
        - path: /
          - pathType: Prefix
        backend:
          service:
            Name: service2
          port:
            number 80
```

# Load Balancing

Load Balancing settings are bootstrapped to an Ingress Controller. These settings can thus apply to all Ingresses.



Ingress does not expose health checks directly. They can be exposed using readiness probes.

# Update an Ingress

To update an existing Ingress to be added to a new Host, edit the resource.

Demo

Command:  
kubectl describe ingress test

Result:  
Name: test  
Namespace: default  
Address: 178.91.123.132

Demo

```
spec:  
rules:  
- host: foo.bar.com  
  https:  
    paths:  
    - backend  
      service:  
        Name: service1  
        port:  
          number 80  
          path: /foo  
    - pathtype: Prefix  
      host: foo.bar.com  
      https:  
        paths:  
        - backend  
          service:  
            Name: service2  
            port:  
              number 80  
              path: /foo  
              pathtype: Prefix
```

# Understanding the Configuration of Ingress



**Problem Statement:** Learn how to configure Ingress in Kubernetes.

# Assisted Practice: Guidelines

---

## Steps to demonstrate Ingress in Kubernetes:

1. Create a minimal Ingress configuration
2. Create a resource
3. Check the created Ingress
4. Describe the Ingress

## Ingress Controllers

## Overview

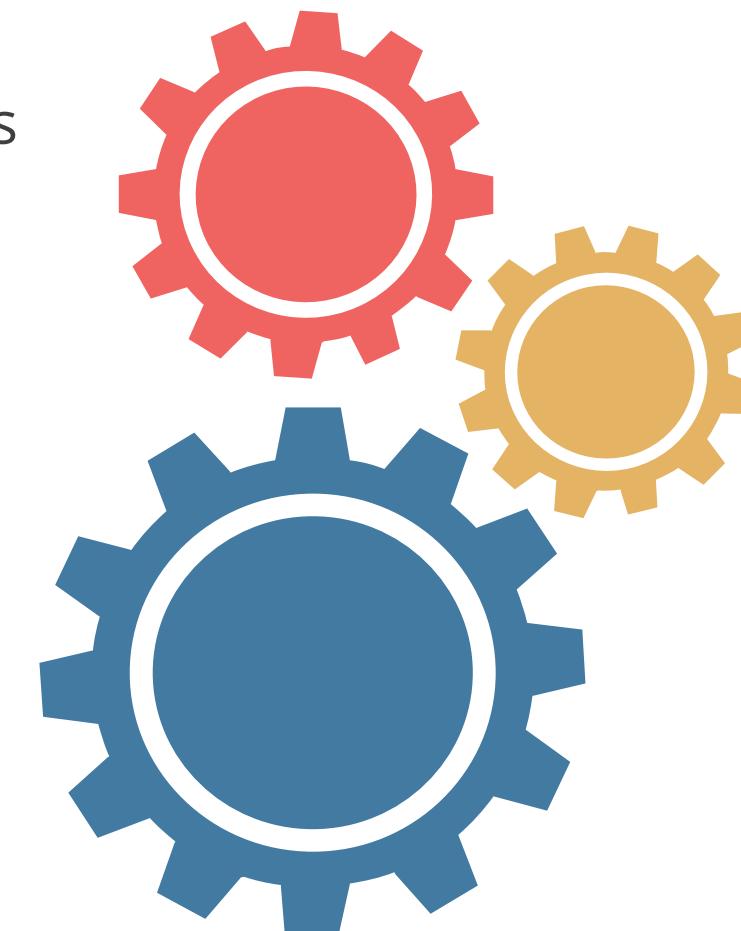
For an Ingress resource to work, the cluster must have an Ingress Controller running.

Kubernetes as a project supports and maintains three Controllers.

AWS Ingress Controllers

GCE Ingress Controllers

nginx Ingress Controllers



# Additional Controllers

Third-party Ingress Controllers include:

1 AKS Application Gateway Ingress Controller for Azure

2 Ambassador API Gateway

3 Apache APISIX Ingress Controller

4 Avi Kubernetes Operator for Load Balancing

5 The Citrix Ingress Controller

6 Contour

7 EnRoute

8 Gloo, an open-source Ingress Controller based on Envoy

# Additional Controllers

9

HAProxy Ingress is an ingress

10

The HAProxy Ingress Controller

11

Skipper HTTP Router and Reverse  
Proxy

12

The Kong Ingress Controller

13

portworxVolume

14

Istio Ingress

15

Voyager

16

The Traefik Kubernetes Ingress  
Provider

# Working with Ingress Controllers



**Problem Statement:** Understand the working of Ingress Controllers in Kubernetes.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

---

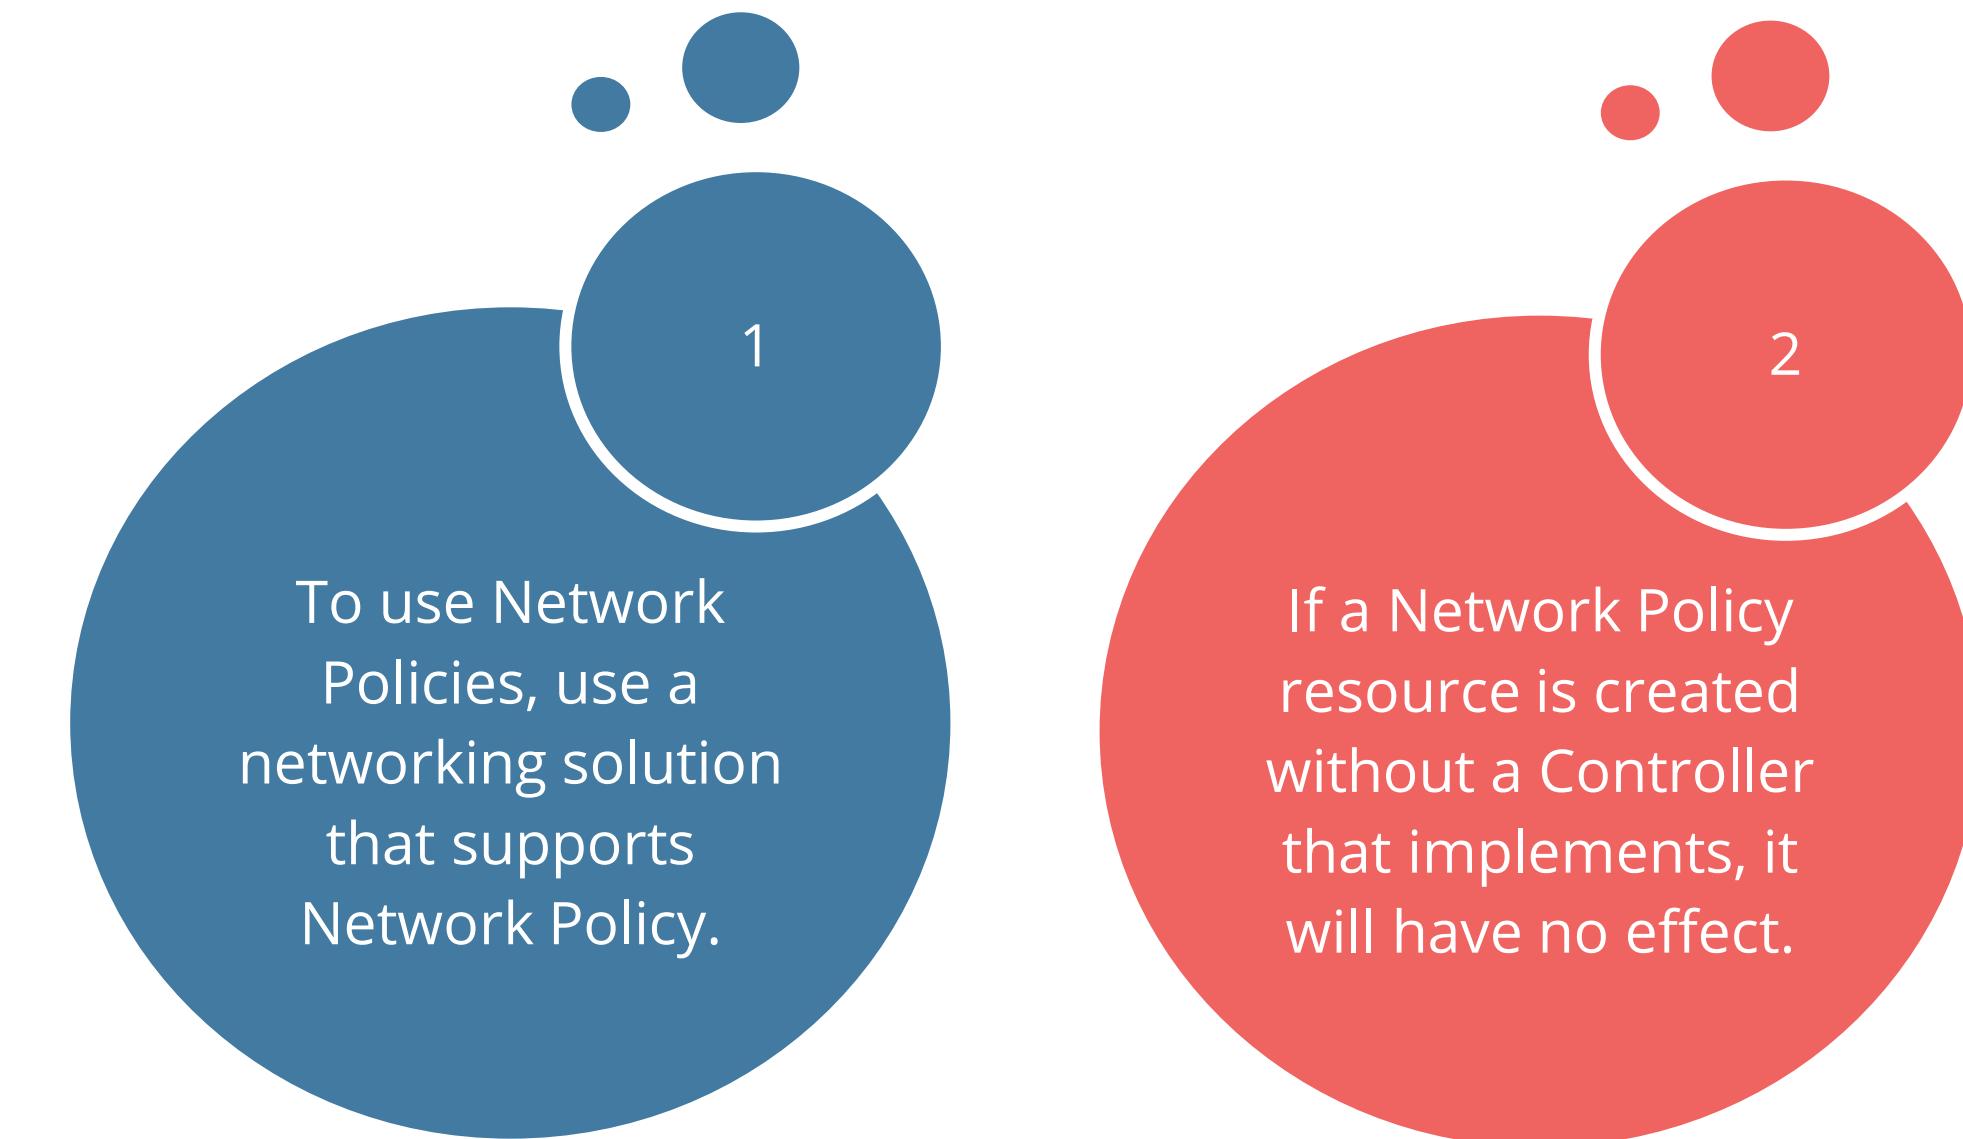
## Steps to demonstrate Ingress Controllers in Kubernetes:

1. Install Ingress Controllers on Bare Metal
2. Verify Installation

## Network Policies

# Network Policies

Network Policies are implemented by the network plugin.

- 
- 1 To use Network Policies, use a networking solution that supports Network Policy.
  - 2 If a Network Policy resource is created without a Controller that implements, it will have no effect.

# Introduction

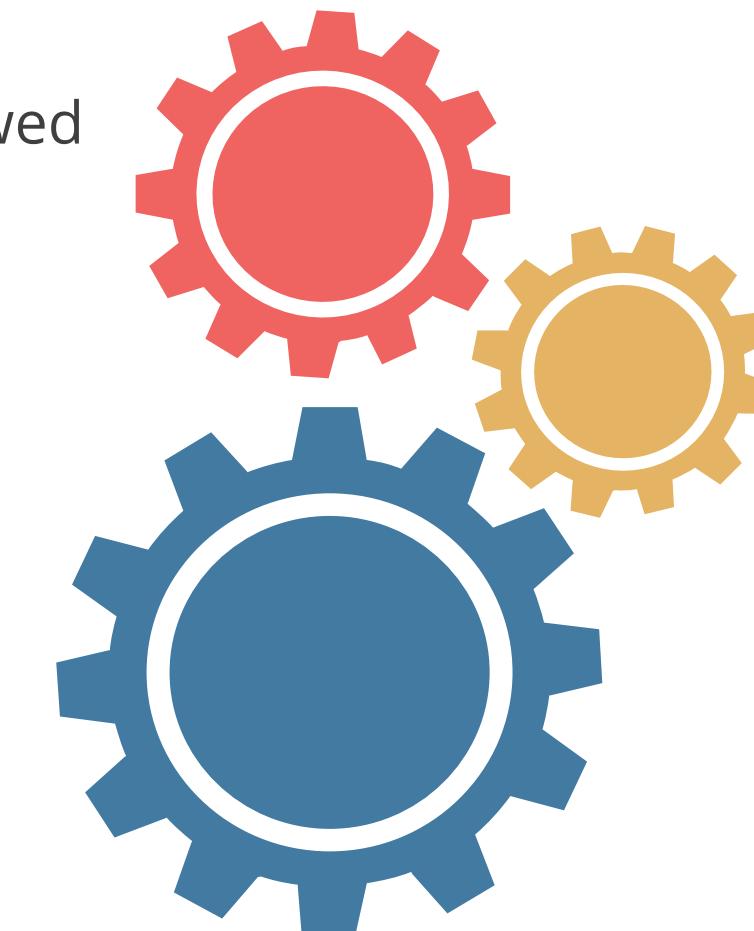
Network Policies are application-centric constructs that enable specifying how a Pod is allowed to communicate with various network entities over the network.

The entities that a Pod can communicate with are identified through a combination of three Identifiers.

Other Pods that are allowed

IP Blocks

Namespaces that are allowed



# Isolated and Non-isolated Pods

By default, Pods are non-isolated; they accept traffic from any source.

When there is a NetworkPolicy in a namespace selecting a particular Pod, the Pod will reject all connections not allowed by the NetworkPolicy.

If a policy (or policies) select(s) a Pod, the Pod is restricted to what is allowed by the union of the Ingress/Egress rules of the policy (or policies).

For a network flow to happen between two Pods, both the Egress policy on the source Pod and the Ingress policy on the destination Pod should allow the traffic.

# Network Policy Resource: Example

Demo

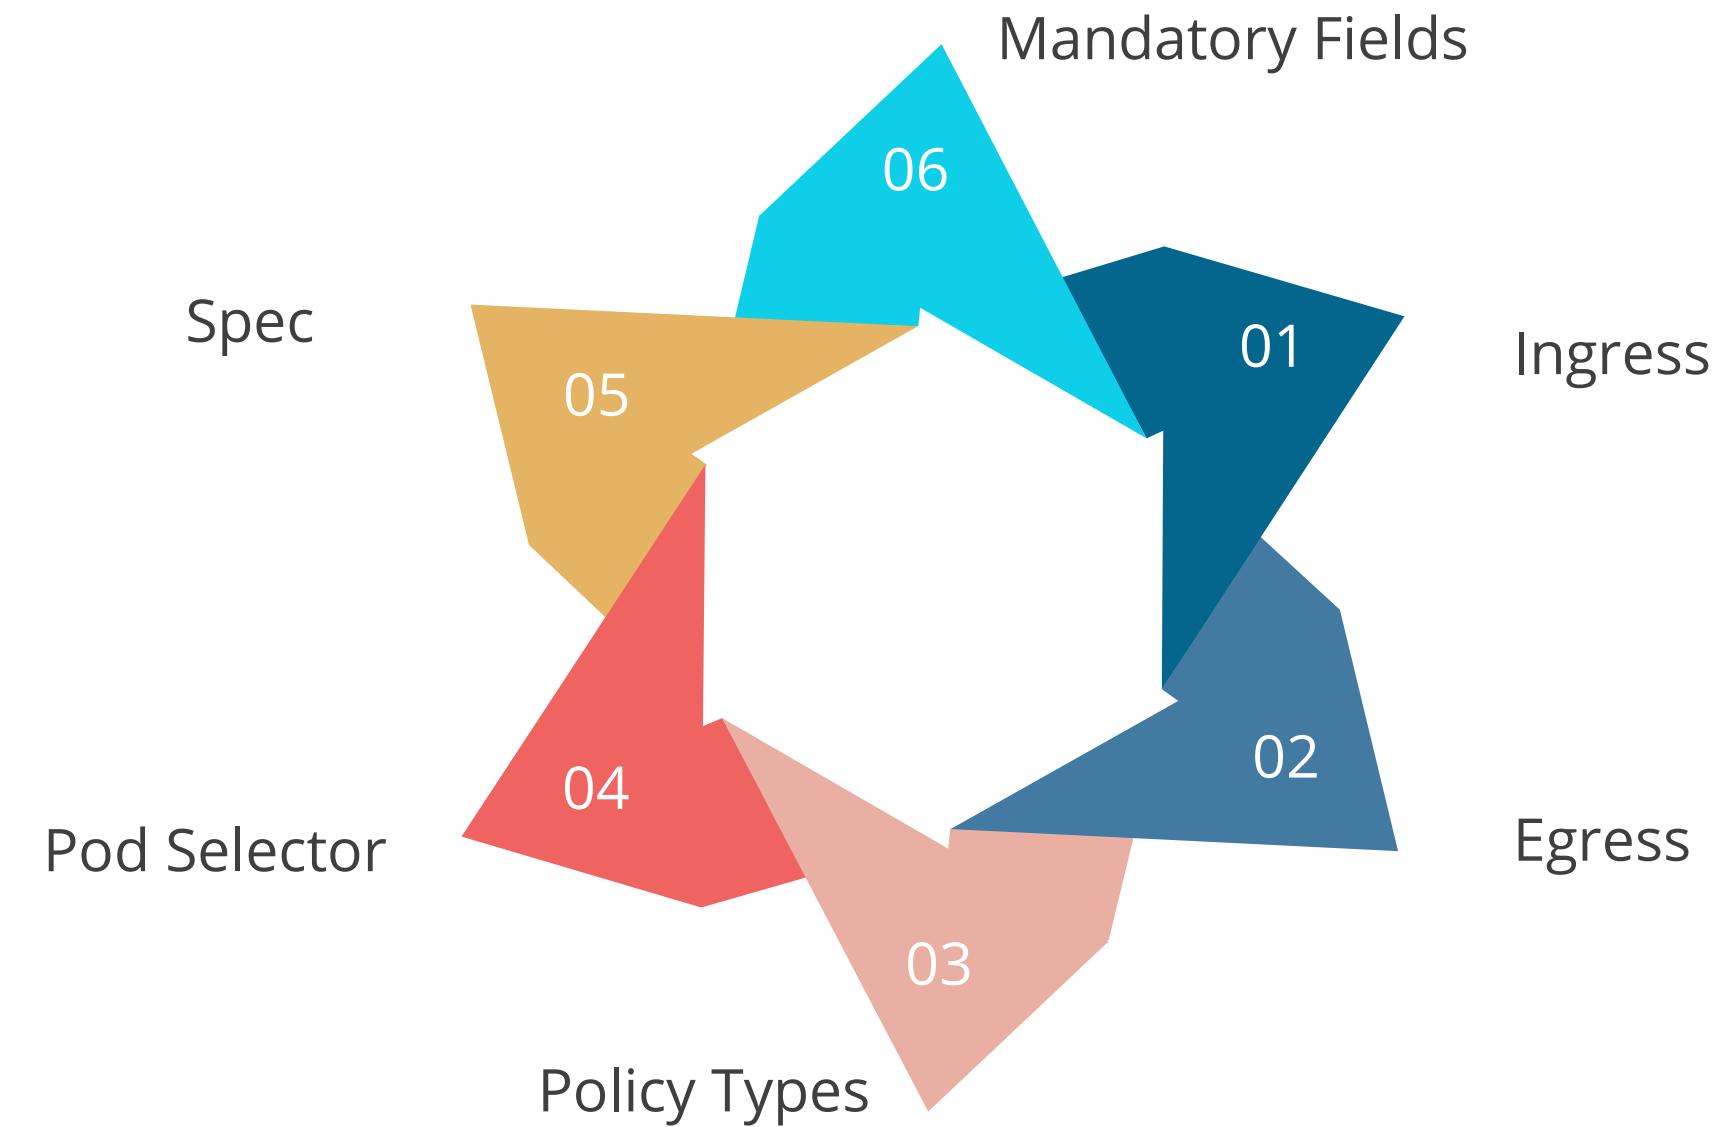
```
apiVersion: networking.k8s.io/v1
Kind: Ingress
Metadata:
  name: test-network-policy
  namespace: default
spec:
  Podselector:
    matchLabels:
      role: db
    policytypes:
      - Ingress
      - Egress
      - ingress:
          - from:
              - ipBlock: 172.17.1.0.0/16
                cidr:
                  except
                  - 172.17.1.0.0/16
    - namespaceSelector:
```

Demo

```
matchLabels:
  project: myproject
  - PodSelector:
    matchLabels:
      role: frontend
      ports:
        - protocol: TCP
          port: 6379
    egress:
      - to:
          - ipBlock
            cidr : 10.0.0.0/24
    ports:
      - protocol: TCP
        - port: 5978
```

# Network Policy Resource

The configuration file has six attributes:



# Behavior of To and From Selectors

## namespaceSelector and podSelector

A single To or From entry, which specifies both namespaceSelector and PodSelector, selects particular Pods within particular namespaces

## PodSelector

Selects Pods that should be allowed as Ingress or Egress destinations within the same namespace as the NetworkPolicy

## namespaceSelector

Selects namespaces for which all Pods should be allowed as Ingress sources or Egress destinations

## ipBlock

Selects cluster-external IP CIDR ranges that can be allowed as Ingress sources or Egress destinations

# Behaviour of To and From Selectors

Demo

```
...
  ingress:
  - fromd:
    namespaceSelector:
      matchLabels:
        user: alice
        PodSelector:
          matchLabels:
            role: client
...
...
```

Policy contains a single From element allowing connections from Pods with the label role=client in namespaces with the label user=alice.

# Default Policies

By default, if no policies exist in a namespace, then all Ingress and Egress traffic is allowed to and from Pods in that namespace.

Demo

```
...  
apiVersion: networking.k8s.10/v1  
Kind: NetworkPolicy  
Metadata:  
  name: default-deny-ingress  
spec:  
  Podselector: {}  
  policyTypes:  
  - Ingress
```

Default **deny all** Ingress traffic

Demo

```
...  
apiVersion: networking.k8s.10/v1  
Kind: NetworkPolicy  
Metadata:  
  name: allow-all-ingress  
spec:  
  Podselector: {}  
  Ingress:  
  - {}  
  policyTypes:  
  - Ingress
```

Default **allow all** Ingress traffic

# Default Policies

Demo

```
...  
apiVersion: networking.k8s.10/v1  
Kind: NetworkPolicy  
Metadata:  
  name: default-deny-egress  
spec:  
  Podselector: {}  
  Ingress:  
    - {}  
  policyTypes:  
    - Ingress
```

Demo

```
...  
apiVersion: networking.k8s.10/v1  
Kind: NetworkPolicy  
Metadata:  
  name: allow-all-egress  
spec:  
  Podselector: {}  
  egress:  
    - {}  
  policyTypes:  
    - egress
```

Demo

```
...  
apiVersion: networking.k8s.10/v1  
Kind: NetworkPolicy  
Metadata:  
  name: default-deny-all  
spec:  
  Podselector: {}  
  policyTypes:  
    - egress  
    - Ingress
```

Default **deny all** Egress traffic

Default **allow all** Egress traffic

Default **deny all** Ingress and all Egress traffic

# Working with Network Policy Resources



**Problem Statement:** Create and work with network policies in Kubernetes.

ASSISTED PRACTICE

# Assisted Practice: Guidelines

---

## Steps to demonstrate Network Policies in Kubernetes:

1. Create a network policy
2. Create a network policy configuration
3. Create the resource
4. Verify the installation

## Adding Entries to Pod /etc/hosts With HostAliases

# Default Hosts File Content

Adding entries to a Pod's etc or hosts file provides Pod-level override of hostname resolution when DNS and other options are not applicable. These custom entries can be added with the HostAliases field in PodSpec.

## Demo

Command to start a Nginx Pod that is assigned to a Pod IP is as follows:  
kubectl run nginx --image nginx

**Result:**

Pod/nginx created

Examine a Pod IP - Command:

kubectl get Pods --output=wide

**Result:**

| NAME  | READY | STATUS  | RESTARTS | AGE | IP         | NODE    |
|-------|-------|---------|----------|-----|------------|---------|
| nginx | 1/1   | Running | 0        | 13s | 10.200.0.4 | worker0 |

# Default Hosts File Content

This is how the Host File Content will look:

Demo

Command:

```
kubectl exec nginx -- cat /etc/hosts
```

Result:

```
# Kubernetes-managed hosts file.  
127.0.0.1      localhost  
::1      localhost ip6-localhost ip6-loopback  
fe00::0      ip6-localnet  
fe00::0      ip6-mcastprefix  
fe00::1      ip6-allNodes  
fe00::2      ip6-allrouters  
10.200.0.4    nginx
```

# Additional Entries with HostAliases

In addition to the default boilerplate, more entries can be made to the hosts file. For example, to resolve foo.local, bar.local to 127.0.0.1 and foo.remote, bar.remote to 10.1.2.3, configure HostAliases for a Pod under .spec.hostAliases.

## Demo

```
apiVersion: v1
Kind: Pod
Metadata:
  name: hostaloases-Pod
spec:
  restartPolicy: Never
  hostAliases:
    - ip: "127.0.0.1"
      hostnames:
        - "foo.local"
        - "bar.local"
    - ip: "10.1.2.3"
      hostnames:
        - "foo.remote"
        - "bar.remote"
  containers:
    - name: cat-hosts
      image: busybox
      command:
        - cat
      args:
        - "/etc/hosts"
```

# Additional Entries with HostAliases

A Pod can be configured by running the following command:

## Demo

```
kubectl apply -f https://k8s.io/examples/service/networking/hostaliases-Pod.yaml
```

Result:

Pod/hostaliases-Pod created

Examine a Pod's details to see its IPv4 address and its status, using the following command:kubectl get Pod --output=wide

Result:

| NAME            | READY | STATUS    | RESTARTS | AGE | IP         | NODE    |
|-----------------|-------|-----------|----------|-----|------------|---------|
| hostaliases-Pod | 0/1   | Completed | 0        | 6s  | 10.200.0.5 | worker0 |

The hosts file content looks like this:

```
kubectl logs hostaliases-Pod
```

Result:

```
# Kubernetes-managed hosts file.
```

```
127.0.0.1      localhost
::1      localhost ip6-localhost ip6-loopback
fe00::0  ip6-localnet
fe00::0  ip6-mcastprefix
fe00::1  ip6-allNodes
fe00::2  ip6-allrouters
10.200.0.5      hostaliases-Pod
```

```
# Entries added by HostAliases.
```

```
127.0.0.1      foo.local      bar.local
10.1.2.3  foo.remote    bar.remote
```

with the additional entries specified at the bottom.

# Adding Entries to `pod_etc_hosts` with HostAliases



**Problem Statement:** Interpret how to add entries to Pod, etc, or hosts with Host Aliases in Kubernetes.

## Assisted Practice: Guidelines

### **Steps to demonstrate adding entries to Pod, etc, or hosts with HostAliases in Kubernetes:**

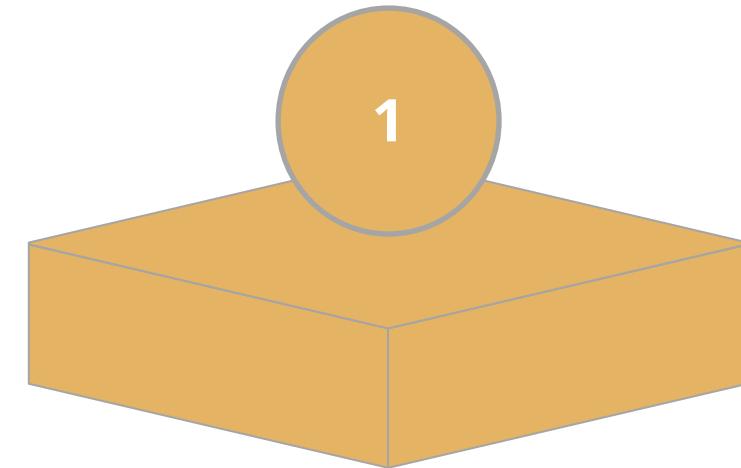
How to add entries to resources with host aliases:

1. Get Default hosts for the Pod webserver that we created earlier
2. Create a configuration with additional entries
3. Create the resource
4. Examine the Pod details
5. Observe the logs

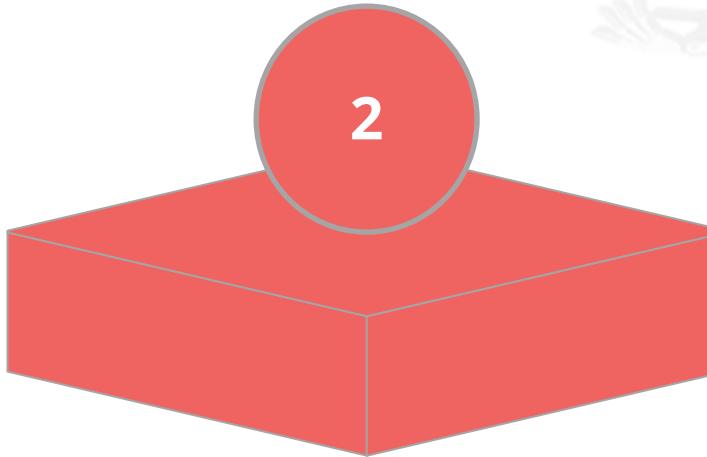
## IPv4/IPv6 Dual-Stack

# Introduction

IPv4/IPv6 Dual-Stack networking enables the allocation of both IPv4 and IPv6 addresses to Pods and Services.

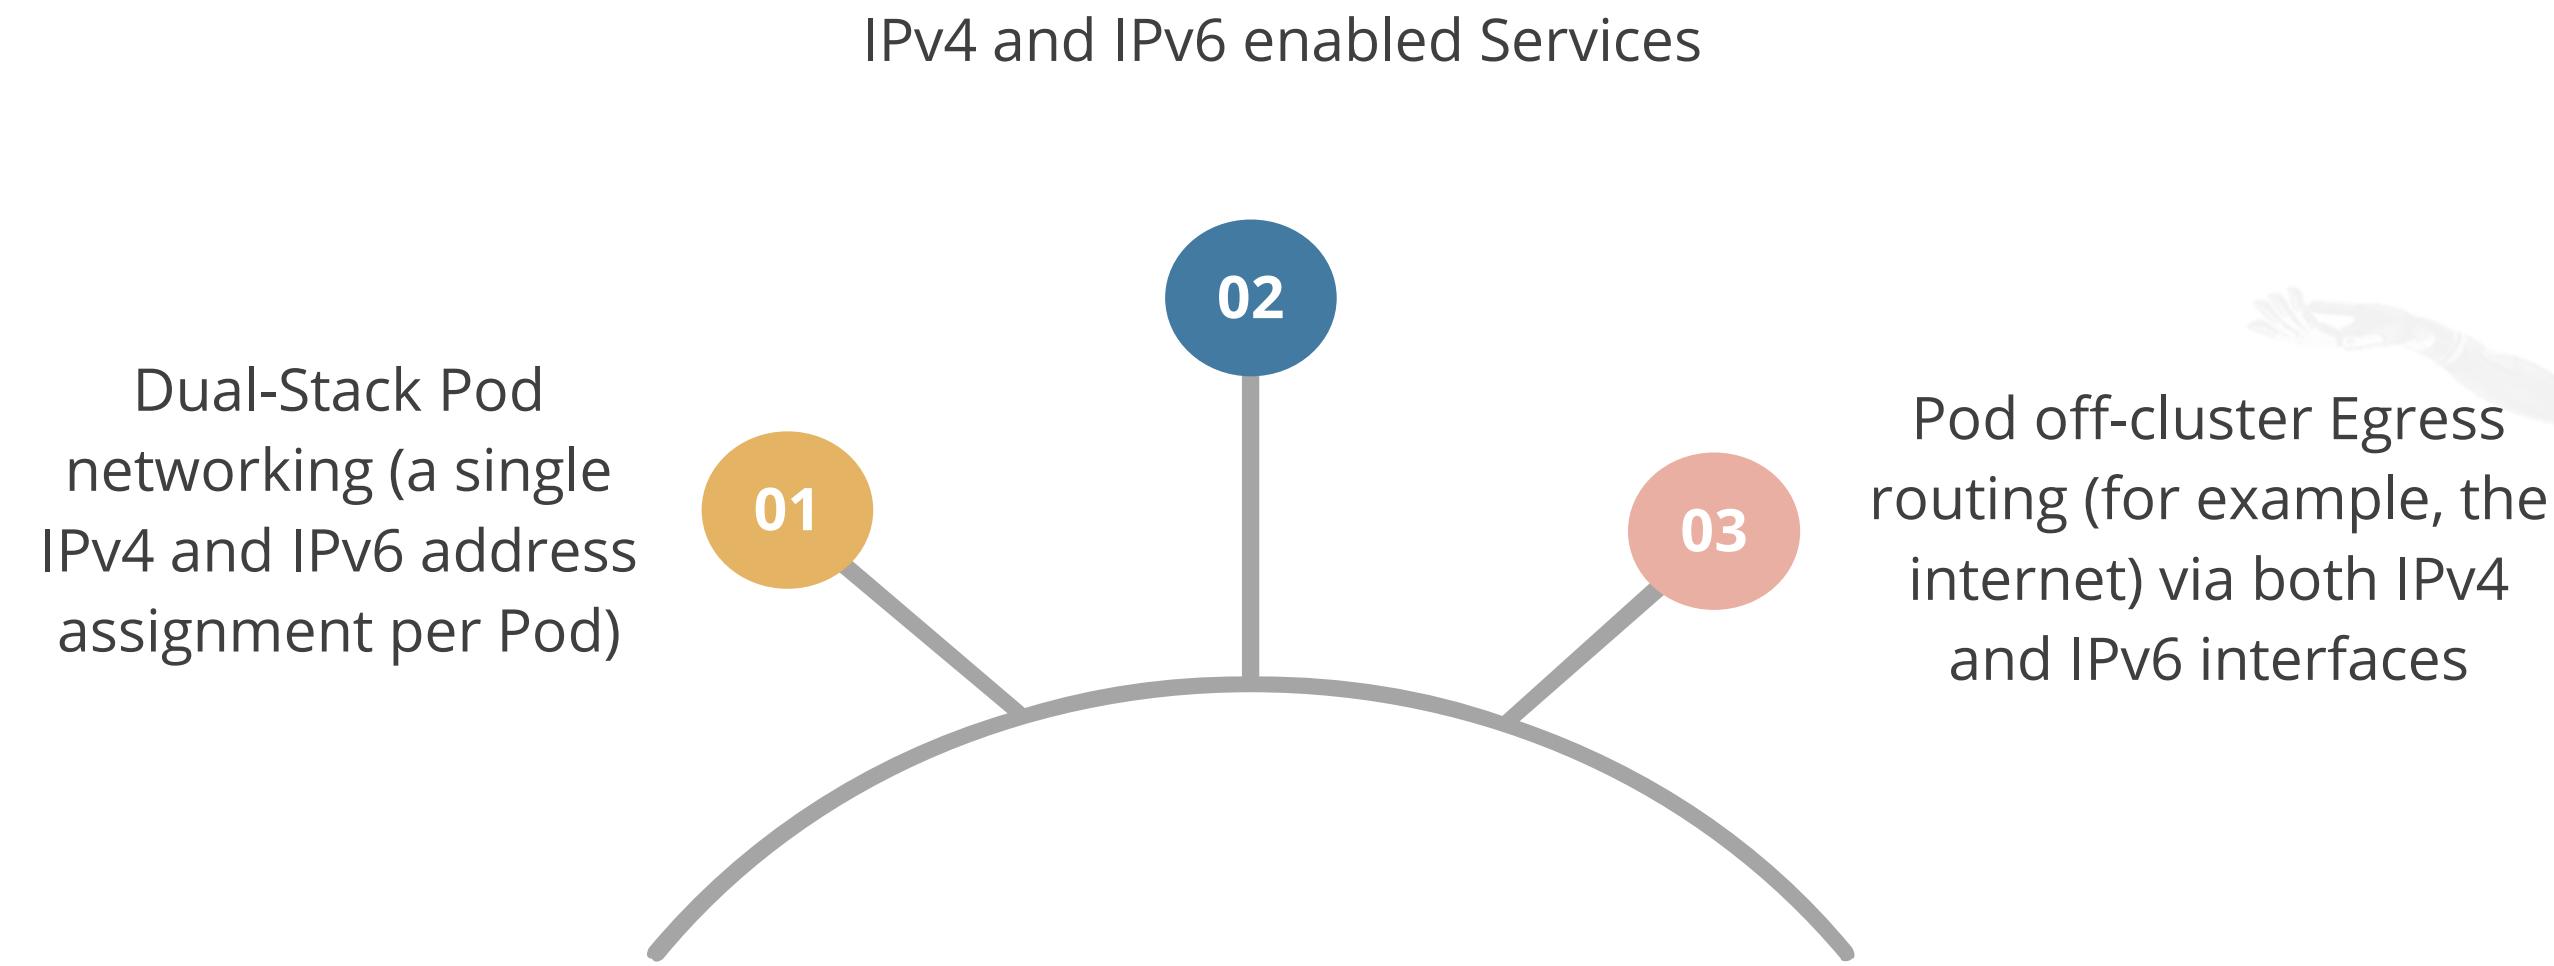


It is enabled by default for the Kubernetes cluster starting in 1.21, allowing simultaneous assignment of both IPv4 and IPv6 addresses.



# Supported Features

The IPv4/IPv6 Dual-Stack on the Kubernetes cluster provides the following Services:

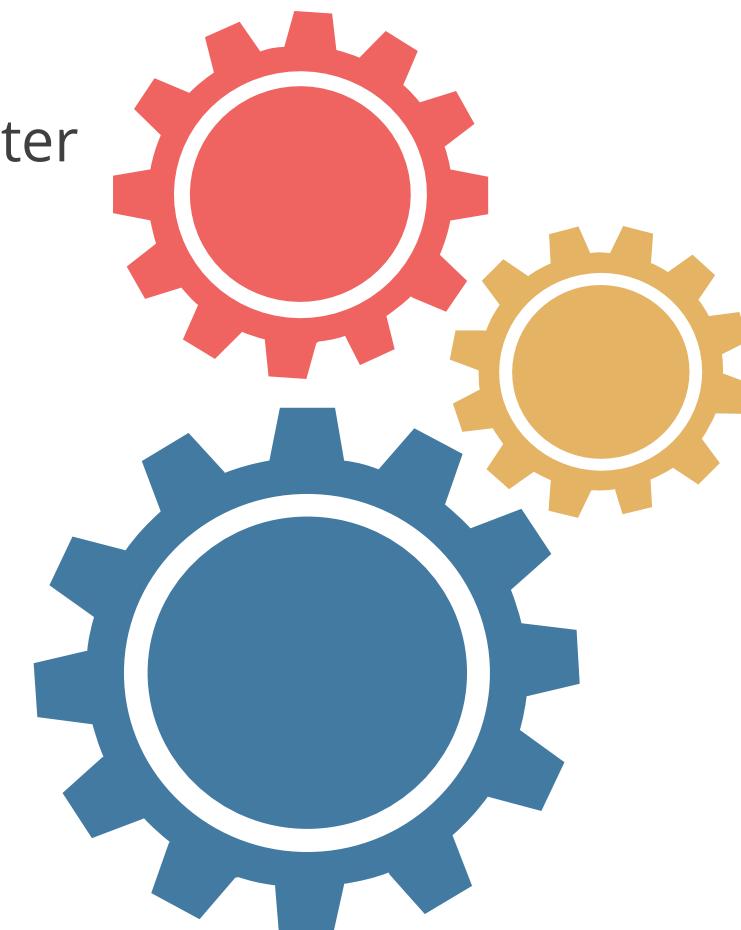


# Prerequisites

There are three prerequisites for utilizing IPv4/IPv6 Dual-Stack Kubernetes clusters:

Kubernetes 1.20 or later

A network plugin that supports Dual-Stack Configure IPv4/IPv6 dual-stack

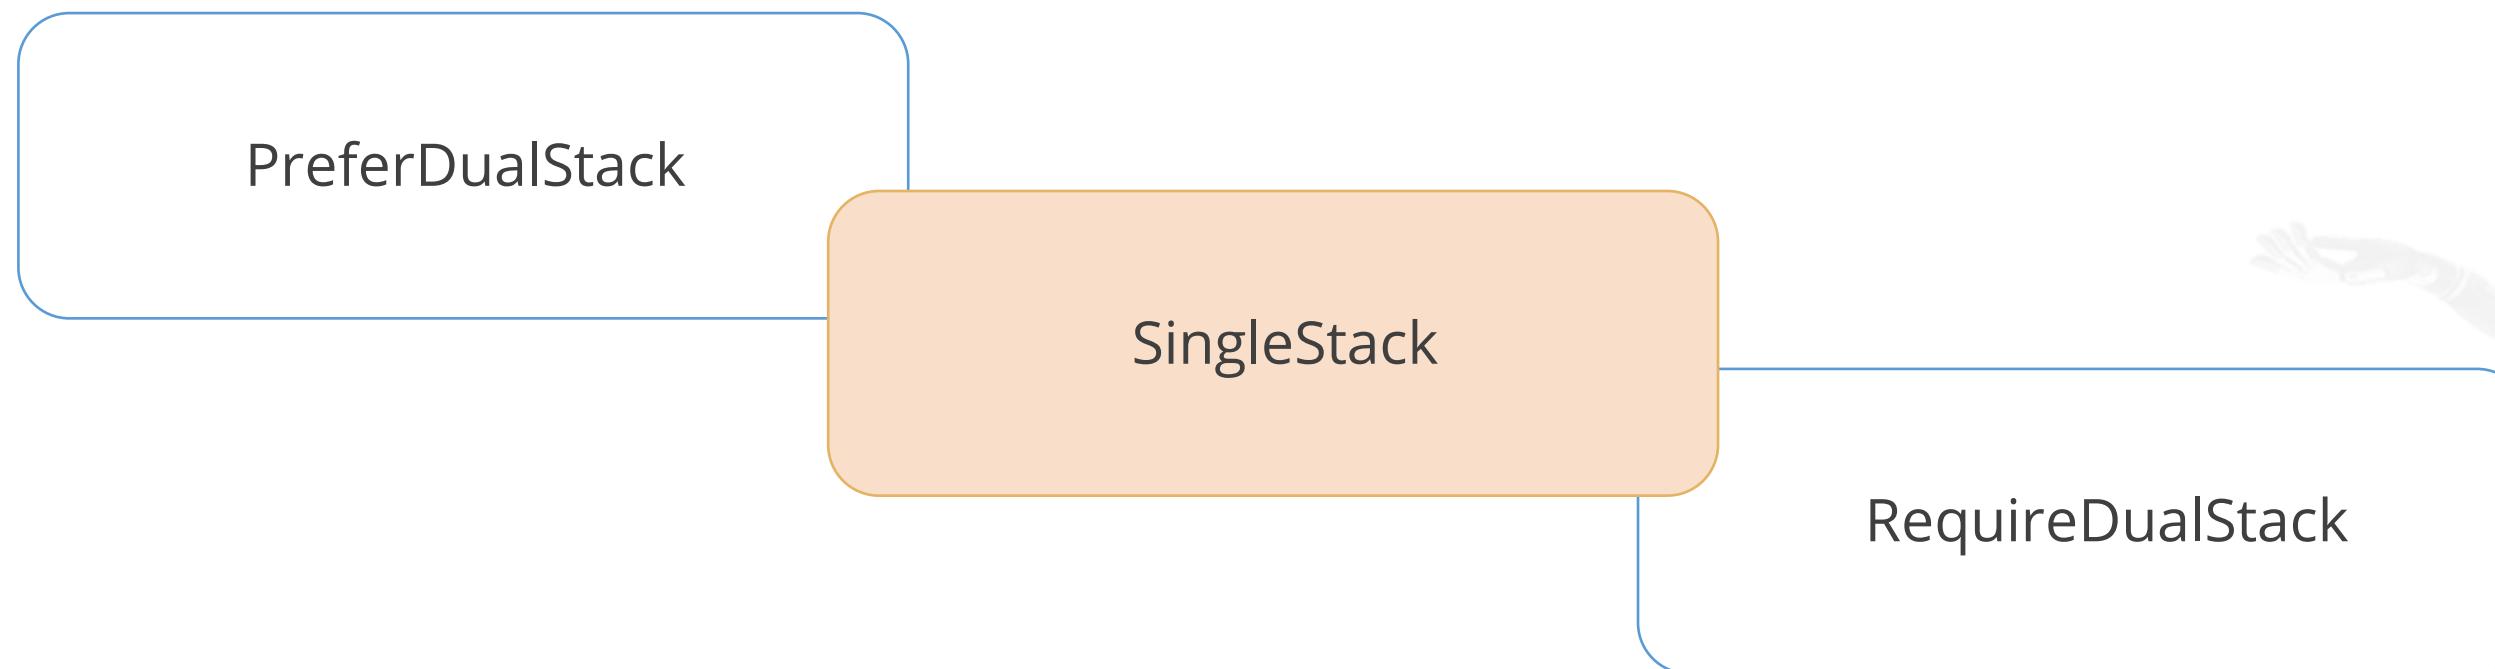


Provider support for Dual-Stack networking

# Services

When a Service is defined, it can be optionally configured as dual-stack.

**.spec.ipFamilyPolicy** field can be used to specify the behavior. It takes one of the following values:



# Dual-Stack Options on New Services

This Service specification does not explicitly define `.spec.ipFamilyPolicy`. Kubernetes **assigns a cluster IP** for the **Service** from the first configured service-cluster-ip-range after the Service is created. It sets the `.spec.ipFamilyPolicy` to SingleStack.

Demo

```
apiVersion: v1
Kind:   service
Metadata:
  name:  my-service
  labels:
    app: MyApp
spec: IPv4
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
```

# Dual-Stack Options on New Services

The Service specification shown below explicitly defines **PreferDualStack** in `.spec.ipFamilyPolicy`.  
`.spec.ClusterIPs` is the primary field and `.spec.ClusterIP` is a secondary field.

Demo

```
apiVersion: v1
Kind:   service
Metadata:
  name:  my-service
  labels:
    app: MyApp
spec:
  ipFamilyPolicy: PreferDualStack
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
```

# Dual-Stack Options on New Services

The Service specification shown below explicitly defines IPv6 and IPv4 in `.spec.ipFamilies` and `PreferDualStack` in `.spec.ipFamilyPolicy`:

```
Demo

apiVersion: v1
Kind:   service
Metadata:
  name:  my-service
  labels:
    app: MyApp
spec:
  ipFamilyPolicy: PreferDualStack
  ipFamilies:
  - IPv6
  - IPv4
  selector:
    app: MyApp
  ports:
  - protocol: TCP
    port: 80
```

# Dual-Stack Defaults on Existing Services

The examples shown below demonstrate the default behavior of Dual-Stack when it is newly enabled on a cluster with Services:

## Demo

```
apiVersion: v1
  Kind:   service
  Metadata:
    name:  my-service
    labels:
      app: MyApp
  spec:
    selector:
      app: MyApp
    ports:
      - protocol: TCP
        port: 80
```

#Validate this behavior by using kubectl to inspect an  
#existing service, using the following command:

```
kubectl get svc my-service -o yaml
```

## Demo

```
apiVersion: v1
  Kind:   service
  Metadata:
    name:  my-service
    labels:
      app: MyApp
  spec:
    clusterIPs: 10.0.197.123
    clusterIPs:
      - 10.0.197.123
    ipFamilies:
      - IPv4
    ipFamilyPolicy: singleStack
    ports:
      - port: 80
        protocol: TCP
        targetPort: 80
    selector:
      app: MyApp
    type: ClusterIP
  status:
    LoadBalancer: {}
```

# Dual-Stack Defaults on Existing Services

Upon enabling dual-stack on a cluster, existing Headless Services with Selectors are configured by the Control Plane. **.spec.ipFamilyPolicy** will be set to SingleStack and **.spec.ipFamilies** will be set to the address family of the first service cluster IP range.

Demo

```
apiVersion: v1
Kind:   service
Metadata:
  name:  my-service
  labels:
    app: MyApp
  spec:
    selector:
      app: MyApp
    ports:
      - protocol: TCP
        port: 80
```

# Dual-Stack Defaults on Existing Services

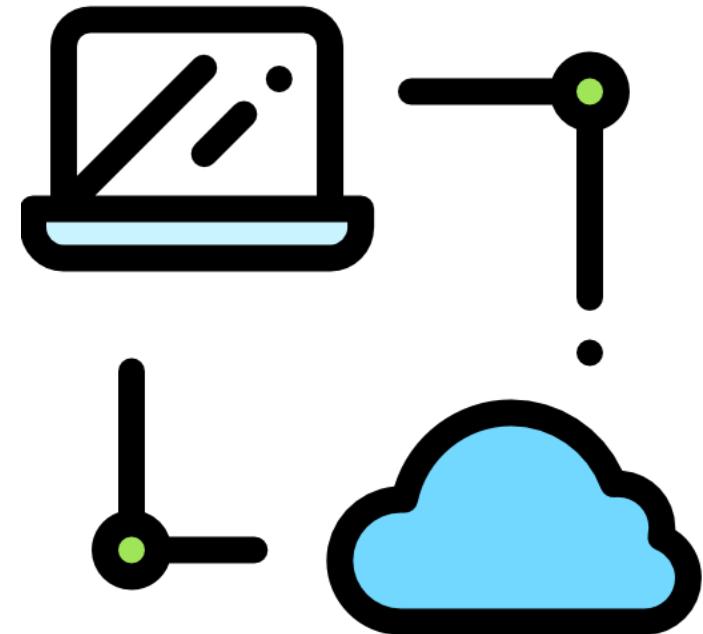
The behavior can be validated by using kubectl to inspect an existing Headless Service with Selectors, using the **kubectl get svc my-service -o yaml** command.

Demo

```
apiVersion: v1
Kind:   service
Metadata:
  labels:
    app: MyApp
  name:  my-service
spec:
  clusterIP: None
  clusterIPs:
  - None
  ipFamilies:
  - IPv4
  ipFamilyPolicy: singleStack
  ports:
  - port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: MyApp
```

## Egress Traffic

To enable Egress traffic in order to reach off-cluster destinations (for example, the public internet) from a Pod that uses non-publicly routable IPv6 addresses, enable the Pod to use a publicly routed IPv6 address.



# Switch Services Between Single-Stack and Dual-Stack

Services can be switched between Dual-Stack and Single-Stack Services and vice versa.

To change a Service from Single-Stack to Dual-Stack, change `.spec.ipFamilyPolicy` from `SingleStack` to **PreferDualStack** or **RequireDualStack** as desired.

Before:

```
spec:  
ipFamilyPolicy SingleStack
```

After:

```
spec:  
ipFamilyPolicy PreferDualStack
```

# Working with IPv4\_IPv6\_dual\_stack



**Problem Statement:** Understand the working of IPv4/IPv6 dual-stack in Kubernetes.

## Assisted Practice: Guidelines

---

### Steps to demonstrate IPv4/IPv6 dual-stack in Kubernetes:

1. Create a default dual-stack service
2. Create a preferred dual-stack service

## Cluster Networking

## Overview

Networking is a central part of Kubernetes, but it can be challenging to understand exactly how it is expected to work.

There are four distinct networking problems to address:

- 1 Highly-coupled Container-to-Container communications
- 2 Pod-to-Pod communications
- 3 Pod-to-Service communications
- 4 External-to-Service communications

# Kubernetes Network Model

Every Pod is assigned an IP address. Kubernetes imposes fundamental requirements on any networking implementation.

Kubernetes allows communication between:

1

Pods in the host network of a Node and Pods on all Nodes without NAT

2

Agents on a Node and all the Pods on that Node

# Implement the Kubernetes Networking Model

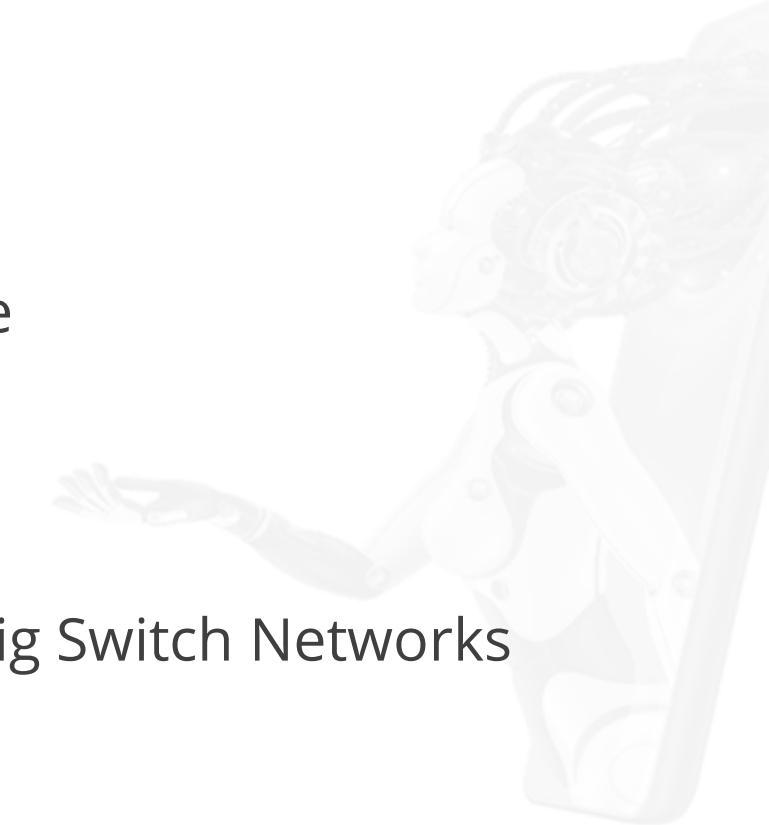
There are a number of ways in which this network model can be implemented:

- |                                       |                                     |
|---------------------------------------|-------------------------------------|
| <b>1</b>   AOS from Apstra            | <b>9</b>   Contrail/Tungsten Fabric |
| <b>2</b>   AWS VPC CNI for Kubernetes | <b>10</b>   DANM                    |
| <b>3</b>   Azure CNI for Kubernetes   | <b>11</b>   Knitter                 |
| <b>4</b>   Cilium                     | <b>12</b>   Kube-OVN                |
| <b>5</b>   CNI-Genie from Huawei      | <b>13</b>   Kube-router             |
| <b>6</b>   cni-ipvlan-vpc-k8s         | <b>14</b>   Multus                  |
| <b>7</b>   Coil                       | <b>15</b>   NSX-T                   |
| <b>8</b>   Contiv                     | <b>16</b>   Nuage Networks VCS      |



# Implement the Kubernetes Networking Model

- |                                |  |
|--------------------------------|--|
| 17   OpenVSwitch               | 23   OVN                                       |
| 18   ACI                       | 24   Antrea                                    |
| 19   Flannel                   | 25   Google Compute Engine                     |
| 20   Jaguar                    | 26   k-vswitch                                 |
| 21   Romana                    | 27   Big Cloud Fabric from Big Switch Networks |
| 22   Weave Net from Weaveworks |  |



## Key Takeaways

- A Kubernetes Service is a logical abstraction for a deployed group of Pods in a cluster.
- Kubernetes supports two primary modes of finding a Service, namely, environment variables and DNS.
- Network Policies are application-centric constructs that specify how a Pod is allowed to communicate with various network entities.
- IPv4/IPv6 Dual-Stack networking supports IPv4 and IPv6 addresses to be allocated to Pods and Services.



# Create a Service to Expose an Application Using EndpointSlices



## Problem Statement:

You have deployed a containerized application and now you have to expose it by creating a service. You also have to create EndpointSlices and add entries to Pods.

## Steps to Perform:

1. Connecting application with services
2. Creating EndpointSlices
3. Adding entries to Pods with host aliases